

“...self-amortizing canals...”

— Mr. Banks in *Mary Poppins* (the movie)

$$1 + 2 + 3 + 4 + \cdots = -\frac{1}{12}$$

$$\sqrt{1 + 2\sqrt{1 + 3\sqrt{1 + 4\sqrt{\cdots}}}} = 3$$

— Ramanujan (1887–1920)

## Lecture VI AMORTIZATION

Amortization is the algorithmic idea of distributing computational cost over a period of time. The terminology comes from home mortgages: most Americans pay for a home by taking out a long-term<sup>1</sup> loan called a home mortgage. This loan is repaid in monthly installments, over the period of the loan. In an amortized analysis of an algorithm, we likewise spread the cost of an operation over the entire run of the algorithm. Suppose each run of the algorithm amounts to a sequence of operations on a data structure. For instance, to sort  $n$  items, the well-known **heapsort algorithm** (§III.7) is a sequence of  $n$  **insert**’s into an initially empty priority queue, followed by a sequence of  $n$  **deleteMin**’s from the queue until it is empty. Thus if  $c_i$  is the cost of the  $i$ th operation, the algorithm’s running time is  $\sum_{i=1}^{2n} c_i$  since there are  $2n$  priority queue operations in all. In worst case analysis, we ensure that *each* operation is efficient, say  $c_i = O(\log n)$ , leading to the conclusion that the overall algorithm is  $O(n \log n)$ . But an **amortization argument** might be able to obtain the same bound  $\sum_{i=1}^{2n} c_i = O(n \log n)$  *without* ensuring that each  $c_i$  is logarithmic. In this case, we will say that the **amortized cost** of each operation is logarithmic. Thus “amortized complexity” is a kind of average complexity although it has nothing to do with probability. Tarjan [13] gave the first systematic account of this topic.

**¶1. Why amortize?** For the heapsort problem above, we could have ensured that *each* operation is logarithmic time. Nevertheless, we may find it advantageous to consider data structures that achieve logarithmic behavior only in the weaker amortized sense. The extra flexibility afforded by using amortized bounds often lead to simpler or more practical algorithms. Indeed many “amortized” data structures are relatively easy to implement. As a concrete example, consider any balanced binary search tree scheme. The algorithms for such trees must constantly perform considerable book-keeping to maintain its balanced shape. In particular, it must maintain some balance information at each node. In contrast, the splay trees in this chapter provide an amortization scheme for binary search trees that is considerably simpler, needing no extra storage at each node, and is quite lax about balancing. The operative word in amortization is “laziness” – try to defer the book-keeping work to the future, when it might be more convenient to do this work.

<sup>1</sup>Long-term means something like 15 to 30 years. Your home is said to be mortgaged, serving as collateral for the loan. This mortgaging analogy fails in some detail. We will see that amortization is closer to a banking model where you maintain a balance account that must not go negative.

*In algorithmics, we like to turn conventional vices (greediness, laziness, gambling with chance, etc) into virtues*

This chapter is in 3 parts: we begin by introducing the **potential function framework** for doing amortization analysis. Then we introduce two data structures which can be analyzed using this framework: **splay trees** and **Fibonacci heaps**. We give non-trivial applications of these data structures. Splay trees can be used in dynamic string compression and for maintaining the convex hull of a planar point set. Fibonacci heaps can be used to give the optimal implementation of Prim's minimum spanning tree algorithm.

## §1. The Potential Framework

We formulate an approach to amortized analysis using the concept of “potential functions”. Borrowing a concept from Physics, we imagine data structures as storing “potential energy” that can be released to do useful work. First, we view a data structure such as a binary search tree as a persistent object, i.e., it exists in time and has a history. At each moment it has a state which can be changed by operations such as insert, delete, rotation, etc. The *characteristic property* of potential functions is that they are a function of the *current* state of the data structure alone; it is independent of the states that preceded the current state. Invariably the potential function depends only on the shape of the underlying data structure, a binary tree in this case. For instance, consider two binary search trees  $T_1$  and  $T_2$  of the same shape. The first is obtained by inserting the keys 1, 2, 3 (in this order) into an empty tree; the second is obtained by inserting the keys 3, 1, 4, 2, followed by deletion of 4. The potential of  $T_1$  and  $T_2$  will be the same, despite their different history.

¶2. A “Counter Example”. Let  $C$  be a binary counter and the only operation on  $C$  is to increment its value. Starting with a counter value of 0, our counter  $C$  goes through the following sequence of states as we repeatedly increment it:

$$(0) \rightarrow (01) \rightarrow (010) \rightarrow (011) \rightarrow (0100) \rightarrow (0101) \rightarrow (0110) \rightarrow (0111) \rightarrow \dots \quad (1)$$

Here, we use a convention of prepending a 0 bit to the standard binary notation for a positive integer.

Our problem is to bound the total cost of a sequence of  $n$  increments, starting from an initial counter value of 0. If the  $i$ th increment has cost  $c_i$ , then the total cost is  $\sum_{i=1}^n c_i$ . It is easy to see that  $c_i = O(\lg n)$  since the length of the binary numbers are  $O(\lg n)$ . This is also a lower bound on the worst case  $c_i$  since the increment  $(0 \underbrace{11 \dots 1}_k) \rightarrow (01 \underbrace{00 \dots 0}_k)$  costs  $\Theta(k)$ , and it is easily checked that  $k$  can reach  $\lg n$ . Using a worst-case analysis, we could only conclude that the total cost  $\sum_{i=1}^n c_i$  is  $O(n \lg n)$ .

We now give a better bound using amortized analysis. Assume  $C$  is represented by a linked list of 0 and 1's, with direct access to the head and tail of the list. This representation determines our **cost model**: *the cost of an increment operation is the number of bits of  $C$  that we need to flip*. Moreover, *the number of bits to flip is equal to the length of the unique suffix of  $C$  of the form  $01^*$* .

Note that  $01^*$  refers to any string that begins with a 0 and is followed by zero or more occurrences of 1. If the suffix of  $C$  is  $01^k$  for some  $k \geq 0$ , then the increment operation will change this suffix to  $10^k$ . E.g., if  $C = (0101)$  in (1), then its  $01^*$  suffix is 01. Since this suffix has 2 bits, it costs 2 units to incrementing  $C$  to (0111). By convention,  $C$  always begins with a 0-bit. So a suffix of the form  $01^*$  always exists, and it's length is at least 1. For instance, if the value of  $C$  is 27, then  $C = (011011)$ . After incrementing, we get  $C = (011100)$ . The cost is 3 since the suffix 011 is changed to 100.

*not your usual  
counter example...*

**Special rule:** if the prefix  $01^*$  of  $C$  is equal to  $C$  itself, then after flipping the bits of  $01^*$ , we also prepend a 0 bit to  $C$ . E.g., when we increment  $C = (0111)$ , the prefix is  $0111$ . After flipping the prefix bits we get  $C = (1000)$ , and this rule tells us to prepend a 0 bit to finally get  $C = (01000)$ . Note that this rule is applied in the very first increment, when we go from  $C = (0)$  to  $C = (01)$ . The prepending cost is  $O(1)$  and it can<sup>2</sup> be absorbed into the overall (positive) cost to flip bits. Therefore we simply say that the cost for incrementing  $C = (0111)$  is 4, not 5.

To begin our amortized analysis, we associate with  $C$  a **potential**  $\Phi(C)$  that is equal to the number of 1's in its list representation. In our preceding example where the value of  $C$  is 27, its potential is  $\Phi(011011) = 4$ . After incrementing  $C$ ,  $\Phi(C) = \Phi(011100) = 3$ . So the potential of  $C$  decreased by 1 after this particular increment. In fact the only increments in which the potential increases is when the suffix is 0. Informally, we “store”  $\Phi(C)$  units of work (or energy) in  $C$ . Each increment of  $C$  has a **cost** that is defined by our cost model above, but we independently **charge** the increment a certain number of work units. In our current simple example, we levy a constant charge of 2 units. It is an interplay among the three quantities

potential, cost and charge

that characterizes any particular amortization analysis.

To analyze the increment operation, we consider two cases.

- (i) Suppose the least significant bit of  $C$  is 0. Then the increment operation just changes this bit to 1. Note that the potential increases by 1 by this operation. We see that our charge of 2 units is just right – one unit to pay for the cost, and one unit to pay for the increase in potential.
- (ii) Suppose an increment operation changes a suffix  $\underbrace{0111 \cdots 11}_k$  of length  $k \geq 2$  into  $\underbrace{1000 \cdots 00}_k$ : the cost incurred is  $k$ . Notice that the potential  $\Phi$  decreases by  $k - 2$ . This decrease “releases”  $k - 2$  units of work that can pay for  $k - 2$  units of the incurred cost. The remaining 2 units of cost can be paid for by our charge of 2 units.

Thus, in both cases (i) and (ii), charging 2 units of work for an operation is sufficient to account for the cost. Summed over the  $n$  operations, the total charge of  $2n$  is sufficient to pay for the total cost of all the operations. We conclude that the amortized cost is  $O(1)$  per operation.

**§3. Abstract Formulation.** We now formulate an amortization analysis that captures the essence of the above Counter Example. Suppose we are analyzing the cost of a sequence

$$p_1, p_2, \dots, p_n \tag{2}$$

of **requests** on a data structure  $D$ . We view a data structure  $D$  as comprising two parts: its **contents** and its **structure**, where the structure represents some organization of the contents. The (structure, contents) pair is called the **state** of the data structure. The term “request” is

<sup>2</sup>Alternatively, we can separately account for the prepending cost: it is at  $\Theta(\log n)$  since we do this operation  $\lg n$  times.

meant to cover two types of operations: **updates** that modify the contents of  $D$ , and **queries** that computes a function of the contents but do not modify the contents. Thus there is no logical necessity to modify the structure of  $D$  in a query. *Nevertheless, it may be advantageous to modify the structure during queries.*

EXAMPLE 1 Let  $D$  be a binary search tree. The structure of  $D$  is the binary tree, and its contents are the keys stored in each node of the binary tree. Insertion and deletion operations are the update requests on  $D$ , while looking up a key is a query request. In Lookups, it is clear that the  $D$  need not change. Nevertheless, to ensure a favorable complexity over a sequence of such operations, we may perform some rotations: we saw this with AVL trees in Chapter III. But a simpler idea illustrating amortization is to simply rotate the searched-for node so that its depth is reduced! This is the “move-to-front” heuristic. ■

The data structure  $D$  is dynamically changing: each request transforms the current state of  $D$ . Let  $D_i$  be the state of the data structure *after* request  $p_i$ . Let  $D_0$  be the initial state.

We assume a **complexity model** whereby request  $p_i$  has a positive **cost**, denoted  $\text{COST}(p_i)$ . This problem before us is to bound the complexity of executing the sequence (2). To carry out an amortization argument, we must invent a **charging scheme** and a **potential function**. Unlike the complexity model, the charging scheme and potential function are not part of the problem specification. They are artifacts of our analysis and their invention usually requires some ingenuity.

A **charging scheme** is just any systematic way to associate a real number  $\text{CHARGE}(p_i)$  to each operation  $p_i$ . E.g., for our Counter Example, we define  $\text{CHARGE}(p_i) := 2$  for each  $p_i$ . Informally, we “levy” a **charge** of  $\text{CHARGE}(p_i)$  on the operation. We emphasize that this levy need not bear any obvious relationship to  $\text{COST}(p_i)$ . The **credit** of this operation is defined to be the “excess charge”,

$$\text{CREDIT}(p_i) := \text{CHARGE}(p_i) - \text{COST}(p_i). \quad (3)$$

E.g., suppose an operation has cost 3. If we charge this operation 5 units, then its credit is 2. But if we charge this operation 2 units, the credit is  $-1$ . In view of (3), specifying a charging scheme is equivalent to specifying a credit scheme. If the credit is a negative number, it is really a “debit”.

A **potential function** is a function  $\Phi$  that assigns a *non-negative* real number to every possible state of  $D$ . For simplicity, we require that

$$\Phi(D_0) = 0$$

where  $D_0$  is the initial state of  $D$ . Call  $\Phi(D_i)$  the **potential** of state  $D_i$ . Let the *increase in potential* of the  $i$ th step be denoted by

$$\Delta\Phi_i := \Phi(D_i) - \Phi(D_{i-1}). \quad (4)$$

The amortization analysis amounts to verifying the following inequality at every step:

$$\text{CREDIT}(p_i) \geq \Delta\Phi_i. \quad (5)$$

Call (5) the **credit-potential invariant**.

The idea is that credit is stored as “potential” in the data structure. Since the potential function and the charging scheme are defined independently of each other, the truth of the invariant (5) is not a foregone conclusion. It must be verified for each case.

*OK, we are mixing financial and physical metaphors here. Since mortgages is a financial term, perhaps the credit/debit ought to be put into a “bank account” and  $\Phi$  is the “current balance”.*

After the credit-potential invariant has been verified, we are allowed to assert that the charge for an operation is its **amortized cost**. This is justified by the following simple result:

**Theorem 1 (Justification of Amortized Cost)**

Assuming the credit-potential invariant (5), the total cost of the requests (2) is upper bounded by the total charge:

$$\sum_{i=1}^n \text{COST}(p_i) \leq \sum_{i=1}^n \text{CHARGE}(p_i).$$

*Proof.*

$$\begin{aligned} \sum_{i=1}^n \text{COST}(p_i) &= \sum_{i=1}^n (\text{CHARGE}(p_i) - \text{CREDIT}(p_i)) && \text{(by the definition of credit)} \\ &\leq \sum_{i=1}^n \text{CHARGE}(p_i) - \sum_{i=1}^n \Delta \Phi_i && \text{(by the credit-potential invariant)} \\ &= \sum_{i=1}^n \text{CHARGE}(p_i) - (\Phi(D_n) - \Phi(D_0)) && \text{(by telescoping)} \\ &\leq \sum_{i=1}^n \text{CHARGE}(p_i) && \text{(since } \Phi(D_n) \geq 0, \Phi(D_0) = 0). \end{aligned}$$

**Q.E.D.**

When invariant (5) is a strict inequality, it means that some credit is discarded and the analysis of the theorem is not tight. For our “counter” example, the invariant is tight in every case! This means that our preceding derivation is an equality at each step, except possibly for the last step. In fact, it shows that

$$(\text{Total Cost}) = (\text{Total Charge}) - (\Phi(D_n) - \Phi(D_0)).$$

This equation allows an initial potential  $\Phi(D_0)$  that is not necessarily equal to 0 (this is like having a pre-paid credit card). We conclude:

**Lemma 2** *The exact cost of incrementing a counter with initial value  $m$  to final value  $n$  is exactly equal to*

$$\sum_{i=m+1}^n c_i = 2(n - m) - (\Phi_n - \Phi_m)$$

where  $\Phi_n$  is the number of 1's in the binary representation of  $n$ .

E.g., the cost to count from 0 to 25 (i.e.,  $C = (011001)$ ) is exactly  $2(25) - \Phi(C) = 50 - 3 = 47$ . Hope you find this result quite remarkable.

The distinction between “charge” and “amortized cost” should be clearly understood: the former is a definition and the latter is an assertion. The assertion is justified by Theorem 1. A charge can only be called an amortized cost if the overall scheme satisfies the credit-potential invariant.

*Get this!*

**§4. Another Amortization Scheme for Counters** The potential function can be generalized in several ways: it need not be defined just for the data structure  $D$ , but could be defined for one or more abstract features of  $D$ . Thus, we might have a potential function  $\Phi_j$  for the  $j$ th feature ( $j = 1, 2, \dots$ ). The charge for an operation could be split up among these potential functions  $\Phi_j$ .

To illustrate this, we give an alternative amortized cost of incrementing binary counters. Let us set up a **charge account** at each bit position of the binary counter: let  $A_i$  be the account at the  $i$ th smallest position ( $i = 0$  corresponds to the least significant bit,  $i = 1$  to the next significant bit, etc). Each unit of work changes the value of a particular bit of the counter; if the  $i$ th bit is changed, we charge one unit to the account  $A_i$ . For example, over a sequence of  $n$  increments, the account  $A_0$  is charged  $n$  times, the account  $A_1$  is charged  $\leq n/2$  times, and in general,  $A_i$  is charged  $\leq n/2^i$  times. Hence the total charge is at most

$$n(1 + \frac{1}{2} + \frac{1}{4} + \cdots) \leq 2n.$$

Hence the amortized cost per increment is  $\leq 2$ .

Note that this charging scheme is slightly simpler than the potential method, since we charge each operation the *actual* cost of the operation! In other words, the credit of each operation is 0. Nevertheless, it does not lead to the the exact cost (Lemma 2) of the potential method. We will return to these ideas in a later chapter on the Union Find data structure.

### EXERCISES

**Exercise 1.1:** Our model and analysis of counters can yield the exact cost to increment from any initial counter value to any final counter value. Show that the number of work units to increment a counter from 68 to 125 is (exactly!) 110.

◇

**Exercise 1.2:** Recall that our cost model for incrementing a binary counter.

- (a) What is the *exact number* of work units to count from 0 to 99?
- (b) What is the *exact expression in  $n$*  for the work units to count from 0 to  $n$ ?
- (c) What is the *exact expression in  $m$  and  $n$*  for the work units to count from  $m$  to  $n$  ( $m \leq n$ )?
- (d) Use your formula in (c) to determine the exact cost to increment from 100 to 200.

◇

**Exercise 1.3:** Let us now change the cost model in the binary counter problem. Assume the cost of flipping a bit depend on the value of the bit, given as follows:

$$\text{Cost}(0 \rightarrow 1) = 2, \quad \text{Cost}(1 \rightarrow 0) = 1.$$

For instance, to increment the counter value  $C = (011011)$  to  $C = (011100)$ , it costs  $\text{Cost}(011 \rightarrow 100) = 2 + 1 + 1 = 4$  units.

- (a) Give an amortized cost upper bound for incrementing such a binary counter.
- (b) What is the cost to increment from 17 to 65? How tight is your bound?

◇

**Exercise 1.4:** We now generalize the previous question on alternative cost models for incrementing counters. Suppose the cost to flip a bit is given by

$$\text{Cost}(0 \rightarrow 1) = a, \quad \text{Cost}(1 \rightarrow 0) = b.$$

where  $a > 0$  and  $b > 0$  are constants. E.g., the cost to increment the counter value  $C = (011011)$  to  $C = (011100)$  is  $\text{Cost}(011 \rightarrow 100) = a + 2b$  units.

- (a) Give a amortized cost upper bound for incrementing such a binary counter. Try to make your bound tight.

**HINT:** A tight amortization scheme would reproduce our known results where  $a = b = 1$ . Note that you need to define a potential function and a charging scheme.

- (b) What is the exact cost to increment from 17 to  $n = 64$ ?

◇

**Exercise 1.5:** Give an amortized analysis for incrementing a 3-ary counter. A 3-ary counter is a sequence of **trits** (like bits) where each trit is 0, 1 or 2. E.g.,

$$(0) \rightarrow (01) \rightarrow (02) \rightarrow (010) \rightarrow (011) \rightarrow (012) \rightarrow (020) \rightarrow \dots$$

The cost model is that it takes unit work to change the value of one trit. Give an amortized cost is  $O(1)$  per operation. *For full credit, make the the analysis tight.* ◇

**Exercise 1.6:** In a ternary counter, we have ‘trits’ (0, 1, 2) instead of ‘bits’. Changing each trit costs 1 unit. E.g., incrementing the counter (021012) gives (021020) and this costs 2 units. What is the exact cost when I increment my ternary counter from 0 to 83?

◇

**Exercise 1.7:** Generalize the previous question to give an amortized analysis for incrementing a  $k$ -ary counter ( $k \geq 2$ ). The cost model is this: changing a  $k$ -ary digit costs 1 unit. Call this the **uniform cost model**.

(a) Give an amortized analysis for incrementing a 3-ary counter. Note that the 3-ary counter is a sequence of “trits” where each trit is a 0, 1 or 2. The cost model is that it takes unit work to change the value of one trit. Give an amortized cost is  $O(1)$  per operation. *For full credit, make the the analysis tight.*

(b) Generalize part(a) to the incrementing of a  $k$ -ary counter for any  $k \geq 2$ . Your amortized cost should take  $k$  into account. ◇

**Exercise 1.8: Exercise 1.9:** In the following, we assume the uniform cost model for binary and ternary counters.

(a) What is the exact cost to count from 130 to 1030 using a binary counter.

(b) Solve part(a) again but using a ternary counter this time.

◇

◇

**Exercise 1.10:** Let us change the cost model for the ternary counter: Let  $C(i \rightarrow j)$  denote the cost of changing a trit value from  $i$  to  $j$ . Our new cost model is:

$$C(0 \rightarrow 1) = C(1 \rightarrow 2) = 1, \quad C(2 \rightarrow 0) = 2.$$

What is the amortized cost for counting in this model?

◇

**Exercise 1.11:** A simple example of amortized analysis is the cost of operating a special kind of pushdown stack. Our stack  $S$  supports the following two operations:  $S.\text{push}(K)$  simply add the key  $K$  to the top of the current stack. But  $S.\text{pop}(K)$  will keep popping the stack



as long as the current top of stack has a key smaller than  $K$  (the bottom of the stack is assumed to have the key value  $\infty$ ). Let  $a, b$  be positive numbers. The cost for push operation is  $a$  and the cost for popping  $m \geq 0$  items is  $(m + 1)b$ . E.g., starting with an empty  $S$ , if we execute

$S.\text{push}(4), S.\text{push}(1), S.\text{push}(2), \text{pop}(3),$

then successively,  $S$  will be  $(4), (4, 1), (4, 1, 2), (4)$ . The corresponding costs are  $a, a, a, 3b$

- (a) Use our potential framework to give an amortized analysis for a sequence of such push/pop operations, starting from an initially empty stack.
- (b) Give an expression for the exact cost of a sequence of operations in terms of  $a, b, u, o, p$  where  $u$  is the number of pushes,  $o$  is the number of pops and  $p$  is the number of items left in the stack at the end. Assume you begin with an empty stack.

REMARK: Such a stack is used in implementing Graham's algorithm for the convex hull of a set of planar points (see Section 5 on convex hull in this chapter).  $\diamond$

**Exercise 1.12:** Let us generalize the example of incrementing binary counters. Suppose we have a collection of binary counters, all initialized to 0. We want to perform a sequence of operations, each of the type

$\text{inc}(C), \quad \text{double}(C), \quad \text{add}(C, C')$

where  $C, C'$  are names of counters. The operation  $\text{inc}(C)$  increments the counter  $C$  by 1;  $\text{double}(C)$  doubles the counter  $C$ ; finally,  $\text{add}(C, C')$  adds the contents of  $C'$  to  $C$  while simultaneously set the counter  $C'$  to zero. Show that this problem has amortized constant cost per operation.

We must define the cost model. The length of a counter is the number of bits used to store its value. The cost to double a counter  $C$  is just 1 (you only need to prepend a single bit to  $C$ ). The cost of  $\text{add}(C, C')$  is the number of bits that the standard algorithm needs to look at (and possibly modify) when adding  $C$  and  $C'$ . E.g., if  $C = 11, 1001, 1101$  and  $C' = 110$ , then  $C + C' = 11, 1010, 0011$  and the cost is 9. This is because the algorithm only has to look at 6 bits of  $C$  and 3 bits of  $C'$ . Note that the 4 high-order bits of  $C$  are not looked at: think of them as simply being “linked” to the output. Here is where the linked list representation of counters is exploited. After this operation,  $C$  has the value 11, 1010, 0011 and  $C'$  has the value 0.

HINT: The potential of a counter  $C$  should take into account the number of 1's as well as the bit-length of the counter.

*You couldn't do this with arrays!*

**Exercise 1.13:** In the previous counter problem, we define a cost model for  $\text{add}(C, C')$  that depends only on the bit patterns in  $C$  and  $C'$ . In particular, the cost of  $\text{add}(C, C')$  and  $\text{add}(C', C)$  are the same. How can you implement the addition algorithm so that the cost model is justified? HINT: recall that counters are linked lists, and you must describe your algorithm in terms of list manipulation.  $\diamond$

**Exercise 1.14:** Generalize the previous exercise by assuming that the counters need not be initially zero, but may contain powers of 2.  $\diamond$

**Exercise 1.15:** Joe Smart reasons that if we can increment counters for an amortized cost of  $\mathcal{O}(1)$ , we should be able to also support the operation of “decrementing a counter”, in



addition to those in the previous exercise. This should have an amortized cost of  $\mathcal{O}(1)$ , of course.

(a) Can you please give Joe a convincing argument as to why he is wrong?

(b) Joe's intuition about the symmetry of decrement and increment is correct if we change our complexity model. Vindicate Joe by showing a model where we can increment, decrement, add and double in  $\mathcal{O}(1)$  per operation. HINT: we allow our counters to store negative numbers. We also need a more general representation of counters.

(c) In your solution to (b), let us add another operation, testing if a counter value is 0. What is the amortized complexity of this operation?  $\diamond$

**Exercise 1.16:** Suppose we want to generate a lexicographic listing of all  $n$ -permutations (See Chapter 5 on generating permutations). Give an amortized analysis of this process.  $\diamond$

END EXERCISES

## §2. Splay Trees

The **splay tree data structure** of Sleator and Tarjan [12] is a practical approach to implementing all ADT operations listed in §III.2. Splay trees are just ordinary binary search trees – there are no structural requirements (cf. AVL trees) on splay tree. For instance, a splay binary tree could be a tree consisting of a single path (effectively a list), as shown in the leftmost tree in Figure 1. What distinguishes them as “splay trees” are the special splay algorithms used to implement standard binary search tree operations such as lookup, insert and delete. These algorithms invariably call a special **splay** operation. Like rotations in Chapter III, its role is purely to restructure a binary tree.

**§5. Splaying.** The splay operation, applied to an arbitrary node of the tree, will bring this node to the root position. Splaying may be traced to an idea called the **move-to-front heuristic**: suppose we want to repeatedly access various items in a list, and the cost of accessing the item is proportional to its distance from the front of the list. The heuristic says *it is a good idea to move an accessed item to the front of the list*. Intuitively, this move will facilitate future accesses to this item. Of course, there is no guarantee that we would want to access this item again in the future. But even if we never again access this item, we have not lost much by moving the item to the front: the cost of this (preemptive) move has already been paid for (by “charging” its cost to the cost of looking up the item). Thus, this move-to-front heuristic can be justified by amortization arguments; alternatively, one can also use probabilistic analysis (see §10). The analogue of the move-to-front heuristic in maintaining binary search trees is this: after accessing a key  $K$  in a tree  $T$ , we must move the node containing  $K$  to the root. This can be done by rotations, of course. What if we looked up  $K$ , and it is not in  $T$ ? It is still essential to perform this heuristic; otherwise, it is impossible (why?) to achieve a sublinear-time amortized bound. In this case, we move the successor or predecessor of  $K$  to the root. Recall that the **successor** of  $K$  in  $T$  is the smallest key  $K'$  in  $T$  such that  $K \leq K'$ ; the **predecessor** of  $K$  in  $T$  is the largest key  $K'$  in  $T$  such that  $K' \leq K$ . Thus  $K$  does not have a successor (resp., predecessor) in  $T$  if  $K$  is larger (resp., smaller) than any key in  $T$ . Also, the successor and predecessor<sup>3</sup> coincide with  $K$  iff  $K$  is in  $T$ . We characterize the **splay** operation as follows. Let the tree  $T'$  be the result of splaying  $T$  at a key  $K$ . We can invoke the splay operation in

<sup>3</sup>When  $K = K'$ , the successor/predecessor is usually called *improper* successor/predecessor.

one of two<sup>4</sup> forms:

Self-modifying form:  $T.\text{splay}(\text{Key } K) \text{ or } T.\text{splay}(\text{Node } u) \quad \triangleleft T \text{ modifies itself}$  (6)

Functional form:  $T' \leftarrow \text{splay}(\text{Key } K, \text{Tree } T) \quad \triangleleft T \text{ is unchanged}$  (7)

The self-modifying form (6) arise from object-oriented programming languages: we treat  $T$  as an instance of the splay tree class, and **splay** is a method (function) of the class. The functional form (7) keeps the original tree  $T$  and returns its splayed version  $T'$ . We normally implement the self-modifying form, but the functional form is useful for expositional purposes as it allows us to talk about  $T$  and  $T'$  simultaneously. Thus:

1.  $T$  and  $T'$  are equivalent binary search trees.
2.  $T'$  has the property that the key (say  $K'$ ) at the root of  $T'$  is *either* the successor or predecessor of  $K$  in  $T$ .

We are indifferent as to whether  $K'$  is the successor or predecessor of  $K$ . In particular, if  $K$  is smaller than any key in  $T$ , then  $K'$  is the smallest key in  $T$ . But if  $K$  is larger than any key in  $T$  then  $K'$  would be the largest key in  $T$ . If  $T$  is non-empty, then any key  $K$  must have at least a successor or predecessor in  $T$ . For example, starting from the BST in Figure 1(a), splaying the key  $K = 4$  yields the BST in Figure 1(c).

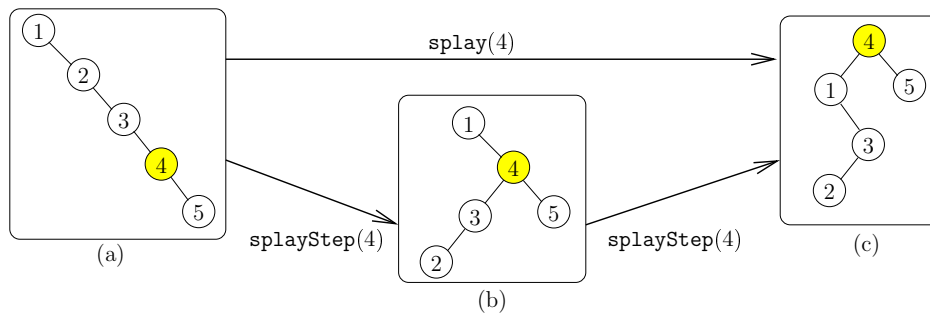


Figure 1: Splaying key 4 (with intermediate step)

So far, we have described  $\text{splay}(K, T)$  where  $K$  is a key. We describe a variant of **splay** where  $K$  is replaced by a node  $u$  in  $T$ :  $\text{splay}(u, T)$ . The operation of  $\text{splay}(u, T)$  is to bring the node  $u$  to the root of  $T$  by repeated rotations; it returns the new tree  $T'$ . The splay with key variant can be reduced to splay by node variant as follows:

```
splay(K,T):
  u ← lookUp(K,T)
  Return splay(u,T)
```

Here,  $\text{lookUp}(K, T)$  is the traditional Lookup of Chapter III: it returns a node  $u$  in which  $u.\text{key}$  is the successor or predecessor of  $K$  in  $T$ . Below we will have a splay variant of Lookup that behaves differently.

<sup>4</sup>These two forms are analogous to the Destructive versus Conservative graph algorithms discussed in §III.21.

¶6. **Correct Implementation of Splaying.** Let  $u$  be a node in a splay tree  $T$ . The operation  $\text{splay}(u, T)$  amounts to repeated application of a single operation called  $\text{splayStep}(u)$ :

```
splay(u, T)
  While ( $u \neq \text{root}$ )
    splayStep(u)
```

Termination is guaranteed because  $\text{splayStep}(u)$  always reduce the depth of any non-root  $u$ . There is an obvious interpretation of  $\text{splayStep}(u)$ : it is simply “ $\text{rotate}(u)$ ”. Call this “naive splaying”. In the next section, we shall show what this version is inadequate for our amortization goal.

So what is the correct implementation of  $\text{splayStep}$ ? Recall in Chapter III that a node  $u$  is called an **outer grandchild** if it has a grandparent  $v$  such that either  $u = v.\text{left}.\text{left}$  or  $u = v.\text{right}.\text{right}$ . We call  $u$  an **inner grandchild** if it has a grandparent but it is not an outer grandchild.

```
splayStep(Node u):
  There are three cases.
    Base Case. If  $u$  has no grandparent,
      then  $\text{rotate}(u)$  (see Figure 11).
    Case I. Else, if  $u$  is an outer grandchild,
      do  $\text{rotate}(u.\text{parent})$ , followed by  $\text{rotate}(u)$ . See Figure 2.
    Case II. Else,  $u$  is an inner grandchild and
      we do a double rotation ( $\text{rotate}(u)$  twice). See Figure 2.
```

Note that the Base Case occurs when  $u$  or  $u.\text{parent}$  is the root. If  $u$  is the root,  $\text{rotate}(u)$  is a no-op; otherwise,  $\text{rotate}(u)$  brings  $u$  to the root. In Figure 1, we see two applications of  $\text{splayStep}(4)$ . Sleator and Tarjan calls the three cases of  $\text{splayStep}$  the zig (base case), zig-zig (case I) and zig-zag (case II) cases. It is easy to see that the depth of  $u$  decreases by 1 in a zig, and decreases by 2 otherwise. Hence, if the depth of  $u$  is  $h$ , the splay operation will halt in  $\lceil h/2 \rceil$   $\text{splayStep}$ ’s. Recall in §III.6, we call the zig-zag a “double rotation”.

We illustrate the fact that  $\text{splay}(K, T)$  may return the successor or predecessor: let  $T_0$  be the splay tree in Figure 3. If we call  $\text{splay}(6, T_0)$ , the result will be  $T_1$  in the figure; so the new root has key 7 that is a successor of  $K$ . But if we call  $\text{splay}(6, T_1)$ , the result will be the tree  $T_2$  in the figure; so the new root has key 5 that is a predecessor of  $K$ . What if you call  $\text{splay}(6, T_2)$ ?

¶7. **Reduction of ADT operations to Splaying.** Recall the operations of the fully mergeable dictionary ADT (§III.2):

Lookup, Insert, Delete, DeleteMin, Merge, Split. (8)

We already know how to implement all these operations using AVL trees (say) with worst case cost of  $O(\log n)$ . But we now propose to implement them using a new class of Splay Algorithms.





Figure 2: SplayStep at  $u$ : Cases I and II.

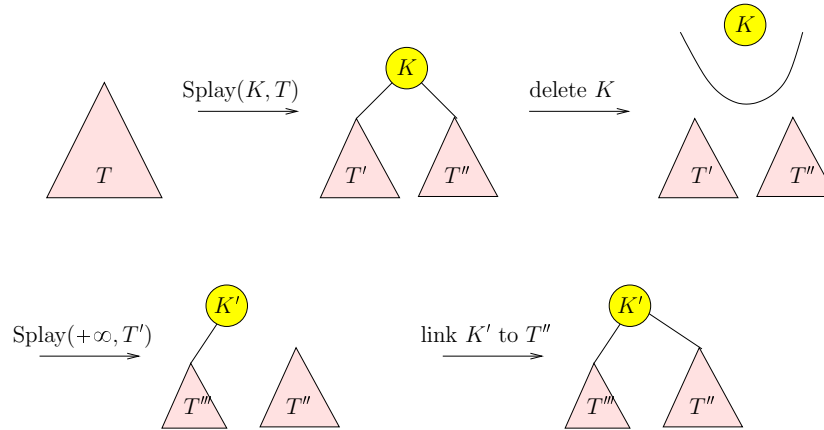
Figure 3: Splaying may return successor or predecessor

The new algorithms are quite simple: each ADT operation is reducible to one or two splaying operations plus some easy  $O(1)$  operations. Here are the self-modifying versions of these operations. We will illustrate them using the BST in Figure 4.

- $T.\text{lookup}(\text{Key } K)$ : first perform  $T.\text{splay}(K)$ . Then examine the root of the resulting tree to see if  $K$  is at the root. **lookup** is a success iff  $K$  is at the root. Note that the tree  $T$  has been modified by splaying, quite unlike the **lookup** algorithm in Chapter III. See Figure 5 for an example of (a) successful lookup and (b) unsuccessful lookup on the tree in Figure 4.
- $T.\text{insert}(\text{Item } X)$ : perform the standard binary search tree insertion of  $X$  into  $T$ . This insertion would fail if  $X.\text{key}$  is already in  $T$ . Regardless of success or failure, the attempted insertion yields a node  $u$  whose key is equal to  $X.\text{key}$  (this key may either be newly inserted or was already there). Now call  $T.\text{splay}(u)$ .
- $T_1.\text{merge}(\text{Tree } T_2)$ : The precondition for this operation is that each key in  $T_1$  must be less than any key in  $T_2$ . First do  $T_1.\text{splay}(+\infty)$ . Here  $+\infty$  is an artificial key larger than

Figure 4:  $T_{15}$ : BST Example for Splay AlgorithmsFigure 5: (a)  $\text{lookUp}(13)$ , (b)  $\text{lookUp}(8.5)$  on  $T_{15}$ 

any real key in  $T_1$ . As a result, the root of  $T_1$  has no right child. We then make  $T_2$  the right subtree of  $T_1$ . The merge operation is illustrated in the bottom row of Figure 6, since it is called by the delete algorithm.

Figure 6: Steps in  $\text{delete}(K, T)$ , including  $T'.\text{merge}(T'')$ 

- $T.\text{delete}(\text{Key } K)$ : first perform  $T.\text{splay}(K)$ . If the root of the resulting tree does not contain  $K$ , there is nothing to delete. Otherwise, delete the root and merge the left  $T_L$  and right  $T_R$  subtrees, as just described above:  $T_L.\text{merge}(T_R)$ . This is illustrated in Figure 6.
- $X \leftarrow T.\text{deleteMin}()$ : we perform  $T.\text{splay}(-\infty)$ . Now the minimum item  $X$  is at the root of  $T$ , and it has no left subtree. We update  $T$  to be the right subtree of  $X$ , and return  $X$ .

- $T' \leftarrow T.\text{Split}(\text{Key } K)$ : By definition, the new tree  $T'$  will contain all keys of  $T$  that are  $> K$ . In particular, if  $K$  occurs in  $T$ , it will remain in  $T$ . We perform  $T.\text{splay}(K)$  so that the root of  $T$  now contains the successor or predecessor of  $K$  in  $T$ . Split off the right subtree of  $T$  as the new tree  $T'$ . But we must check the key  $K'$  at the root of  $T$ : if  $K' > K$ , then we must give this root to  $T'$ , and otherwise, it remains the root of  $T$ .

Although we will regard the above as our official splay algorithms, it is instructive to see some other variants:

- Alternative  $T.\text{insert}(\text{Item } X)$ : we could first  $T.\text{splay}(X.\text{key})$ . Let  $K$  be the key at the root after this splay. If  $K = X.\text{key}$ , then we cannot insert  $X$ . Otherwise, insert  $X$  as the right or left-child of the root, depending on whether  $K < X.\text{key}$  or  $K > X.\text{key}$ .
- Alternative  $T.\text{delete}(\text{Key } K)$ : we could first do the standard deletion algorithm. This deletion identifies a node  $u$  whose child has been “cut” (removed). We then do  $\text{splay}(u, T)$ .

It is interesting to compare the above ADT implementations with our treatment of Binary Search Trees in Chapter 3: there, we reduce all ADT operations to a single operation, rotation. Now, all operations are reduced to splaying. We shall see that splaying itself is reduced to rotations.

**§8. Inadequacy of Naive Splaying.** Recall that naive splaying refers to implementing  $\text{splayStep}(u)$  as a call  $\text{rotate}(u)$ . Naive splaying does ensure the correctness of our splay algorithms in §7. But it is insufficient to achieve our amortization goal, namely *each of the operations in §7 has an amortized cost of  $O(\log n)$* .




Figure 7: Naive splaying on  $T_5$

Consider the binary tree  $T_n$  that contains the keys  $1, 2, \dots, n$  and whose shape is just a right-path. The case  $T_5$  is illustrated in Figure 7(a). It is easy to verify that we can get  $T_n$  by inserting the keys in the following order:

$$n, n-1, n-2, \dots, 2, 1.$$

The insertion algorithm follows the prescription in §7: first do standard insertion and then splay the inserted node.

Next, let us perform a sequence of Lookups on  $T_n$ : first do  $\text{lookup}(n, T_n)$ . Using our splay-based lookup algorithm, we first  $\text{splay}(n, T_n)$ , which produces a tree rooted at  $n$  with a left child that is  $T_{n-1}$ . This is illustrated in Figure 7(b). We can repeat this process by performing lookups on  $n-1, n-2, \dots, 1$ . As seen in Figure 7(f), the final result is again  $T_n$ . The cost for this sequence of operations is  $\Theta$ -order of  $\sum_{i=2}^n i = \Theta(n^2)$ . Therefore it is impossible to achieve an amortized cost of  $O(\log n)$  on each operation in this Lookup sequence.

¶9. **Top Down Splaying.** The above splay algorithms require two passes over the splay path (once to find the node for splaying, and then once more to splay to the root). The following explores a one-pass algorithm denoted `topSplay(Key  $K$ , Node  $u$ )`. Instead of going down the LookUp path to visit nodes, we want to bring these nodes up to the root. This is called “top-down splaying”.

Before giving the correct solution, it is instructive to see a false start. Suppose we are looking for the key  $K$  and  $u$  is the root. If  $u.Key = K$ , we are done. If not, we must next visit the left or right child ( $u_L$  or  $u_R$ ) of  $u$ . But instead of going down the tree, we simply rotate  $u_L$  or  $u_R$  to the root! For instance, if  $u.Key > K$ , then we must search for  $K$  in the subtree rooted at  $u_L$ , and so we do the rotation `rotate( $u_L$ )` to bring  $u_L$  to the root. So our initial idea amounts to a repeated left- or right-rotation at the root. Unfortunately, this simple approach has a pitfall. To fix this, let us store some state information: we have 4 possible states in our algorithm. Again,  $u$  will be the root node and  $u_L, u_R$  denote the left and right child of  $u$  (these children may be null).

*Find the pitfall!*

- State 0: Both  $u_L$  and  $u_R$  have not been visited.
- State 1:  $u_L$ , but not  $u_R$ , has been visited. Moreover,  $u_L.key < K$ .
- State 2:  $u_R$ , but not  $u_L$ , has been visited. Moreover,  $u_R.key > K$ .
- State 3: Both  $u_L$  and  $u_R$  have been visited. Moreover,  $u_L.key < K < u_R.key$ .

We have a global variable **State** that is initialized to 0. Here are the possible state transitions. From state 0, we must enter either state 1 or state 2. From states 1 or 2, we can either remain in the same state or enter state 3. Once state 3 is entered, we remain in state 3. Unfortunately this `topSplay` algorithm does not have amortized logarithmic behavior (Exercise).




Figure 8: Top Splay: Cases 0, I and II

Our solution is to maintain 3 splay trees,  $L, C, R$ , corresponding to the Left-, Center- and Right-trees. Refer to Figure 8(a). Initially Left- and Right-trees are empty, and the Center-tree is the input tree. We assume the keys in these trees satisfy  $L < C < R$ . Inductively, assume the key  $K$  that we are looking for is in  $C$ . There are three cases: suppose  $u$  is the root.



- CASE (0): This is the base case. The key  $K$  we are looking for is equal to the key at  $u$ . We attach the left and right children of  $u$  at the rightmost tip and leftmost tip of  $L$  and  $R$ , resp. We then make  $u$  the root of a tree with left and right children  $L$  and  $R$ . See Figure 8(b). In case  $K$  is equal to the key at a child of  $u$ , we just rotate that child before applying the preceding transformation.
- CASE (I): Suppose the key  $K$  is found in the subtree rooted at an outer grandchild  $w$  of  $u$ . By symmetry, assume that  $w$  is the left child of  $v$ , where  $v$  is the left child of  $u$ . In this case, we transfer the nodes  $u, v$  and the right subtrees of  $u$  and  $v$  to the left-tip of  $R$ , as shown in Figure 8(c).
- CASE (II): Suppose key  $K$  is found in the subtree rooted at an inner grandchild  $w'$  of  $u$ . By symmetry, assume  $w'$  is right child of  $v$ , where  $v$  is the left child of  $u$ . In this case, we transfer the node  $u$  and its right subtree to the left-tip of  $R$ , and the node  $v$  and its left-subtree to the right-tip of  $L$ , as shown in Figure 8(d).

The correctness of this procedure is left as an easy exercise.

---

## EXERCISES

**Exercise 2.1:** (a) Describe the shape of the splay tree  $T_n$  that results from the following sequence of operations, starting from an initially empty tree:

`insert(1, 2, ..., n),    lookUp(1)`

Note that `insert(1, ..., n)` represents a sequence of  $n$  insertions starting with `insert(1)`.

HINT: First show us  $T_8$  and  $T_9$ . Then tell us the general shape for an arbitrary  $n$ .

- (b) Consider  $T_n$  from part(a). Please show the result of  $T_n.\text{Insert}(1.5)$  for two cases:  $n = 8$  and  $n = 9$ .
- (c) Consider  $T_n$  from part(a). Please show the result of  $T_n.\text{Delete}(2)$  for two cases:  $n = 8$  and  $n = 9$ .

◇

**Exercise 2.2:** Comment on the following assertions (open ended):

- (a) Every splay tree is a binary search tree.
- (b) Every binary search tree is a splay tree.

◇

**Exercise 2.3:** Let  $T$  be a binary tree with nodes  $u_1, \dots, u_5$  where  $u_1$  is root,  $u_{i+1}$  is the child of  $u_i$ . However,  $u_{i+1}$  is a left-child if  $i$  is odd, and a right-child if  $i$  is even.

- (a) Attach the five keys 1, 2, 3, 4, 5 to the nodes in  $T$  so that the result is a splay tree. Draw the resulting tree, denoted  $T_a$ .
- (b) Splay 3 in  $T_a$  and show the resulting tree, denoted  $T_b$ .
- (c) Insert 6 in  $T_a$  and show the resulting  $T_c$ .
- (d) Insert 3.5 in  $T_a$  and show the resulting  $T_d$ .
- (e) Delete 3 in  $T_a$  and show the resulting  $T_e$ .
- (f) Delete 2 in  $T_a$  and show the resulting  $T_f$ .

◇

**Exercise 2.4:** In this question, we ask which BSTs can arise from splay operations. Define an **IL splay tree** to be one that can be obtained by a sequence of insertions (I) and lookups

(L) into an initially empty splay tree. Note that we omit deletes (D). If we include deletes, we have an **ILD splay tree**. Remember that it is the **shape** of a BST that matters, and there is no need to indicate explicit keys in your BSTs. Furthermore, *assume that insertion and deletion is implemented by first doing the standard BST insertion/deletion, followed by a splay*. This assumption will facilitate our enumeration of ILD shapes: For instance, given any shape of size  $n$ , it is easy to see all the shapes of size  $n + 1$  that can result from an insertion: just add a leaf, and splay from that leaf. There are  $n + 1$  ways to insert this leaf.

(a) Show that not every BST on 3 nodes is an IL splay tree.

(b) Show that every BST on 3 nodes is an ILD splay tree.

(c) Show that any BST on 4 nodes is an IL splay tree.

(d) Show that any BST on 5 nodes is an IL splay tree.

HINT: Furthermore, it is sufficient to enumerate shapes up to left-right symmetry (e.g., for BSTs with 3 nodes, there are only distinct 3 shapes, up to left-right symmetry). Moreover, when we insert into a given

(e) Prove or disprove: every BST on  $n > 5$  nodes is an IL splay tree.  $\diamond$

**Exercise 2.5:** In the text, we splay the sequence of keys  $n, n - 1, n - 2, \dots, 1$  in the tree  $T_n$  (Figure 7) to show that for naive splaying is not amortized  $O(\log n)$ . See  $T_5$  in Figure 7(a). Please re-do this example using regular splaying.

(a) First illustrates what happens for  $n = 5$  and  $n = 6$ .

(b) What is the final result after splaying on the sequence  $n, n - 1, \dots, 2, 1$  on  $T_n$ ? Try to prove this rigorously (using EGVs!).

(c) What is the total number of rotations (count each splayStep as two rotations)?  $\diamond$

**Exercise 2.6:** Perform the following splay tree operations, starting from an initially empty tree.

`insert(3, 2, 1; 6, 5, 4; 9, 8, 7), lookUp(3), delete(7), insert(12, 15, 14, 13), Split(8).`

Show the splay tree after each operation (do include intermediate splaySteps).  $\diamond$

**Exercise 2.7:** Show the result of `merge( $T_1, T_2$ )` where  $T_1, T_2$  are the splay trees shown in Figure 9.  $\diamond$




Figure 9: Splay trees  $T_1, T_2$

**Exercise 2.8:** Consider the insertion of the following sequence of keys into an initially empty tree:  $1, -1, 2, -2, 3, -3, \dots, n, -n$ . Let  $T_k$  be the splay tree after inserting  $k$  (for  $k = 1, -1, 2, -2$ , etc).

(i) Show  $T_n$  and  $T_{-n}$  for  $n = 1, 2, 3$ .

(ii) State and prove a conjecture about the shape of  $T_n$ .

◇

**Exercise 2.9:** Consider the insertion of the following sequence of keys into an initially empty tree:  $3, 2, 1; 6, 5, 4; 9, 8, 7; \dots; 3n, 3n-1, 3n-2$ , where we have written the keys in groups of three to indicate the pattern. Let  $T_n$  be the splay tree after the  $n$ th insertion (so  $T_0$  is the empty tree).

(a) Show  $T_{3n}$  for  $n = 1, 2, 3, 4$ .

(b) State and prove a conjecture about the shape of  $T_n$ .

(a', b') Repeat (a) and (b), but use the following sequence of input keys

$$2, 3, 1; 5, 6, 4; 8, 9, 7; \dots; 3n-1, 3n, 3n-2.$$

◇

**Exercise 2.10:** Consider the insertion of the following sequence of keys into an initially empty tree:  $2, 3, 1; 5, 6, 4; 8, 9, 7; \dots; 3n-1, 3n, 3n-2$ , where we have written the keys in groups of three to indicate the pattern. Let  $T_n$  be the splay tree after the  $n$ th insertion (so  $T_0$  is the empty tree).

(a) Show  $T_{3n}$  for  $n = 1, 2, 3, 4$ .

(b) State and prove a conjecture about the shape of  $T_n$ .

◇

**Exercise 2.11:** (Open ended) Prove that if we have any “regular” sequence of insertions, as in the previous exercises, the result is a “regular” splay tree. Part of this problem is to capture various notions of regularity. Let us capture the previous two exercises: let  $G(n) = (g_1(n), g_2(n), \dots, g_k(n))$  where  $k$  is a constant and  $g_i(n)$  is a integer polynomial in  $n$ . E.g.  $G(n) = (n, -n)$  and  $G(n) = (3n, 3n-1, 3n-2)$  captures the regularity of the previous two exercises. Assume that the sequences  $G(n)$  and  $G(m)$  have different members for  $m \neq n$ . Then we want to show that the insertion of the sequence  $G(1); G(2); G(3); \dots; G(n)$  yields a splay tree  $T_n$  that is “regular”, but in what sense?

◇

**Exercise 2.12:** Fix a key  $K$  and a splay tree  $T_0$ . Let  $T_{i+1} \leftarrow \text{splay}(K, T_i)$  for  $i = 0, 1, \dots$

(a) Under what conditions will the  $T_i$ 's stabilize (become a constant)? How many splays will bring on the stable state?

(b) Under what conditions will the  $T_i$ 's not stabilize? How many splays will bring on this condition?

◇

**Exercise 2.13:** Let  $T$  be a binary search tree in which every non-leaf has one child. Thus  $T$  has a linear structure with a unique leaf.

(a) What is the effect of `lookUp` on the key at the leaf of  $T$ ?

(b) What is the minimum number of `lookUp`'s to make  $T$  balanced?

◇

**Exercise 2.14:** In our operations on splay trees, we usually begin by performing a splay. This was not the case with our insertion algorithm in the text. But consider the following variant insertion algorithm. To insert an item  $X$  into  $T$ :

1. Perform `splay(X.key, T)` to give us an equivalent tree  $T'$ .

2. Now examine the key  $K'$  at root of  $T'$ : if  $K' = X.\text{key}$ , we declare an error (recall that keys must be distinct).

3. If  $K' > X.\text{key}$ , we install a new root containing  $X$ , and  $K'$  becomes the right child

of  $X$ ; the case  $K' < X.\text{key}$  is symmetrical. In either case, the new root has key equal to  $X.\text{key}$ . See Figure 10.

- (a) Prove that the amortized complexity of this insertion algorithm remains  $O(\log n)$ .  
 (b) Compare the amortized complexity of this method with the one in the text. Up to constant factors, there is no difference, of course. But which has a better constant factor?

◇

**Exercise 2.15:** As in the previous Exercise, carry out the details of the variant deletion algorithm noted in the text.

◇




Figure 10: Alternative method to insert a key  $K$ .

**Exercise 2.16:** (Top Down Splaying)

- (a) What is the “pitfall” mentioned for the initial implementation of the top-down splaying algorithm.  
 (b) Show that the amortized complexity of the second attempt cannot be  $O(\log n)$ . To be specific, here is the code:

```

topSplay( $K, u$ )
  Input:  $K$  is a key and  $u$  a node; we are looking for  $K$  in the subtree  $T_u$  rooted at  $u$ 
  Output: We return “found” or “not found”. In any case, as side effect, in place of  $u$ 
          will be a node containing the predecessor or successor of  $K$  in the tree  $T_u$ .
1.  If ( $u.key = K$ ) Return(“found”).
2.  case(State)
      State=0:
        If ( $u.key > K$ )
          rotate( $u.left$ );
          State  $\leftarrow$  2.
        else
          rotate( $u.right$ );
          State  $\leftarrow$  1.
      State=1:
        If ( $u.key > K$ )
           $v \leftarrow u.left.right$ ;
          If ( $v = \text{nil}$ ) Return(“not found”).
          rotate2( $v$ );
          State  $\leftarrow$  3;
        else
           $v \leftarrow u.right$ ;
          If ( $v = \text{nil}$ ) Return(“not found”).
          rotate( $v$ );  $\triangleleft$  Remain in State 1.
      State=2:
        ...  $\triangleleft$  Omitted: symmetrical to State 1.
      State=3:
        If ( $u.key > K$ )
           $v \leftarrow u.left.right$ 
        else
           $v \leftarrow u.right.left$ .
        If ( $v = \text{nil}$ ) Return(“not found”).
        rotate2( $v$ )
         $\triangleleft$  End Case Statement
3.  topSplay( $K, v$ ).  $\triangleleft$  Tail recursion

```

◇


END EXERCISES

### §2.1. Splay Analysis

Our main goal is to prove:

**Theorem 3 (Splay Theorem)** *The amortized cost of each splay operation is  $O(\log n)$  assuming at most  $n$  items in the tree and we started out with an empty tree.*

Before proving this result, let us show that it is a “true” amortization bound. More precisely, we show that the worst case bound is linear, not logarithmic. To see this, consider the repeated



550 insertion of the keys  $1, 2, 3, \dots, n$  (in this order) into an initially empty tree. It is easy to see  
 551 by induction that the result is a left-path  $L_n$  of length  $n - 1$ . Finally, if we perform a lookup  
 552 on key 1 in this tree, we must expend  $\Theta(n)$  units of work.

553 ¶10. **Amortization Scheme for Splay Trees.** Our data structure  $D$  is a collection of  
 554 splay trees, and  $D$  can be modified by various “requests”  $p$  on  $D$ . Following the abstract model  
 555 of ¶3, there are 4 elements:

556 (1) We need a cost model which determines the  $\text{COST}(p, D)$  of each request  $p$  on  $D$ . Looking  
 557 at all the list of ADT operations in ¶7, we see that each operation is reduced to one or  
 558 two splays, plus some  $O(1)$  time work. Hence we will focus on the splay request. Such a  
 559 request can be reduced to a sequence of “splay steps”, plus some  $O(1)$  time work. Thus we  
 560 may designate  $\text{COST}(p, D) = k + 1$  if there are  $k \geq 1$  splay steps.

(2) We must design a potential function  $\Phi(D)$  for  $D$ , viewed as a forest of splay trees. Let  
 $\text{SIZE}(u)$  denote, as usual, the number of nodes in the subtree rooted at  $u$ . Define its  
**potential** to be

$$\Phi(u) = \lfloor \lg \text{SIZE}(u) \rfloor.$$

561 Note that  $\text{SIZE}(u) = 1$  iff  $u$  is a leaf. Thus  $\Phi(u) = 0$  iff  $u$  is a leaf. If  $S = \{u_1, u_2, \dots, u_k\}$   
 562 is a set of nodes, we may write  $\Phi(S)$  or  $\Phi(u_1, u_2, \dots, u_k)$  for the sum  $\sum_{u \in S} \Phi(u)$ . If  $T$  is a  
 563 splay tree  $\Phi(T)$  is defined as  $\Phi(S)$  where  $S$  is the set of nodes in  $T$ . An empty tree  $T$  has  
 564 no potential,  $\Phi(T) = 0$ . Finally,  $\Phi(D)$  is just the sum of  $\Phi(T)$  where  $T \in D$ .

565 E.g., if  $T$  is a left-path with nodes  $S = \{u_1, \dots, u_5\}$  such that  $\text{SIZE}(u_i) = i$  for each  $i$ . Then  
 566  $\Phi(S) = \sum_{i=1}^5 \Phi(i) = 0 + 1 + 1 + 2 + 2 = 6$ .

567 (3) We must specify a  $\text{CHARGE}(p, D)$  on each request  $p$  on  $D$ . If  $p$  is a splay operation, we can  
 568 charge  $O(\log n)$  if the splay tree has size  $n$ .

569 (4) Finally, we must show that the credit-potential invariant in (5) holds for each step. We can  
 570 then invoke Theorem 1 to conclude that  $\text{CHARGE}(p, D)$  can be viewed as the amortized  
 571 cost of  $p$  on  $D$ .

572 **Lemma 4 (Key Lemma)** Let  $\Phi$  be the potential function before we apply  $\text{splayStep}(u)$ , and  
 573 let  $\Phi'$  be the potential after. The credit-potential invariant is preserved if we charge the  $\text{splayStep}$

$$3(\Phi'(u) - \Phi(u)) \tag{9}$$

574 units of work in Cases I and II. In the Base Case, we charge one extra unit, in addition to the  
 575 charge (9).

576 Theorem 3 follows easily from the Key Lemma: suppose that splaying at  $u$  reduces to a  
 577 sequence of  $k$   $\text{splaySteps}$  at  $u$  and let  $\Phi_i(u)$  be the potential of  $u$  after the  $i$ th  $\text{splayStep}$ . Then

$$\begin{aligned} \text{COST}(\text{splay}(u)) &\leq \sum_{i=1}^k \text{CHARGE}(\text{splay}(u)) && \text{(by Theorem 1)} \\ &= 1 + \sum_{i=1}^k 3[\Phi_i(u) - \Phi_{i-1}(u)] && \text{(by Key Lemma)} \\ &= 1 + 3[\Phi_k(u) - \Phi_0(u)] && \text{(by telescoping)} \\ &\leq 1 + 3 \lg n && (\Phi_k(u) \leq \lg n, \ 0 \leq \Phi_0(u)) \end{aligned}$$

578 Note that the “1+” comes from the fact that the last  $\text{splayStep}$  may belong to the base case.  
 579 To finish off the argument, we must account for the cost of looking up  $u$ . But it easy to see that  
 580 this cost is proportional to  $k$  and so it is proportional to the overall cost of splaying. This only  
 581 increases the constant factor in our charging scheme. This concludes the proof of the Splay  
 582 Theorem.

¶11. Proof of Key Lemma. The following is a useful remark about rotations:

**Lemma 5** Let  $\Phi$  be the potential function before a rotation at  $u$  and  $\Phi'$  the potential function after. Then the increase in potential of the overall data structure is at most

$$\Phi'(u) - \Phi(u).$$

The expression  $\Phi'(u) - \Phi(u)$  is always non-negative.




Figure 11: Rotation at  $u$ .

*Proof.* We refer to Figure 11. The increase in potential is

$$\begin{aligned} \Delta\Phi &= \Phi'(u, v) - \Phi(u, v) \\ &= \Phi'(v) - \Phi(u) \quad (\text{as } \Phi'(u) = \Phi(v)) \\ &\leq \Phi'(u) - \Phi(u) \quad (\text{as } \Phi'(u) \geq \Phi'(v)). \end{aligned}$$

It is obvious that  $\Phi'(u) \geq \Phi(u)$ .

**Q.E.D.**

*Proof of Key Lemma.* The Base Case is almost immediate from Lemma 5: the increase in potential is at most  $\Phi'(u) - \Phi(u)$ . This is at most  $3(\Phi'(u) - \Phi(u))$  since  $\Phi'(u) - \Phi(u)$  is non-negative. The charge of  $1 + 3(\Phi'(u) - \Phi(u))$  can therefore pay for the cost of this rotation and any increase in potential.

The remaining two cases illustrated in Figure 2 are reproduced here for the reader's convenience: let the sizes of the subtrees  $A, B, C, D$  be  $a, b, c, d$ , respectively.




Figure 12: Case I of  $\text{SplayStep}(u)$



Consider Case I (see Figure 13): the increase in potential is

$$\begin{aligned}\Delta\Phi &= \Phi'(u, v, w) - \Phi(u, v, w) \\ &= \Phi'(v, w) - \Phi(u, v) \quad (\text{as } \Phi'(u) = \Phi(w)) \\ &\leq 2(\Phi'(u) - \Phi(u)) \quad (\text{as } 2\Phi'(u) \geq \Phi'(v, w), \quad 2\Phi(u) \leq \Phi(u, v)).\end{aligned}$$

Since  $\Phi'(u) \geq \Phi(u)$ , we have two possibilities: (a) If  $\Phi'(u) > \Phi(u)$ , then the charge of  $3(\Phi'(u) - \Phi(u))$  can pay for the increased potential *and* the cost of this splay step. (b) Next suppose  $\Phi'(u) = \Phi(u)$ .

$$\begin{aligned}\Phi'(u) = \lfloor \lg(3 + a + b + c + d) \rfloor &= \Phi(u) = \lfloor \lg(1 + a + b) \rfloor && (\text{by assumption}) \\ 3 + a + b + c + d &< 2(1 + a + b) && (\text{else, } \Phi'(u) > \Phi(u)) \\ 2 + c + d &< 1 + a + b && (\text{subtract } 1 + a + b \text{ from both sides}) \\ 2(2 + c + d) &< 3 + a + b + c + d && (\text{add } 2 + c + d \text{ to both sides}) \\ \Phi'(w) = \lfloor \lg(1 + c + d) \rfloor &< \lfloor \lg(3 + a + b + c + d) \rfloor = \Phi(u).\end{aligned}$$

Also,

$$\Phi'(v) \leq \Phi'(u) = \Phi(u) \leq \Phi(v).$$

Combining the last two inequalities, we conclude that

$$\Phi'(w, v) < \Phi(u, v).$$

Hence  $\Delta\Phi = \Phi'(w, v) - \Phi(u, v) < 0$ . Since potentials are integer-valued, this means that  $\Delta\Phi \leq -1$ . Thus the change in potential releases at least one unit of work to pay for the cost of the splay step. Note that in this case, we charge nothing since  $3(\Phi'(u) - \Phi(u)) = 0$ . Thus the credit-potential invariant holds.




Figure 13: Case II of SplayStep(u)

Consider Case II (see Figure 13): The increase in potential is again  $\Delta\Phi = \Phi'(v, w) - \Phi(u, v)$ . Since  $\Phi'(v) \leq \Phi(v)$  and  $\Phi'(w) \leq \Phi(u)$ , we get

$$\Delta\Phi \leq \Phi'(u) - \Phi(u).$$

If  $\Phi'(u) - \Phi(u) > 0$ , then our charge of  $3(\Phi'(u) - \Phi(u))$  can pay for the increase in potential and the cost of this splay step. Hence we may assume otherwise and let  $t = \Phi'(u) = \Phi(u)$ . In this case, our charge is  $3(\Phi'(u) - \Phi(u)) = 0$ , and for the credit potential invariant to hold, it suffices to show

$$\Delta\Phi < 0.$$

It is easy to see that  $\Phi(v) = t$ , and so  $\Phi(u, v) = 2t$ . Clearly,  $\Phi'(v, w) \leq 2\Phi'(u) = 2t$ . If  $\Phi'(v, w) < 2t$ , then  $\Delta\Phi = \Phi'(v, w) - \Phi(u, v) < 0$  as desired. So it remains to show that  $\Phi'(v, w) = 2t$  is impossible. For, if  $\Phi'(v, w) = 2t$  then  $\Phi'(v) = \Phi'(w) = t$  (since  $\Phi'(v), \Phi'(w)$  are both no larger than  $t$ ). But then

$$\Phi'(u) = \lfloor \lg(\text{SIZE}'(v) + \text{SIZE}'(w) + 1) \rfloor \geq \lfloor \lg(2^t + 2^t + 1) \rfloor \geq t + 1,$$

a contradiction. Here,  $\text{SIZE}'$  denotes the size after the splay step operation. This proves the Key Lemma.

All our splay tree operations are ultimately reduced to splay operation, plus some simple primitive  $O(1)$  time operations. The Key Lemma bounds the potential change in the splay operation. But some of these primitive operations may change the potential of our splay trees via the deletion or insertion of a leaf: our next lemma bounds such potential changes.

**Lemma 6** *Let  $T^+$  be a binary tree with  $n + 1$  items, and  $T$  is obtained from  $T^+$  by deleting some leaf  $x$ . Alternatively,  $T^+$  can be viewed as the result of inserting a new leaf  $x$  into  $T$ . Then  $\Phi(T^+) - \Phi(T) \leq \lg n$ .*

*Proof.* Let  $(u_0, u_1, \dots, u_m)$  denote the path from the root of  $T^+$  to  $x = u_m$ . Let  $n_i$  be the size of the subtree at  $u_i$  in  $T$ , and so  $n > n_0 > n_1 > \dots > n_m = 0$ . Therefore the potential of  $u_i$  in  $T^+$  is  $\lfloor \lg(n_i + 1) \rfloor$ , and the increase in potential in going from  $T$  to  $T^+$  is

$$\Phi(T^+) - \Phi(T) = \sum_{i=0}^{m-1} [\lg(n_i + 1)] - [\lg(n_i)].$$

Observe that  $\lfloor \lg(n_i + 1) \rfloor - \lfloor \lg(n_i) \rfloor$  is either 0 or 1. Moreover, it is 1 iff  $n_i + 1$  is a power of 2. There are at most  $\lg n$  values of  $i$  for which it is equal to 1. Hence

$$\Phi(T^+) - \Phi(T) \leq \lg n.$$

**Q.E.D.**

**¶12. Amortized Cost of Splay Implementation of ADT Operations.** We conclude with an amortized cost statement for splay trees.

**Theorem 7** *Starting from an initially empty splay tree, any sequence of  $m$  requests of the types*

*lookUp, insert, merge, delete, deleteMin, Split,*

*and involving a total of  $n$  items, has total time complexity of  $O(m \log n)$ . Thus, the amortized cost is  $O(\log n)$  per request.*

*Proof.* This follows almost immediately from Theorem 3 since each request can be reduced to a constant number of splay operations plus  $O(1)$  extra work. The splay operations are charged  $O(\lg n)$  units, and the extra work is charged  $O(1)$  units. But two important details must not be overlooked: sometimes, the extra  $O(1)$  work increases the potential by a non-constant amount, and this increase must be properly charged. This situation happens in two situations.

(A) When inserting a new key: the key will become a leaf of the tree, followed by splaying at this leaf. While the splaying cost is accounted for, the act of creating this leaf may also increase the potential of every node along the path to the leaf. By Lemma 6, this increase is at most  $\lg n$ .

(B) When merging two trees  $T_1, T_2$ . In this case, we first perform `splay(+∞,  $T_1$ )`. If  $u$  is the root of the resulting tree, then  $u$  has no right child and we simply attach  $T_2$  as the right subtree of  $u$ . This “attachment” will increase the potential of  $u$  by at most  $1 + \lg(|T_1| + |T_2|) - \lg(|T_1|) = 1 + \lg(1 + |T_2|/|T_1|) < 1 + \lg n$ . Thus we just have to charge this operation an extra  $1 + \lg n$  units. Note that deletion is also reduced to merging, and so its charge must be appropriately increased. In any case, all charges remain  $O(\lg n)$ , as we claimed. **Q.E.D.**

Note that the above argument does not immediately apply to `topSplay`: this is treated in the Exercises. Sherk [11] has generalized splaying to  $k$ -ary search trees. In such trees, each node stores an ordered sequence of  $t - 1$  keys and  $t$  pointers to children where  $2 \leq t \leq k$ . This is similar to  $B$ -trees except we do not require all the leaves to be in the same level.

**¶13. Application: Splaysort** Clearly we can obtain a sorting algorithm by repeated insertions into a splay tree, followed by repeated `deleteMins` — this is analogous to the heapsort algorithm (¶III.7). This algorithm is known as **splaysort**. It is not only theoretically optimal with  $O(n \log n)$  complexity, but has been shown to be quite practical [7]. Splaysort has the ability to take advantage of “presortedness” in the input sequence. For instance, running splaysort on the input sequence  $x_1 > x_2 > \dots > x_n$  will take only  $O(n)$  time in the worst case. One way to quantify presortedness is to count the number of pairwise inversions in the input sequence. E.g., if the input is already in sorted order, we would like to (automatically) achieve an  $O(n)$  running time. However, if  $x_1 < x_2 < \dots < x_n$  is our input, the advantage of presortedness is lost for our version of splaysort. The Exercises discuss some ways to overcome this.

Quicksort (Lectures II and VIII) is regarded as one of the fastest sorting algorithms in practice. But algorithms like splaysort may run faster than Quicksort for “well presorted” inputs. Quicksort, by its very nature, deliberately destroy any property such as presortedness in its input.

---

## EXERCISES

**Exercise 2.17:** Where in the proof is the constant “3” actually needed in our charge of  $3(\Phi'(u) - \Phi(u))$ ?  $\diamond$

**Exercise 2.18:** Adapt the proof of the Key Lemma to justify the following variation of `SplayStep`:

```
VARSPPLAYSTEP( $u$ ):
    (Base Case) If  $u$  is a child or grandchild of the root,
        then rotate once or twice at  $u$  until it becomes the root.
    (General Case) else rotate at  $u.parent$ , followed by two rotations at  $u$ .
```

**Exercise 2.19:** Let us define the potential of node  $u$  to be  $\Phi(u) = \lg(\text{SIZE}(u))$ , instead of  $\Phi(u) = \lfloor \lg(\text{SIZE}(u)) \rfloor$ . In other words, we avoid the floor function to make the potential function continuous.  $\diamond$

(a) How does this modification affect the validity of our Key Lemma about how to charge **splayStep**? In our original proof, we had 2 cases: either  $\Phi'(u) - \Phi(u)$  is 0 or positive. But now,  $\Phi'(u) - \Phi(u)$  is always positive. Thus it appears that we have eliminated one case in the original proof. What is wrong with this suggestion?

(b) Consider Case I in the proof of the Key Lemma. Show that if  $\Phi'(u) - \Phi(u) \leq \lg(6/5)$  then  $\Delta\Phi = \Phi'(w, v) - \Phi(u, v) \leq -\lg(6/5)$ . HINT: the hypothesis implies  $a+b \geq 9+5c+5d$ .

(c) Do the same for Case II.  $\diamond$

**Exercise 2.20:** In the previous question, we computed a constant  $K > 0$  such that can pay for our amortization. Optimize this constant.  $\diamond$

**Exercise 2.21:** Simulate the Splaysort algorithm on the following input sequence of numbers:

89, 32, 11, 52, 46, 76, 29, 99, 23, 60, 42, 24.

Show the splay tree at the end of each operation (either insert or deleteMin). In each step, indicate the cost of splaying (i.e., the length of the splay path). 1G  $\diamond$

END EXERCISES

## §2.2. Application to Splay Compression


Suppose we want to transmit a string of ASCII characters  $X = x_1x_2 \cdots x_m$ . This string has  $8m$  bits. If  $\Sigma = \{x_1, \dots, x_m\}$  is a small subset of the ASCII character set, then we can theoretically transmit a binary string  $E(X)$  of length  $nm$  where  $n = \lceil \lg |\Sigma| \rceil \leq 8$ . But besides  $E(X)$ , we probably need to also send some auxilliary string to encode the map  $x \in \Sigma \mapsto \{0, 1\}^n$ . Is there a way to transmit  $X$  *without such auxilliary data*? The answer is yes, because dynamic Huffman coding (§V.5) is such a solution.

In this section, we present another method called **splay compression** that also does not need auxilliary data. Compared to dynamic Huffman coding, its implementation is relatively simple, based on the concept of splaying (external) binary search trees.

**¶14. Representation of Prefix-free Codes by External BST.** Recall that we can represent any prefix-code  $C : \Sigma \rightarrow \{0, 1\}^*$  by a code tree (¶V.19). Such a tree is illustrated in Figure 14 where, for instance,  $C(\mathbf{e}) = 001$  and  $C(\mathbf{h}) = 0100$ .

The code tree is an external BST (¶III.40). whose leaves contain the letters of the alphabet  $\Sigma$ , arranged in the sorting order of  $\Sigma$ . In Figure 14, the sorting order on  $\Sigma$  is  $\mathbf{e} < \mathbf{h} < \mathbf{l} < \mathbf{o} < \mathbf{r} < \mathbf{w}$ . For dynamic encoding, we will further need a special letter  $*$   $\notin \Sigma$  whose sorting order is less than any  $x \in \Sigma$ . Thus the letter  $*$  is stored in the left-tip of the code tree's root. The internal nodes of the BST contain the letters of  $\Sigma$  where they are used as search keys for the purposes of lookup. Such an external BST is said to have the partner property (¶III.40).

**¶15. One-Pass Transmission and Standard Encodings.** To transmit the string  $X = x_1x_2 \cdots x_m \in \Sigma^*$  the classical Huffman code algorithm makes 2 passes over the  $X$ , once to

Figure 14: A code tree for the alphabet  $\Sigma = \{e, h, l, o, r, w\}$ 

compute the frequency function of  $X$ , and once more to encode each symbol in  $X$ . Like dynamic Huffman coding, we seek a protocol that makes only one pass over  $X$ . As each letter  $x_i \in X$  is scanned, the transmitter **emits** a binary sequence, which we denote by  $e(x_1, \dots, x_i)$  (not “ $e(x_i)$ ” because what is emitted potentially depends on the prefix  $x_1 x_2 \dots x_{i-1}$ ). Therefore, the output of the transmitter protocol on the string  $X$  can be written as

$$E(X) := e(x_1)e(x_1 x_2) \dots e(x_1 x_2 \dots x_m). \quad (10)$$

The receiver protocol, on receiving  $e(X_i)$  should be able to reconstruct  $x_i$ . This is the “online” property of our protocol. In particular, the emitted string  $e(x_1 \dots x_i)$  is self-limiting in the sense that the receiver *knows* when the last bit of  $e(x_1 \dots x_i)$  is received.

But what does it mean to “reconstruct the letter  $x_i$ ”? Here, we must assume an injective map  $ASC$  of the form

$$ASC : \Sigma \rightarrow \{0, 1\}^N \quad (11)$$

for some fixed  $N > 1$ . This map, known to both the transmitter and the receiver, will be called the **standard encoding** of  $\Sigma$ . Thus, it is important to distinguish between the symbol  $x \in \Sigma$  and the string  $ASC(x) \in \{0, 1\}^N$ . Nevertheless, in our discourse, we often treat  $x$  and  $ASC(x)$  as interchangeable. There are other standard encodings besides the ASCII code: e.g., UTF (Unicode Transformation Format) and UCS (Universal Character Set). Please see discussions in ¶V.31 (page 50) and following. Based on the standard encoding, we actually view the original string, not as  $X = x_1, \dots, x_m$ , but as

$$ASC(X) = ASC(x_1)ASC(x_2) \dots ASC(x_m).$$

Now the reconstruction problem is clear: the receiver must reproduce the string  $ASC(X)$ .

A “stateless” transmitter protocol can simply emit  $ASC(x_i)$  for the  $i$ th character. This transmits a total of  $Nn$  bits when  $|X| = n$ . The goal of dynamic string encoding is to design beat this bound of the stateless protocol. For instance, we would expect dynamic protocols to transmit  $X = \text{aaaaa}$  much more efficiently than  $X = \text{abcde}$ . This is true of our splay protocol. As a sneak preview, our splay protocol would transmit the string  $X = \underbrace{\text{aaa} \dots \text{aa}}_n$  as

$$E(X) = ASC(a) \underbrace{111 \dots 11}_{n-1}$$

using only  $7 + n$  bits, instead of  $8n$  bits. As  $n \rightarrow \infty$ , this transmission rate approaches 1 bit per character in  $X$ . This is clearly optimal. This is a data compression ratio of 8 for the ASCII character set. For this reason, we also view the splay solution as a string compression algorithm.

¶16. **External Splay Trees.** We recall the algorithms for lookup, insertion and deletion for external BSTs (§III.42). These algorithms are even simpler than the corresponding algorithms for standard BST. Next, to turn an external BST into an “external splay tree”, we only need to stipulate the following additional “splay feature” to these algorithms:

Let  $x$  be the external node that we reach in a lookup/insertion/deletion:  
 In case of lookup/insertion, we must splay the parent  $u$  of  $x$ .  
 In case of deletion, both  $x$  and its parent  $u$  would have been deleted; we must splay the parent of  $u$ .

The external BSTs which are maintained using these splay algorithms are called **external splay trees**. Observe that we only splay internal nodes, never an external node. The result remains an external BST.

*Why can't we splay external nodes?*

What does the above splaying rule do? Suppose  $u$  is the parent of a leaf  $x$  in the external splay tree. Note that splaying  $u$  will bring  $u$  to the root. But what happens to  $x$ ? Figure 15 shows an example of splaying  $u$ , which reduces the depth of  $x$  from 6 to 4.




Figure 15:  $\text{splay}(u)$  transforms the initial depth-pair  $(5, 1)$  to  $(0, 4)$ .

**Lemma 8** Let  $x$  be an external node in external BST  $T$ . If  $u$  is the parent of  $x$  with  $\text{depth}(u) = d \geq 1$ , then after  $T.\text{splay}(u)$ ,

- The depth of  $x$  is at most  $1 + \lfloor d/2 \rfloor$ .
- The depth of  $x$  is reduced by at least  $\lfloor d/2 \rfloor$ .

*Proof.* Let  $\text{depth}(u, x)$  denote the depth of  $x$  in the subtree rooted at  $u$ . Consider the “depth pair” of  $u, x$ ,

$$D(u, x) := (\text{depth}(u), \text{depth}(u, x)) = (a, b).$$

Consider how this pair of integers is transformed by performing splaying  $u$ . Suppose  $D'(u, x)$  is the depth pair after  $\text{splayStep}(u)$ . Assume  $a \geq 1$  in  $D(u, x) = (a, b)$ . The  $\text{splayStep}$  falls under one of 3 cases:

- When  $a \geq 2$ , this is Case I or Case II. From Figure 16, we see that

$$D'(u, x) = (a - 2, b) \text{ or } (a - 2, b + 1).$$

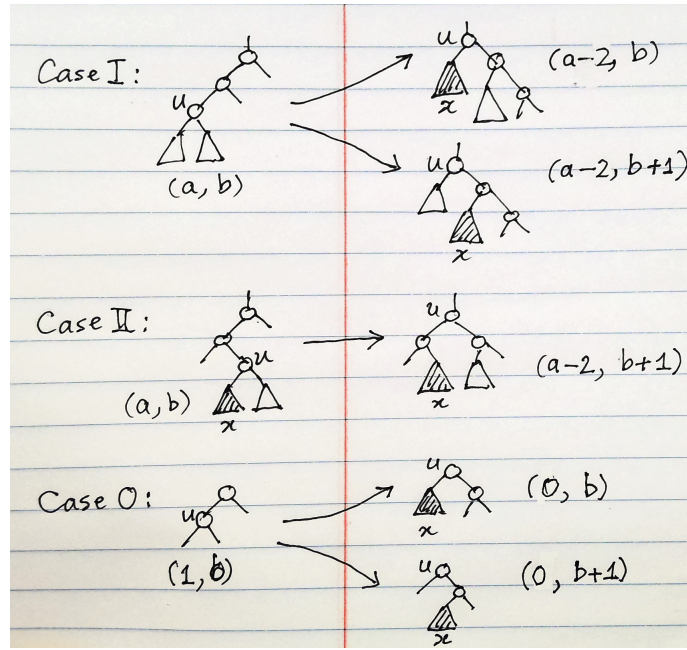


Figure 16: Transformations of depth-pairs

- When  $a = 1$ , this is Case 0. From Figure 16 shows that  $D'(u, x) = (0, b)$  or  $(0, b + 1)$ .

Thus, in all 3 cases,  $b$  increases by at most 1. If the initial depth-pair is  $(d, 1)$ , there are  $\lceil d/2 \rceil$  splaySteps because  $d$  will be reduced to 0 in these many splaySteps. Since each splayStep increases  $\text{depth}(u, x)$  by at most 1, the final value of  $\text{depth}(u, x)$  is at most  $1 + \lceil d/2 \rceil$ . Since  $u$  is finally the root, the final depth of  $x$  is  $\leq 1 + \lceil d/2 \rceil$ .

Since the initial depth of  $x$  is  $1 + d$  and final depth is  $\leq 1 + \lceil d/2 \rceil$ , we conclude that  $\text{splay}(u)$  has reduced the height of  $x$  by  $\geq (1 + d) - (1 + \lceil d/2 \rceil) = \lfloor d/2 \rfloor$ . **Q.E.D.**

This lemma shows that splaying the parent of  $x$  in an external BST only reduces the depth of  $x$  by half. Thus, splaying is a “logarithmic form” of the move-to-front heuristic because we need  $\lg d$  splays to move  $x$  to a depth  $\leq 2$ . Note that  $\text{depth}(x) \geq 2$  if  $x$  does not represent the smallest or largest key in the BST.

**¶17. The Splay Protocols.** We now describe the transmitter and receiver protocols for transmitting a string  $X = x_1 x_2 \cdots x_m \in \Sigma^*$ . The “state information” that is represented by an external BST  $T$ . Let  $T_i$  be the state of  $T$  just before  $x_i$  is scanned. Then the external nodes of  $T_i$  stores the set  $S_i \cup \{*\}$  where

$$S_i = \{x_1, \dots, x_i\} \quad (i = 0, \dots, n).$$

In particular,  $S_0$  is the empty set, and  $T_0$  has only one external node containing  $*$ , which is also the root. Let

$$C_i : S_i \cup \{*\} \rightarrow \{0, 1\}^*$$

denote the prefix-free code represented by  $T_i$ .



## TRANSMISSION PROTOCOL:

- The transmission protocol at the  $i$ th step ( $i = 1, \dots, n$ ) falls under one of two cases, depending on whether  $x_i \in S_{i-1}$  or not. To decide the case, we perform a LookUp of  $x_i$  in  $T$ , ending up at an external node  $b$ . We also record the binary string  $\beta$  representing the path from root to  $b$ .
- EASY CASE:  $x_i$  is found ( $b.\text{key} = x_i$ ). We can emit the string  $\beta$  which represents  $C_{i-1}(x_i)$ . Then we splay  $p$ , the parent of  $b$ .
- HARD CASE:  $x_i$  is not found ( $b.\text{key} \neq x_i$ ). We take these actions:

1. Emit the sequence  $C_{i-1}(*)$ .  
This sequence is the LookUp path on  $*$ .  
Let  $q$  be the the parent of the  $*$ -node.
2. Emit  $\text{ASC}(x_i)$ .  
This tells the receiver about the new symbol  $x_i$ .
3. Insert  $x_i$  into the splay tree.  
Using node  $p$ , parent of  $b$  from the initial (failed) Lookup on  $x_i$ , this insertion takes  $O(1)$  time.
4. Splay  $p$
5. Splay  $q$ .

## RECEIVER PROTOCOL:

- The receiver maintains an external BST  $R$  that mirrors the transmitter's external BST  $T$ . In the  $i$ th step, the receiver receives the  $i$ th emission from the transmitter, updates  $R$  so that  $R$  will be identical to  $T$ , and then it must **output** its own string: this output should be  $\text{ASC}(x_i)$ .
- In general, the received emission represents a path from the root of  $R$  to an external node  $b$ . If  $b.\text{key} \neq *$ , this is the EASY CASE, and the receiver can just output  $\text{ASC}(b.\text{key})$ . Of course, the parent  $q$  of  $b$  must be splayed.
- If  $b.\text{key} = *$ , this is the HARD CASE, and the receiver knows that the next  $N$  bits represent  $\text{ASC}(x_i)$ . So the receiver can immediately output  $\text{ASC}(x_i)$ . To update  $R$ , the receiver must first insert  $x_i$  into  $R$ , and splay the parent  $p$  of the newly inserted  $x_i$ . Then it must splay  $q$ , the parent of  $b$ .
- It should be clear that the receiver splay tree  $R$  is now perfectly synchronized with the transmitter splay tree  $T$ . It is interesting to notice an asymmetry in the HARD CASE: the transmitter protocol discovers node  $p$  before node  $q$ , while the receiver protocol discovers node  $q$  before node  $p$ . But they both splay  $p$  before  $q$ .
- The first step is special: the receiver knows that this path is the empty string and so the first  $N$  bits is  $\text{ASC}(x_1)$  which it can insert into  $R$ . It can also output  $\text{ASC}(x_1)$ .

Let us see what this protocol will transmit in case  $X = \underbrace{aaa \cdots aa}_n$ . At the beginning, a LookUp on the first character  $a$  in  $X$  will fail, and the path to the  $*$ -node is the empty string. So the transmitter just transmits  $\text{ASC}(a)$ . But each subsequent occurrence of  $\text{ASC}(a)$  in  $X$  will cause a single binary bit 1 to be transmitted. This confirms our earlier remark that  $E(\underbrace{aaa \cdots aa}_n) = \text{ASC}(a) \underbrace{111 \cdots 11}_{n-1}$ .

**¶18. Analysis.** To bound the length of the transmitted string  $E(X)$ , we exploit the key result about splay trees, *viz.*, each splay operation has amortized logarithmic time.

**Lemma 9** *If  $X = x_1x_2 \cdots x_m$  has  $n$  distinct characters, then the total cost of transmitting  $X$  is  $O((m+n)\log n)$ . Moreover, the length of  $E(X)$  is  $O((m+n)\log n)$ .*

*Proof.* To transmit  $X$ , we perform  $m$  LookUps. In  $n$  of these LookUps, we must also insert a new leaf and splay the \*-node. This results in  $m+n$  splays, each with an amortized cost of  $O(\log n)$ . Also, when we attach a new leaf to a splay tree, we increase the potential of the tree by  $O(\log n)$  (Lemma 6). This yields the  $O((m+n)\log n)$  total cost. Clearly, the length of  $E(X)$  cannot exceed the cost to compute  $E(X)$ . **Q.E.D.**

It would be an interesting exercise to determine the implicit constant in the  $O((m+n)\log n)$ .

---

EXERCISES

**SIMULATION HINT** for external Splay trees: you will usually find it unnecessary to draw external nodes: just indicate their presence by an edge to no-where. This will save half of your artistic effort.

**Exercise 2.22:** Splay Compression (15+5 Points)

Consider the ASCII string representing a baby's magical incantation:

$$X = \text{abadacadaba}$$

Note that  $X$  has length 11. We want to compute the binary string  $E(X)$  obtained by our splay compression protocol (see ¶VI.2.2, page 26).

- (a) First, show the sequence of splay trees,

$$T_0 \rightarrow T_1 \rightarrow T_2 \rightarrow \cdots \rightarrow T_{11}$$

where  $T_i$  is the splay tree after processing the  $i$ th character in  $X$ .

Note that a transition  $T_i \rightarrow T_{i+1}$  may take one or two splays. Please introduce as many intermediate trees between  $T_i$  and  $T_{i+1}$  as there are splaySteps.

- (b) What is the total length of  $E(X)$ ? Also, show us the breakdown of  $E(X)$  into the 11 bit substrings that were emitted after the scan of each character of  $X$ .

In the notation of the text (¶VI.15, p.27), these substrings are

$$e(\mathbf{a}), e(\mathbf{ab}), e(\mathbf{aba}), \dots$$

You may write  $\text{ASC}(x)$  for ASCII representation of any character  $x$ .

◇

**Exercise 2.23:** Consider the string

$$X = (\text{abcba})^m$$

where  $a, b, c$  are distinct ASCII characters and  $m \geq 2$ .

(a) How many bits have you transmitted after processing the first 5 characters ( $abcba$ ) in string  $X$ ? Show the external splay tree at this stage.

(b) What is the total number of bits transmitted for  $X$ ? ASSUME that after inserting a new character  $x$ , you must first splay the parent of  $*$  leaf, then splay the parent of  $x$  leaf.

◇

**Exercise 2.24:** How many bits are transmitted by our Splay Compression algorithm for the string  $a^n b^n c^n$  where  $a, b, c$  are distinct characters?

◇

**Exercise 2.25:** Consider the transmission of the string “hello world!” using the Dynamic Huffman tree algorithm and the splay tree algorithm.

(a) First describe the bit string that is transmitted using the Dynamic Huffman tree algorithm. Show the successive Huffman trees resulting from processing successive input characters, and indicated the transmitted bits. How many bits in total are transmitted?

(b) Do the same, but using the Splay compression.

(c) Draw some conclusion about the relative merits of the two methods.

◇

**Exercise 2.26:** Our protocol above assumes a setting where both parties knows when the transmission has begun, and when it has stopped. Design a hand-shaking protocol that takes care of this detail. Make any reasonable assumptions you need.

◇

**Exercise 2.27:** We want to transmit an ASCII string  $X$  using our external splay tree protocol. Here  $|X| = 3n$ . As  $n \rightarrow \infty$ , how many bits per characters are transmit by our protocol in the following scenarios:

(a)  $X = aabaab \cdots aab = (aab)^n$

(b)  $X = X_1 X_2 \cdots X_n$  where each  $X_i \in \{aaa, bbb\}$ ?

(c)  $X = X_1 X_2 \cdots X_n$  where each  $X_i \in \{aaa, bbb, aba, bab\}$ ?

◇

**Exercise 2.28:** (a) Show the result of inserting the following characters,  $h, e, l, o, w, r, d$  (in this order) into an external splay tree that initially contains the character  $*$ .

Assume the characters have the usual alphabetic sorting order but  $*$  is smaller than any other character. REMARK: it suffices to show the external splay tree after each insertion.

(b) Show the tree after deleting  $w$  from the final tree in part(a).

◇

**Exercise 2.29:** Please use our external splay tree to transmit the string **Hello World!** using the standard sorting order of the ASCII characters. How many bits are transmitted?

◇

**Exercise 2.30:**

(a) Program the above transmission and receiving protocols based on external splay trees in your favorite programming language.

(b) Do the same for the Dynamic Huffman tree protocol (§V.4).

(c) Compare the performance of the splay tree protocol with that of the Dynamic Huffman Tree for Lincoln’s speech at Gettysburg (Exercise, §V.4).

◇

**Exercise 2.31:** Here is an ASCII string  $X = b^m a^n c^p$  where  $m, n, p$  are positive integers. So  $|X| = m + n + p$ . The ASCII ordering on the characters is  $a < b < c$  (of course). Recall

our protocol for transmitting strings using External Splay trees.

(a) If  $m = 10, n = 2, p = 20$ , what is the length of  $E(X)$ ? The answer is a single number, but explain how you get this number.

(b) TRUE or FALSE: there exist constants  $\alpha, \beta, \gamma, \delta$  such that the length of  $E(X)$  is equal to

$$|E(X)| = \alpha m + \beta n + \gamma p + \delta$$

840

◇

**Exercise 2.32:** We want a dynamic transmission protocol based on external splay trees. Suppose the string  $X$  to be transmitted is ‘bilingual’. That is,  $X$  is a string over two character sets  $\Sigma_A$  and  $\Sigma_B$ , with standard encodings  $ASC_\mu : \Sigma_\mu \rightarrow \{0, 1\}^{N_i}$  ( $\mu = A$  or  $\mu = B$ ) for each character set.

(a) Give a solution to this problem using external splay trees.

(b) Suppose there are  $n_\mu$  ( $\mu = A, B$ ) characters from  $\Sigma_\mu$  in  $X$  (so  $|X| = n_A + n_B$ ). Also the number of distinct characters among the  $n_\mu$  characters is  $s_\mu$  (so  $s_\mu \leq n_\mu$ ). Show that the transmitted string has length  $O(s_A \log n_A + s_B \log n_B)$

848

◇

**Exercise 2.33:** What is the “worst case” behavior of our external splay tree protocol? One notion of worst case is when  $E(X_i)$  has  $n$  bits when the tree  $T$  has  $n + 1$  external nodes. Can you force this to happen arbitrarily often?

851

◇

**Exercise 2.34:** In our protocol, we do not use deletion in the external splay tree  $T$ . To introduce deletion, we introduce an upper bound  $B$  on the number of external nodes in  $T$ . When we need to insert a new character  $c'$  when  $T$  already has  $B$  external nodes, we need to pick a character  $c$  in  $T$  to delete. For this purpose, we maintain the “last access time”  $LAT(c)$  for each  $c$  in  $T$ . E.g., if  $x_i = c$  then we set  $LAT(c) = i$  after we transmit  $x_i$ . Then, before we insert  $c'$ , we first delete the character  $c$  whose  $LAT(c)$  is minimum.

858

◇

**Exercise 2.35:** Simplified<sup>5</sup> Dynamic Huffman Code.

let  $T$  be the external BST in Figure 17 whose leaves store the set  $\Sigma = \{*, e, h, l, o, r, w\}$  of characters with the sorting order

$$* < e < h < l < o < r < w.$$

859

Given  $x \in \Sigma$ , we can compute its code  $C(x)$  by doing  $\text{Lookup}(x)$  in  $T$ , and  $C(x)$  is just the corresponding search path. E.g.,  $C(h) = 0100$ , and  $C(w) = 11$ . Thus,  $T$  represents a prefix-free code for  $\Sigma$ .

860

861




Figure 17: An external BST

<sup>5</sup>This is simplified because we do not do insertion and deletion from  $T$ .

Let  $s \in \Sigma^*$  be a string to be transmitted. Assuming that receiver also has a copy of the tree  $T$ , the receiver would be able to decode our message. We look at two scenarios: a static and a dynamic way to use  $T$  for transmission.

1. We first assume that  $T$  is static (unchanging). How many bits does it take to transmit  $C(s)$  where

$$s = \text{hellloooo} \quad (12)$$

Please show the encoded string  $C(s)$ .

2. Next, assume that  $T$  is dynamic. After each  $\text{Lookup}(x)$ , you will immediately transmit the code  $C(x)$ . *But you further splay the parent of  $x$ .* Hence the next time you do  $\text{Lookup}(x)$ , the code  $C(x)$  may be different! Note that the receiver can do the same actions. How many bits does it take to transmit the string (12) of part(a)? Again, show us the encoded binary string  $C(s)$ .

◇

**Exercise 2.36:** In our usual transmission model, each protocol maintains an external BST  $T$  that represents a prefix-free code  $C_T : S \rightarrow \{0, 1\}^*$ . The 2 copies of  $T$  are synchronized. We explore the scenario of limited memory: this means  $|S| \leq k$  for some fixed  $k \geq 0$  known to both parties. So when  $|S| = k$ , the only way to insert a new character  $\mathbf{b}$  requires a simultaneous deletion of some  $\mathbf{a} \in S$ . The transmitter can emit the sequence

$$C_T(*)ASC(\mathbf{a})ASC(\mathbf{b}).$$

When receiver notices that  $\mathbf{a}$  is in  $S$ , it is a signal to delete  $\mathbf{a}$  and insert  $\mathbf{b}$  into  $S$ . How should  $\mathbf{a}$  be chosen?

◇

---

 END EXERCISES

## §2.3. Application to Convex Hulls

In many elementary problems of Computational Geometry, we are interested in storing and searching a finite set  $X$  of  $d$ -dimensional points:  $X \subseteq \mathbb{R}^d$  where  $d \geq 1$ . Of course  $d = 1$  is the case of binary search trees (BSTs). Here we will assume  $d = 2$ , so  $X$  is a planar set of points. Can we still use BST's in this case? It depends on what kind of queries we need. In Chapter III, we indicated a 2-dimensional BST in which we make “box queries” of the form: return the subset  $X \cap B$  for any given box  $B = [a, b] \times [c, d]$ . In the present section, we study a situation where we use a variant of our 1-dimensional BST. This is possible because  $X$  has very special linear structure.

**¶19. Geometric Predicates.** In plane Euclidean Geometry, the most basic objects are points which are elements of  $\mathbb{R}^2$ . If  $p \in \mathbb{R}^2$ , we can also write  $p = (p_x, p_y)$ . Thus we have already introduced a coordinate system on our plane. A plane with some fixed coordinate system is called the Cartesian Plane (after Descartes).

**Introduction to Analytic Euclidean Geometry.** Classical Euclidean Geometry is axiomatic: beginning with a small set of axioms, we derive various properties. The most famous of these axioms are the 5 Postulates of Euclid. Starting with Descartes, we realized that such axiomatic reasoning can be reduced to algebraic calculations! This dream of reducing all human reasoning to calculations has inspired many philosophers, mathematicians and scientists, all the way to our present day. In Geometry, Descartes' idea gave birth of Analytic Geometry.

Let  $p, q \in \mathbb{R}^2$ . If  $p \neq q$ , then there is a unique line  $\overline{p, q}$  passing through  $p$  and  $q$ . In general, a line  $L$  is represented by a linear equation

$$\ell(x, y) = ax + by + c = 0 \quad (13)$$

with coefficients  $a, b, c \in \mathbb{R}$  with  $ab \neq 0$ . Call  $(a, b, c)$  a coefficient triple for  $L$ . A point  $p$  lies on  $L$  iff  $\ell(p_x, p_y) = 0$ . E.g., if  $L$  passes through the points  $(0, 1)$  and  $(2, 3)$ , then  $\ell(x, y) = x - y + 1$  (i.e.,  $a = 1, b = -1, c = 1$ ) is the linear equation for  $L$ . Note that if  $(a, b, c)$  is a coefficient triple for  $L$ , then so is  $(da, db, dc)$  for any  $d \neq 0$ . We could write  $a : b : c$  to denote the equivalent triples.

A **directed line** is a line with one of two directions. Let  $\overrightarrow{p, q}$  denote the line  $\overline{p, q}$  with the direction from  $p$  to  $q$ . Note that  $\overrightarrow{q, p}$  then represents the same line (since  $\overline{p, q} = \overline{q, p}$ ) with the opposite direction, from  $q$  to  $p$ . Note that  $(a, b, c)$  is a coefficient triple for  $\overrightarrow{p, q}$  iff  $(-a, -b, -c)$  is a triple for  $\overrightarrow{q, p}$ .

We proceed intuitively here, taking advantage of the fact that humans (not computers) have a lot of geometric intuitions which are in turn based on our experiences in the physical world of space and time. Imagine the following physical motion involving 3 points  $p, q, r$ : starting at  $p$ , we move towards  $q$  and then move towards  $r$ . We say that  $(p, q, r)$  represents a “left-turn” if it is necessary to make a left-turn at  $q$  in the preceding motion.

So we want a predicate  $\text{LeftTurn}(p, q, r)$  that is true iff  $(p, q, r)$  represents a left turn. This is one of the most important predicates in planar geometry, and it often called the **orientation predicate**. But we prefer the self-descriptive name **LeftTurn Predicate**. Here is the full definition:

$$\text{LeftTurn}(p, q, r) = \begin{cases} +1 & \text{if the movement } (p, q, r) \text{ make a left turn at } q \\ 0 & \text{if the points } p, q, r \text{ are collinear} \\ -1 & \text{if the movement } (p, q, r) \text{ make a right turn at } q. \end{cases} \quad (14)$$

By definition, a set of points of points  $S$  is **collinear** if all the points in  $S$  lies in a single line. For instance, if  $p = q$  then the set  $\{p, q, r\}$  is collinear. Note that we call **LeftTurn** a “predicate” even though it is a 3-valued function. In logic, predicates are usually 2-valued (i.e., true or false). Having 3 values is a common phenomenon in geometry, and so we call any 3-valued predicate a **geometric predicate**. Also, let **logical predicate** refer to any 2-valued function.

For any  $p, q, r \in \mathbb{R}^2$ , define the  $3 \times 3$  matrix

$$M = M(p, q, r) := \begin{bmatrix} p_x & p_y & 1 \\ q_x & q_y & 1 \\ r_x & r_y & 1 \end{bmatrix}.$$

Let  $\det(M)$  denote the determinant of any square matrix  $M$ . For instance, you can easily check

that  $\det(M(p, q, r))$  reduces to a  $2 \times 2$  determinant

$$\det(M(p, q, r)) = \det \begin{bmatrix} q_x - p_x & q_y - p_y \\ r_x - p_x & r_y - p_y \end{bmatrix} = (q_x - p_x)(r_y - p_y) - (r_x - p_x)(q_y - p_y).$$

Also, for any real number  $x$ , the sign function is

$$\text{sign}(x) := \begin{cases} 0 & \text{if } x = 0, \\ +1 & \text{if } x > 0, \\ -1 & \text{if } x < 0. \end{cases}$$

Descartes's idea to turn geometric reasoning to calculations is illustrated in a small way by the following lemma:

### Lemma 10

$$\text{LeftTurn}(p, q, r) = \text{sign}(\det(M(p, q, r))).$$

We leave the proof to an Exercise. We can define two related predicates using **LeftTurn**:

$$\begin{aligned} \text{RightTurn}(p, q, r) &:= -\text{LeftTurn}(p, r, q) \\ \text{Collinear}(p, q, r) &:= [\text{LeftTurn}(p, q, r) = 0] \end{aligned} \quad (15)$$

where  $[\dots]$  is the logical predicate that is true iff “ $\dots$ ” is true. Thus  $\text{RightTurn}(p, q, r) = +1$  iff  $\text{LeftTurn}(p, r, q) = -1$ , and  $\text{Collinear}(p, q, r) = \text{true}$  iff  $\text{LeftTurn}(p, q, r) = 0$ .

We say a set  $X$  of points is **degenerate** if there exist three distinct points  $p, q, r \in X$  that are collinear; otherwise,  $X$  is **nondegenerate**. We will assume our input set  $X$  is nondegenerate. The motivation for this assumption is pedagogical: the non-degenerate cases are easier to understand, and there are relatively few non-degenerate cases. We note that there are general techniques in computational geometry for handling degeneracies, but it is beyond our scope.

**¶20. Upper Hulls** First let us see what special structure  $X$  must have. Our problem is to compute the convex hull  $CH(S)$  of a finite set  $S \subseteq \mathbb{R}^2$ . This convex hull is represented by the circular sequence of points. For example, Figure 18(a) shows a circular sequence  $H = [v_1, v_2, \dots, v_9]$ , representing the convex hull  $CH(S)$ , i.e., the bounded yellow region. In general,  $\{v_1, \dots, v_9\}$  is a proper subset of  $S$ . We can decompose  $H$  into an “upper convex chain”  $X = (v_1, v_2, \dots, v_6)$  and a “lower convex chain”  $Y = (v_1, v_9, v_8, v_7, v_6)$ . The chain  $X$  defines an “upper hull”  $U_X$  illustrated by the unbounded yellow region in Figure 18(b).

Formally, an ordered sequence

$$X = (v_1, v_2, \dots, v_k), \quad (k \geq 2) \quad (16)$$

is **upper chain** if the  $v_i$ 's satisfy

$$v_1 <_x v_2 <_x \dots <_x v_k \quad (17)$$

where<sup>6</sup>  $a <_x b$  means  $a.x < b.x$ . Moreover, every consecutive triple  $(v_{i-1}, v_i, v_{i+1})$  is a **right-turn** (see above definition).

Consider the sequence

$$X = (v_1 - v_2 - \dots - v_k). \quad (18)$$

<sup>6</sup>Similarly, we may write  $a \leq_x a$  or  $a =_x a$ . We can also replace  $x$  by the  $y$ -coordinate and write  $a <_y a$ , etc.






Figure 18: (a) Convex hull represented by  $H = [v_1, \dots, v_9]$ ; (b) Upper hull represented by  $X = (v_1, \dots, v_6)$ .

Let  $Path(X)$  denote the continuous polygonal path from  $v_1$  to  $v_k$  that is the union of  $k - 1$  line segments. The **upper hull** of  $X$  is the set

$$U_X = \{a \in \mathbb{R}^2 : (\exists p \in Path(X)[a \leq_y p])\}$$

comprising those points  $a$  that lie vertically below some point  $p$  in  $Path(X)$ . In case every consecutive triple in (16) is a left-turn we call  $X$  a **lower chain** and its **lower hull**  $L_X$  is similarly defined.

Alternatively, we can introduce 2 special points

$$v_S := (0, -\infty), \quad v_N := (0, +\infty) \quad (19)$$

called the **South** and **North poles**. Then the upper hull  $U_X$  is just the convex hull of the set  $\{v_1, \dots, v_k\} \cup \{v_S\}$ . Likewise, if (18) represents a lower convex chain, then the lower hull is the convex hull of  $L_X = \{v_1, \dots, v_k\} \cup \{v_N\}$ .

**¶21. Tangent Queries** We may surmise that an upper chain such as (16) can be stored in a BST using the values  $v_i.x$  as keys. This is almost correct, but we will actually store the entire point  $v_i$  as a key (with the BST order based only on  $v_i.x$ ). We need the extra information provided by  $v_i.y$  because of the 2-dimensional nature of our queries.

We define the concept of a **left-tangent point** of  $p$  (relative to the upper chain  $X = (v_1, \dots, v_k)$  in (18)). The concept is undefined if  $p$  lies inside the upper hull  $U_X$ . Moreover, if  $p \leq_x v_1$ , we shall define the left-tangent point of  $p$  to be the special point  $v_S$  (south pole in (19)). Otherwise, assume that

$$v_1 <_x p \quad \text{and} \quad p \notin U_X. \quad (20)$$

In this case, there is a smallest index  $i = 1, \dots, k$  in (18) such that

$$v_i <_x p \quad \text{and} \quad \begin{cases} \text{LeftTurn}(p, v_i, v_{i-1}) = +1 & \text{if } i > 1, \\ \text{LeftTurn}(p, v_1, v_2) \in \{0, +1\} & \text{if } i = 1. \end{cases} \quad (21)$$

We call this point  $v_i$  the **left-tangent point** of  $p$  (relative to  $X$ ). Note that if  $i = 1$ , we allow  $\text{LeftTurn}(p, v_1, v_2) = 0$ . See further discussion of this in the box paragraph below on degeneracy. In Figure 18(b), the left-tangent point of  $p$  is  $v_3$  since  $v_3 <_x p$  and  $\text{LeftTurn}(p, v_3, v_2) = +1$ .

The left-tangent point of  $p$  is only defined if  $p \notin U_X$  where  $X = (v_1, \dots, v_k)$ . How do we check this? Assuming that  $v_1 \leq_x p \leq v_k$ , then there is a unique  $i = 1, \dots, k-1$  such that  $v_i \leq_x p \leq v_{i+1}$ . Then  $p$  lies outside  $U_X$  iff  $(v_i, v_{i+1}, p)$  is a left-turn.

Similarly, the **right-tangent point** of  $p$  (relative to  $X$ ) is not defined if  $p$  lies inside the upper hull  $U_X$ . Moreover, if  $p \geq_x v_k$ , we shall define the right-tangent point of  $p$  to be the south pole  $v_S$  as before. Otherwise, there is a largest index  $i = 1, \dots, k$  in (18) such that

$$v_i >_x p \quad \text{and} \quad \begin{cases} \text{RightTurn}(p, v_i, v_{i+1}) = +1 & \text{if } i < k, \\ \text{RightTurn}(p, v_k, v_{k-1}) \in \{0, +1\} & \text{if } i = k. \end{cases} \quad (22)$$

The **right-tangent point** of  $p$  is this  $v_i$ . Again, the case  $i = k$  allows the possibility that  $\text{RightTurn}(p, v_k, v_{k-1}) = 0$ , i.e.,  $p, v_k, v_{k-1}$  are collinear. Collectively, the left- and right-tangent points of  $p$  are just called the **tangent points** of  $p$ .

**On Degeneracies in Convex Hulls.** The definitions of left- and right-tangent points allows the possibility that  $\text{LeftTurn}(p, v_1, v_2) = 0$  and (respectively)  $\text{RightTurn}(p, v_k, v_{k-1}) = 0$ . For instance, in Figure 18(b), we see that the left- and right-tangent points of  $p$  are  $v_3$  and  $v_6$ , respectively. We are illustrating the special case where  $(p, v_5, v_6)$  are collinear, i.e.,  $\text{RightTurn}(p, v_6, v_5) = 0$ . Hence the right-tangent point is not  $v_5$  but  $v_6$ , according to the requirement that the index  $i$  in (22) be maximum.

These ideas are related to the notion of **degeneracies** in a finite point set  $S \subseteq \mathbb{R}^2$ . In general, given any 3 distinct points  $a, b, c \in S$ , we expect that they do not lie on a single line, i.e.,  $\text{LeftTurn}(a, b, c) \neq 0$ . But when it happens, we say that  $S$  is **degenerate**. Let us say that point  $b \in S$  is an **extremal** point of  $S$  if there exists a line  $\ell$  passing through  $b$  such that  $S$  lies *strictly to one side* (either the left- or right-side) of  $\ell$ . We can define the **convex hull** of  $S$  to comprise all the extremal points of  $S$ . If we drop the *strictly* requirely, allowing some points of  $S$  to lie on  $\ell$ , then then we say  $b$  is **semi-extremal**. When  $S$  is degenerate, it is possible that some  $b \in S$  is semi-extremal but not extremal. Whether we want to output such points depends on the application.

The following is easily verified:

**Lemma 11** *When  $p \notin U_X$ , the left- and right-tangent points are well-defined.*

The **tangent query** for  $p$  relative to an upper chain  $X$  is this:

$$\text{tangent}(p, X) = \begin{cases} (v_\ell, v_r) & \text{if } p \notin U_X, \\ \uparrow & \text{else.} \end{cases} \quad (23)$$

where  $v_\ell$  and  $v_r$  are the left- and right-tangent points of  $p$ .

**¶22. Updating the Upper Convex Chain** Suppose  $\text{tangent}(p, X) = (v_\ell, v_r)$ . Then we want to update the upper chain from (16) to

$$X' = (v_0, v_1, \dots, v_\ell, p, v_r, v_{r+1}, \dots, v_k). \quad (24)$$

This operation may be denoted

$$X' \leftarrow \text{insert}(p, X). \quad (25)$$

It amounts to replacing the subchain  $(v_{\ell+1}, \dots, v_{r-1})$  by  $p$ . If the subchain is empty, then we are just inserting  $p$  into the  $X$ . Otherwise, we need to split off from  $X$  two subchains

$$X_L = (v_0, v_1, \dots, v_\ell), X_R = (v_r, v_{r+1}, \dots, v_k).$$

Then we merge  $X_L, X_R$  and finally insert  $p$  into the merger.

**¶23. Reduction to Fully Mergeable Dictionary Operations.** If  $T$  is a BST representing  $X$ , then we can write

$$T.\text{insert}(p) \quad (26)$$

to denote the self-modifying operation on  $T$  with the same meaning as “ $X \leftarrow \text{insert}(p, X)$ ” (cf. (25)). We reduce (26) to the operations of a fully mergeable dictionary ADT (¶III.10):

`lookUp, insert, delete, merge, Split.`

Specifically, we want the BST  $T$  to be a splay tree (¶7).

By calling the split operation twice, we produce three subchains,  $T_L : (v_0, \dots, v_\ell)$ ,  $T_M : (v_{\ell+1}, \dots, v_{r-1})$  and  $T_R : (v_r, \dots, v_k)$ . Then we call  $T_L.\text{merge}(T_R)$  and finally insert  $p$  into  $T_L$ .

Each node  $u \in T$  has a member  $u.\text{key}$  which is equal to a point  $v_j$  in the upper chain  $X$ . We also assume that  $u$  has members  $u.\text{succ}$  and  $u.\text{pred}$  pointing to successor and predecessor of  $u$ . We next show how to implement the requests

$$\text{insert}(p, T) \quad \text{and} \quad \text{tangent}(p, T).$$

Note that both  $\text{insert}(p, T)$  and  $\text{tangent}(p, T)$  must take a common first step to decide if  $p$  is outside the upper hull of  $T$  or not. If not, then both algorithms return  $\uparrow$ . So our first task is:

**¶24. Deciding if  $p$  is inside the upper hull.** We assume the splay tree  $T$  represents the upper chain  $X$ . If

$$p <_x v_0 \quad \text{or} \quad p >_x v_k$$

then clearly,  $p$  lies outside the upper hull  $U_X$ . Otherwise, we do `lookUp` on  $p.x$  in  $T$ . This returns a node  $u \in T$  that is either the successor or predecessor of  $p.x$  in  $T$ . Using the successor and predecessor pointers of  $u$ , we can easily determine the unique  $i$  such that

$$v_i \leq_x p <_x v_{i+1}.$$

If  $p.x = v_i.x$  then  $p$  is inside  $U_X$  iff

$$p.y \leq v_i.y$$

Otherwise, we have  $v_i <_x p <_x v_{i+1}$ . Clearly,

$$p \notin U_X \iff \text{LeftTurn}(v_i, v_{i+1}, p).$$

This concludes our decision procedure to check if  $p \in U_X$ .

This algorithm is based on an *explicit* binary search on the key  $p.x$ . It should be contrasted with an *implicit* binary search for tangent points below.

Assuming that  $p \notin U_X$ , we proceed to determine the tangent points for  $p$ . But keep in mind that we now know which one of the following three conditions hold:

$$(p <_x v_0) \quad \text{or} \quad (p >_x v_k) \quad \text{or} \quad (v_i <_x p <_x v_{i+1}) \quad (27)$$

¶25. **Method One for Tangent Point: Walking Method.** The first method is simpler and more intuitive. By symmetry, let us only discuss finding the left tangent point. If  $p <_x v_0$ , then the left tangent point is  $v_S$ . So we may assume from (27) that we know the node  $v_i$  of the

$$v_i <_x p <_x v_{i+1}.$$

Let  $v_\ell$  denote the left tangent point. Clearly,  $\ell \leq i$ . So, using the predecessor pointers to “walk” along the upper hull, starting from  $v_i$  to  $v_{i-1}, v_{i-2}$ , etc.

In general, for any index  $j$  with the property that  $v_j <_x p$ , we can decide whether  $\ell = j$ ,  $\ell < j$  or  $\ell > j$  by testing the orientation of these two triples:  $(p, v_j, v_{j-1})$  and  $(p, v_j, v_{j+1})$ .




Figure 19: LeftTangent Search from  $p$ : (i)  $\ell = j$ , (ii)  $\ell < j$ , (iii)  $\ell > j$ .

**CASE (i)**  $(p, v_j, v_{j-1})$  is a left-turn, and  $(p, v_j, v_{j+1})$  is not a right-turn.  
Then  $\ell = j$ .

**CASE (ii)**  $(p, v_j, v_{j-1})$  is a right-turn  $(p, v_j, v_{j+1})$  is a left-turn.  
Then  $\ell < j$ .

**CASE (iii)**  $(p, v_j, v_{j-1})$  is a left-turn  $(p, v_j, v_{j+1})$  is a right-turn.  
Then  $\ell > j$ .

These three cases are illustrated in Figure 19. Case (i) is a bit subtle: it includes the case where both triples are left-turns. But it also includes the possibility where the second triple  $(p, v_j, v_{j+1})$  is collinear. This degeneracy might arise for an arbitrary query point  $p$ . Of course, in our “walking method”, we never encounter case (iii); but case (iii) will be needed in our next method.

¶26. **Method Two: Implicit Binary Search.** The previous method for finding tangent points  $(q, r) = (v_\ell, v_r)$  takes linear time in the worst case. But if we were inserting the point  $p$  into the upper hull, then the elements lying strictly between  $q$  and  $r$  would be deleted. Then, in an amortized sense, this linear search has cost  $O(1)$ . But if we view the convex hull as a generic structure that supports arbitrary tangent queries, then this linear cost cannot be amortized into a sublinear cost. Our second method can overcome this problem and achieve an amortized cost of  $O(\lg n)$ .

Assume we want to find the left-tangent point  $q = v_\ell$ . As before, we call the explicit binary search in ¶24 to determine the  $i$  satisfying (27). We search for  $q = v_\ell$  using a new kind of binary search. Instead of walking along the upper hull as before, this search begins at the root of the splay tree  $T$ . In general, we are at some node  $u$  of  $T$  where  $u.\text{key}$  is the point  $v_j$  of the upper chain. It must make a 3-way decision, whether  $j = \ell$ , or  $j < \ell$  or  $j > \ell$ . as in Cases (i)-(iii) above. This is an “implicit” binary search because the 3-way decisions are not based

on any explicitly stored keys, but by combining the data stored in  $v_{j-1}, v_j, v_{j+1}$  but *naturally* also the query point  $p$ . Summarizing:

FINDLEFTTANGENTPOINT( $p, T$ ):

▷ *Implicit Binary Search*

1. Initialize  $u$  to the root of  $T$ ; let  $u.\text{key} = v_j$  for some  $j$ .

2. Repeat:

    If  $v_j \geq_x p$ , set  $u \leftarrow u.\text{leftChild}$ .

    Else, we have three possibilities given in ¶25:

        If CASE (i) holds, Return( $v_j$ ).      ◁ *Left tangent found:  $\ell = j$*

        If CASE (ii) holds, set  $u \leftarrow u.\text{rightChild}$ .      ◁ *Go right:  $j > \ell$*

        If CASE (iii) holds, set  $u \leftarrow u.\text{leftChild}$ .      ◁ *Go left:  $j < \ell$*

    Update  $j$  to the index of the vertex in  $u$ .

We now make the inner loop of the implicit search in the FINDLEFTTANGENT algorithm explicit:

2. Repeat:

2.1 Let  $u_0 = u.\text{pred}$  and  $u_1 = u.\text{succ}$ .      ◁ *These may be NULL.*

2.2 If ( $p <_x u$ )

    If ( $u_0 = \text{NULL}$ ), Return( $v_S$ ).      ◁ *South Pole*

    else  $u \leftarrow u.\text{leftChild}$

2.3 elif (( $u_0 \neq \text{NULL}$ ) and LeftTurn( $u_0, u, p$ ) = 1))

$u \leftarrow u.\text{leftChild}$

2.4 elif (( $u_1 \neq \text{NULL}$ ) and LeftTurn( $u, u_1, p$ ) = -1))

$u \leftarrow u.\text{rightChild}$

2.5 else Return( $u$ ).      ◁ *This is correct, even if  $u_0$  or  $u_1$  are NULL.*

To see why the return statement in line 2.5 is correct, first assume  $u_0$  and  $u_1$  are non-null. Then line 2.2 has ensured that  $p <_x u$ , line 2.3 has verified that  $(u_0, u, p)$  is a right-turn and line 2.4 has verified  $(u, u_1, p)$  is a left-turn. These verifications depend on non-degeneracy assumptions. The reader should next verify correctness in case  $u_0$  or  $u_1$  are null.

¶27. **Convex Hull Algorithm and Complexity** We may conclude with the following. Let  $D$  be our data structure for the convex hull  $H$ : so  $D = (T_L, T_R)$  is a pair of splay trees representing the upper and lower chains of  $H$ . We assume (Exercise) that  $D.\text{tangent}(p)$  can be reduced to  $T_L.\text{tangent}(p)$  and  $T_R.\text{tangent}(p)$ , and similarly  $D.\text{insert}(p)$  can be reduced to  $T_L.\text{insert}(p)$  and  $T_R.\text{insert}(p)$ .

### Theorem 12

(i) Using the data structure  $D$  to represent the convex hull  $H$  of a set of points, the requests  $D.\text{insert}(p)$  and  $D.\text{tangent}(p)$  has an amortized cost of  $O(\log |H|)$  time per request.

(ii) From  $D$ , we can produce the cyclic order of points in  $H$  in time  $O(|H|)$ . In particular, this gives an  $O(n \log n)$  algorithm for computing the convex hull of a set of  $n$  points.

This theorem gave us an  $O(n \log n)$  algorithm for computing the convex of a planar point set. Under reasonable models of computation, it can be shown that  $O(n \log n)$  is optimal. There are over a dozen known algorithms for convex hulls in the literature (see Exercise for some).

The complexity is usually expressed as a function of  $n$ , the number of input points. But an interesting concept of “output sensitivity” is to measure the complexity in terms of  $n$  and  $h$ , where  $h$  is the size of the final convex hull. Note  $h$  is a measure of the output size, satisfying  $1 \leq h \leq n$ . For instance, there is a gift-wrap algorithm for convex hull that takes time  $O(hn)$ . So gift-wrap is faster than  $O(n \log n)$  algorithm when  $h = o(\log n)$ . But Kirkpatrick and Seidel [5] gave an output-sensitive algorithm whose complexity is  $O(n \log h)$ .

Our data structure  $D$  for representing convex hulls is only semi-dynamic because we do not support the deletion of points. If we want to allow deletion of points, then points that are inside the current convex hull must be represented in the data structure. Overmars and van Leeuwen [8] designed a data structure for a fully dynamic convex hull that uses  $O(\log^2 n)$  time for insertion and deletion.

There are many applications and generalizations of the convex hull problem. An obvious extension is to ask for the convex hull of a set of points in  $\mathbb{R}^d$  ( $d \geq 3$ ). Or we can replace “points” by balls or other geometric objects in  $\mathbb{R}^d$ . We encourage the student to explore further.

## EXERCISES

**Exercise 2.37:** What is “convex hull” in 1-dimension? What is the upper and lower hull here?

◇

**Exercise 2.38:** You need basic facts about matrices and determinants, and use some physical intuition. Let  $a, b, c \in \mathbb{R}^2$  and consider the following  $3 \times 3$  matrix,

$$M = M(a, b, c) := \begin{bmatrix} a_x & a_y & 1 \\ b_x & b_y & 1 \\ c_x & c_y & 1 \end{bmatrix}.$$

- (i) Show that  $\det M = 0$  iff  $a, b, c$  are collinear (i.e., all 3 points lie on some line). HINT: Wlog, assume  $a \neq b$ . Express  $c$  in terms of  $a$  and  $b$  if  $c$  lies on the line through  $a, b$ ?
- (ii) Show that  $\det M$  is invariant if you translate all the points, i.e.,  $(a, b, c)$  is replaced by  $(a + t, b + t, c + t)$  for any  $t \in \mathbb{R}^2$ .
- (iii) Show that, up to sign,  $\det M$  is the twice the area of the triangle  $(a, b, c)$ . Refer to Figure 20. HINT: from part (ii), you may assume  $a = (0, 0)$  is the origin. First do the case where  $b, c$  both lie in the first quadrant, and then extend the argument, and do some simple calculations!




Figure 20: Area of  $\Delta(a, b, c)$  is half of  $|\det(M)|$

- (iv) Show that  $\det M > 0$  (resp.,  $\det M < 0$ ) iff  $(a, b, c)$  represents a CCW (resp., CW) non-degenerate triangle. HINT: show this is true in a special case. Then extend by translating and rotating the triangle.

(v) What is the relation between  $\text{sign}(\det M)$  and  $\text{LeftTurn}(a, b, c)$ ?

◇

**Exercise 2.39:** Prove the correctness of the  $\text{FINDLEFTTANGENT}(p, T)$  algorithm.

◇

**Exercise 2.40:** For this question, read the relevant parts of §VI.2.3, p.34 (Application to Convex Hull). Please download the latest version of Chap.VI in Brightspace. In your simulation, you will need to show your calculations for each application of the  $\text{LeftTurn}$  predicate.

**NOTE:** you may wish to program the  $\text{LeftTurn}$  predicate. But even if you program the actual algorithms of parts (c) and (d), you must verbally describe how each step works.

Consider the sequence

$$X = (v_0, v_1, \dots, v_4)$$

of points where

$$\begin{aligned} v_0 &= (-8.25, 4.0), & v_1 &= (-3.22, 9.06), & v_2 &= (1.88, 9.86), \\ v_3 &= (8.65, 6.45), & v_4 &= (10.21, 3.82). \end{aligned}$$

Moreover, assume that  $X$  is represented by the splay tree  $T$  shown in Figure 21.

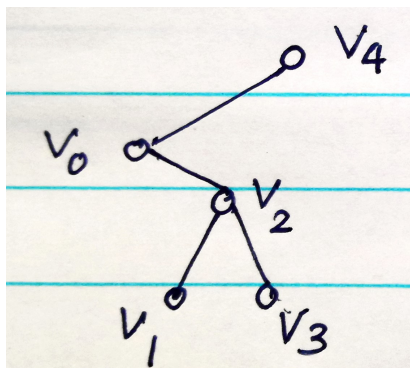


Figure 21: Splay tree for upper chain  $X = (v_0, \dots, v_4)$

- Verify that this is an upper convex chain (see ¶VI.20, page 36). Show your calculations.
- Recall that our splay tree  $T$  maintains successor and predecessor pointers (see ¶VI.23, p.38). Such links are not affected by rotations (see ¶III.27, p.25). Please redraw the splay tree in Figure 21 but show all the successor and predecessor pointers.
- Consider the point

$$p = (8.21, 10.11).$$

Please simulate the method for checking whether or not  $p \in U_X$ . The method is described in ¶VI.24 (p. 38). NOTE: you MUST refer to **actual** numerical data associated with  $p$  and the points  $v_0, \dots, v_4$ .

E.g., Since  $p.x = 8.21$  is greater than  $v_0.x = -8.25$ , we know that  $p$  has a left-tangent point.

- (d) Show how you compute the left-tangent point of  $p = (8.21, 10.11)$  using the Walking Method (§VI.25, p.38). HINT: after part(c), you know the index  $i$  such that

$$v_i <_x p <_x v_{i+1}$$

In fact,  $i = 2$ . Also, we write “ $p <_x q$ ” to mean that  $p.x < q.x$ .

- (e) Repeat part(d), but using the implicit binary search method (§VI.26, p.37).

◇

**Exercise 2.41:** Treatment of Degeneracy. Recall our definition of degeneracy in the previous question.

(i) First define carefully what we mean by the convex hull (and upper hull) of  $X$  in case of degeneracies. You have two choices for this. Also define what we mean by “left-tangent” of  $p$  in case  $p$  lies on a line through two consecutive vertices of the convex hull.

(ii) Modify the `FINDLEFTTANGENT( $p, T$ )` algorithm so that it works correctly for all inputs, degenerate or not. Actually, you need to describe two versions, depending on which way a degenerate convex hull is defined, etc.

◇

**Exercise 2.42:** Let us attend to several details in the convex hull algorithm.

(i) Show how to the non-degeneracy assumptions in the text: recall that we assume consecutive vertices on the convex hull have distinct  $x$ -coordinates, and input set  $X$  is nondegenerate.

(ii) Implement the operation `tangent( $p, H$ )` in terms of `tangent( $p, H_U$ )` and `tangent( $p, H_L$ )`.

(iii) Implement the operation `insert( $p, H_U$ )`.

(iv) When we inserted a new point  $p$ , we split the original tree into  $T_1, T_2, T_3$  and then form a new splay tree rooted at  $p$  with left and right subtrees  $T_1, T_3$ . The cost of forming the new tree is  $O(1)$ . What is the amortized cost of this operation?

◇

**Exercise 2.43:** One of the simplest algorithms for convex hull is the so-called Gift-Wrapping algorithm. Start with  $v_1$  the leftmost point of the convex hull. Now try to find  $v_2, v_3$ , etc in this order. Show that you can find the next point in  $O(n)$  time. Analyze the complexity of this algorithm as a function of  $n$  and  $h$ , where  $1 \leq h \leq n$  is the number of vertices on the convex hull. How does this algorithm compare to  $O(n \log n)$  algorithm?

◇

**Exercise 2.44:** Modified Graham’s algorithm for upper hulls. Let  $S_n = (v_1, \dots, v_n)$  be an input sequence of points in the plane. Assume that the points are sorted by  $x$ -coordinates and satisfy  $v_1 <_x v_2 <_x \dots <_x v_n$ . (Recall “ $a <_x b$ ” means that  $a.x < b.x$ .) Our goal is to compute the upper hull of  $S_n$ . In stage  $i$  ( $i = 1, \dots, n$ ), we have processed the sequence  $S_i$  comprising the first  $i$  points in  $S_n$ . Let  $H_i$  be the upper hull of  $S_i$ . The vertices of  $H_i$  are stored in a push-down stack data structure,  $D$ . Initially,  $D$  contain just the point  $v_1$ .

(a) Describe a subroutine `Update( $v_{i+1}$ )` which modifies  $D$  so that it next represents the upper hull  $H_{i+1}$  upon the addition of the new point  $v_{i+1}$ . HINT: Assume  $D$  contains the sequence of points  $(u_1, \dots, u_h)$  where  $h \geq 1$  and  $u_1$  is at the top of stack, with  $u_1 >_x u_2 >_x \dots >_x u_h$ . For any point  $p$ , let  $LT(p)$  denote the predicate `LeftTurn( $p, u_1, u_2$ )`. If  $h = 1$ ,  $LT(p)$  is defined to be `true`. Implement `Update( $v_{i+1}$ )` using the predicate  $LT(p)$  and the (ordinary) operations of push and pop of  $D$ .

(b) Using part (a), describe an algorithm for computing the convex hull of a set of  $n$  points. Analyze the complexity of your algorithm.



REMARK: The amortized analysis of  $S.Update(p)$  was essentially described in an Exercise (Section 1, this Chapter). Graham's original idea is to sort the vertices by their angular angle about some point  $p_0$  in the interior of the convex hull. We must implement this with care, so as to avoid the actual computation of angles (such computation would be inexact and have robustness problems).  $\diamond$

**Exercise 2.45:** The divide-and-conquer for convex hull is from Shamos: divide the set into two sets  $X_L, X_R$ , each of size about  $n/2$  and the two sets are separated by some vertical line  $L$ . Recursively compute their convex hulls  $H_L, H_R$ . What kind of operation(s) will allow you to compute  $CH(X)$  from  $H_L$  and  $H_R$ ? Show that these operations can be implemented in  $O(n)$  time.  $\diamond$

END EXERCISES

### §3. Fibonacci Heap Data Structure

#### ¶28. Motivation: Complexity of Prim's MST Algorithm.

In ¶V.39 (page 62) we saw that Prim's algorithm on a connected bigraph  $G = (V, E; C)$  works as follows: it maintains a set  $S \subseteq V$  such that we already know the MST of  $G|S$ . We can maintain a priority queue  $Q$  on the set complementary set  $U := V \setminus S$ , where each  $u \in U$  is stored in  $Q$  with priority

$$\min \{C(s-u) : s \in S\}.$$

Then a simple version of Prim's algorithm is as follows

```

T ← Prim(G)
INPUT: connected costed bigraph G = (V, E; C) where V = [1..n]
OUTPUT: a MST T ⊆ E of G represented by array d[1..n] where
    u-v ∈ T iff d[u] = v or d[v] = u.
▷ Initialize S, d and priority:
    S ← ∅
    For all v ∈ V, d[v] ← 1.    < 1 is arbitrary
    For all v ∈ V, priority[v] ← C(1, v).    < priority[v] = 0 iff v = 1
▷ Initialize priority Q:
    for v ∈ V, Q.insert(v, priority[v])
▷ Main Loop:
    While (Q is non-empty)
        v ← Q.deleteMin()
        S ← S ∪ {v}
        For all (u adjacent to v)
            If (u ∉ S and priority[u] > C(v-u))
                priority[u] ← C(v-u)    - (**)
                d[u] ← v
    Return d

```

The line (\*\*) is problematic: we have changed the priority of  $u \in Q$ . This is not a normal operation of a priority queue! But it can be implemented by using two normal operations of a

priority queue: (1) First delete  $u$  from  $Q$ , (2) Reinsert  $u$  into  $Q$  with the new priority. If we do this, we see that the complexity of Prim's algorithm would be

$$T(n, m) = O((m + n) \log n) \quad (28)$$

since each insert and delete in the queue is  $O(\log n)$ . We now introduce a priority queue in which **(\*\*)** can be implemented in amortized  $O(1)$  time.

### ¶29. The mergeable queues ADT.

The **Fibonacci heap data structure** invented by Fredman and Tarjan (1987) gives an efficient implementation of the mergeable queues abstract data type (ADT).

The mergeable queues data structure stores several sets of items, and each set is stored in its own priority queue. As usual, the with priority given by the keys in item. The supported operations are as follows: assume  $Q, Q_i$  are queues,  $x$  is an item and  $k$  a key.

1.	<code>makeQueue()</code> $\rightarrow Q$	returns an empty queue $Q$
2.	<code>insert</code> ( $x, k, Q$ )	insert $x$ with priority $k$ into $Q$
3.	<code>deleteMin</code> ( $Q$ ) $\rightarrow x$	the minimum item $x$ in $Q$ is deleted
4.	<code>decreaseKey</code> ( $x, k, Q$ )	the priority in $x$ is replaced by a smaller priority $k$
5.	<code>union</code> ( $Q_1, Q_2$ )	form the union of the two queues

The first three operations of mergeable queues are just those of the priority queue ADT (§III.2). The two new operations are the ability to decrease the key of an item, and the ability to merge two queues. The full power of mergeable queues are not needed in applications with only one mergeable queue (so the **union** operation is irrelevant). Such is the case in the algorithms of Prim and Dijkstra in the next section.


We briefly add some further clarifications on these operations. First of all, none of the ADT operations require the items in the queue to be sorted. In the union of  $Q_1, Q_2$ , the items in  $Q_2$  are first moved into queue  $Q_1$ , then queue  $Q_2$  is destroyed. Union is sometimes known as the **meld** operation. Thus, the number of queues can increase or decrease over the lifetime of the data structure. The operation `deleteMin`( $Q$ ) returns a minimum item in  $Q$ , and this item is deleted from  $Q$ . This operation is undefined if  $Q$  is empty. The operation `decreaseKey`( $x, k, Q$ ) is undefined if the key  $k$  is greater than the key in  $x$ .

There are additional useful operations that should probably be provided in practice, but is omitted in the above list for the sake of economy: deleting an item, making a singleton queue, and getting the minimum item without deleting it. These can be defined as follows:

<code>delete</code> (Item $x$ , Queue $Q$ )	$\equiv$	<code>decreaseKey</code> ( $x, -\infty, Q$ ); <code>deleteMin</code> ( $Q$ ).
<code>makeQueue</code> (Item $x$ ) $\rightarrow Q$	$\equiv$	<code>makeQueue</code> () $\rightarrow Q$ ; <code>insert</code> ( $x, Q$ ).
<code>min</code> (Queue $Q$ ) $\rightarrow x$	$\equiv$	<code>deleteMin</code> ( $Q$ ) $\rightarrow x$ ; <code>insert</code> ( $x, Q$ ).

¶30. **The Fibonacci heap data structure.** We now consider the implementation details. Each mergeable queue  $Q$  is implemented by a **Fibonacci heap**  $H$ , where  $H$  is a collection of trees  $T_1, \dots, T_m$  with these properties:

- Each tree  $T_i$  satisfies the **min-heap property**: this means that each node has a key that is greater than or equal to the key of its parent. In particular, the root of  $T_i$  has the minimum item in  $T_i$ .
- The roots of these trees are kept in a doubly-linked list, called the **root-list** of  $H$ .
- There are two fields  $H.min$ ,  $H.n$  associated with  $H$ . The field  $H.min$  points to the node with a minimum key, and  $H.n$  is the number of items in  $H$ .
- For each node  $x$  in a tree  $T_i$ , we have four pointers that point to (i) the parent of  $x$ , (ii) one of its children, and (iii) two of its siblings. The sibling pointers are arranged so that all the children of  $x$  appears in a circular doubly-linked list called the **child-list** of  $x$ . If  $y$  is a child of  $x$ , the **sibling-list** of  $y$  is the child-list of  $x$ . Also, we keep track of  $x.degree$  (the number of children of  $x$ ) and  $x.mark$  (a Boolean value to be explained).

Figure 22: A Fibonacci heap  $H = (T_0, \dots, T_3)$ :  $T_0$  in detail

This is illustrated in Figure 22. One of the trees  $T_0$  is shown in detail: the root  $a$  of  $T_0$  has 3 children  $b, c$  and  $d$  and they each point to  $a$ ; on the other hand,  $a$  points only to  $b$ . There are two non-trivial sibling lists:  $(b, c, d)$  and  $(f, g)$ .

**§3.1. Linking, cutting and marking.** We describe some elementary operations used in maintaining Fibonacci heaps.

(a) If  $x, y$  are two roots such that the item in  $x$  is not less than the item in  $y$ , then we can **link**  $x$  and  $y$ ; this simply makes  $y$  the parent of  $x$ . The appropriate fields and structures are then updated:  $x$  is deleted from the root-list, and then inserted into the child-list of  $y$ , the degree of  $y$  incremented, etc. This operation costs  $O(1)$ .

(b) The converse to linking is **cutting**. If  $x$  is a non-root in a Fibonacci heap  $H$  then we can perform  $Cut(x, H)$ : this basically removes  $x$  from the child-list of its parent and inserts  $x$  into the root-list of  $H$ . The appropriate data variables are updated. E.g., the degree of the parent of  $x$  is decremented. Again, this operation costs  $O(1)$ .

(c) We say  $x$  is **marked** if  $x.mark = \text{true}$ , and **unmarked** otherwise. Initially,  $x$  is unmarked. Our rules ensure that a root is always unmarked. We mark  $x$  if  $x$  is not a root and  $x$  loses a child for the first time. If  $x$  loses a child for the second time, it is cut and put in the root-list as a root (and therefore becomes unmarked again).

To do amortized analysis, we define a potential function. The **potential** of a Fibonacci heap  $H$  is defined as

$$\Phi(H) := t(H) + 2 \cdot m(H) \quad (29)$$

where  $t(H)$  is the number of trees in  $H$  and  $m(H)$  is the number marked items in  $H$ . The potential of a collection of Fibonacci heaps is just the sum of the potentials of the individual heaps.

One more definition: let  $D(n)$  denote the maximum degree of a node in a Fibonacci heap with  $n$  items. We will show later that  $D(n) \leq 2 \lg n$ .

**Remark:** The reader may observe how “low-tech” this data structure appears – along with the humble array structure, linked-lists is among the simplest data structures. Yet we intend to achieve the best known overall performance for mergeable queues with Fibonacci heaps. This should be viewed as a testimony to the power of amortization.

### §3.1. Fibonacci Heap Algorithms

We now implement the mergeable queue operations. Our goal is to achieve an amortized cost of  $O(1)$  for each operation except for **deleteMin**, which will have logarithmic amortized cost.

Recall that for each operation  $p$ , we have a **cost**  $\text{COST}(p)$  which will be mostly self-evident in the following description. We must define a **charge**  $\text{CHARGE}(p)$ . The **credit** is thereby determined:  $\text{CREDIT}(p) := \text{CHARGE}(p) - \text{COST}(p)$ . This charging scheme will achieve the stated goal in the previous paragraph:  $\Theta(1)$  charges for all the non-deletion operations, and  $\Theta(\log n)$  for the **deleteMin** operation. Finally, we verify the credit-potential invariant equation (5) for each operation.

**makeQueue()**: we create an empty root-list. The cost is 1, the charge is 1, so credit is 0, and finally  $\Delta\Phi = 0$ . The credit-potential invariant holds trivially.

**insert( $H, x$ )**: we create a new tree  $T$  containing only  $x$  and insert  $T$  into the root-list of  $H$ . Update  $H.\text{min}$  and  $H.n$  in the obvious way. Let us check the credit-potential invariant:

$$\text{COST} = 1, \quad \text{CHARGE} = 2, \quad \text{CREDIT} = 1, \quad \Delta\Phi = 1.$$

**union( $H_1, H_2$ )**: concatenate the two root-lists and call it  $H_1$ . It is easy to update  $H_1.\text{min}$  and  $H_1.n$ . Checking the credit-potential invariant:

$$\text{COST} = 1, \quad \text{CHARGE} = 1, \quad \text{CREDIT} = 0, \quad \Delta\Phi = 0.$$

**deleteMin( $H$ )**: we remove  $H.\text{min}$  from the root-list, and the child-list of  $H.\text{min}$  can now be regarded as the root-list of another Fibonacci heap. These two circular lists can be concatenated in constant time into a new root-list for  $H$ . If  $t_0$  is the old value of  $t(H)$ , the new value of  $t(H)$  is at most  $t_0 + D(n)$ . Next we need to find the new value of  $H.\text{min}$ . Unfortunately, we do not know the new minimum item of  $H$ . There is no choice but to scan the new root-list of  $H$ . While scanning, we might as well<sup>7</sup> spend some extra effort to save future work. This is a process called **consolidation** which is explained next.

<sup>7</sup>OK, we may be lazy but not stupid.

¶32. **Consolidation.** In this process, we are given a root-list of length  $L$  ( $L \leq t_0 + D(n)$  above). We must visit every member in the root-list, and at the same time do repeated linkings until there is at most one root of each degree. We want to do this in  $O(L)$  time. By assumption, each root has degree at most  $D(n)$ .

The basic method is that, for each root  $x$ , we try to find another root  $y$  of the same degree and link the two. So we create a ‘new’ root of degree  $k + 1$  from two roots of degree  $k$ . If we detect another root of degree  $k + 1$ , we link these two to create another ‘new’ root of degree  $k + 2$ , and so on. The way to detect the presence of another root of the same degree is by indexing into an array  $A[1..D(n)]$  of pointers. Initialize all entries of the array to `nil`. Then we scan each item  $x$  in the root-list. If  $k = x.\text{degree}$  then we try to make  $A[k]$  point to  $x$  (we say  $x$  is “inserted” into the array  $A$  when this happens). But we can only insert  $x$  if  $A[k] = \text{nil}$ ; otherwise, suppose  $A[k]$  points to some  $y$ . We then link  $x$  to  $y$  if  $x.\text{key} \geq y.\text{key}$ , and otherwise link  $y$  to  $x$ . Wlog, assume we linked  $x$  to  $y$ . Now,  $y$  has degree  $k + 1$  and we recursively try to “insert  $y$ ”, i.e., make  $A[k + 1]$  point to  $y$ , etc. So each failed insertion leads to a linking, and there are at most  $L$  linking operations. Since each linking removes one root, there are at most  $L$  linkings in all. (This may not be obvious if we see this the wrong way!) Thus the total cost of consolidation is  $O(L)$ .

Returning to `deleteMin`, let us check its credit-potential invariant.

$$\begin{aligned}\text{COST} &\leq 1 + t_0 + D(n), & \text{CHARGE} &= 2 + 2D(n), \\ \text{CREDIT} &\geq 1 + D(n) - t_0, \\ \Delta\Phi &\leq 1 + D(n) - t_0.\end{aligned}$$

We need to explain our bound for  $\Delta\Phi$ . Let  $t_0, m_0$  refer to the values of  $t(H)$  and  $m(H)$  before this `deleteMin` operation. If  $\Phi_0, \Phi_1$  are (respectively) the potentials before and after this operation, then  $\Phi_0 = t_0 + 2m_0$  and  $\Phi_1 \leq 1 + D(n) + 2m_0$ . To see this bound on  $\Phi_1$ , note that no node can have degree more than  $D(n)$  (by definition of  $D(n)$ ) and hence there are at most  $1 + D(n)$  trees after consolidation. Moreover, there are at most  $m_0$  marked nodes after consolidation. Therefore  $\Delta\Phi = \Phi_1 - \Phi_0 \leq 1 + D(n) - t_0$ , as desired.

`decreaseKey`( $x, k, H$ ): this is the remaining operation and we will exploit the marking of items in a crucial way. The idea of the `decreaseKey` algorithm is basically to maintain heap property, namely that the parent should not have larger keys than its children. But by decreasing the key of  $x$  to  $k$ , we may cause a violation of this heap property. To remove the violation, we can cut the link of  $x$  to its parents. With this in mind, consider the details of this operation: First, we decrease the key of  $x$  to  $k$  (first checking that  $k \leq x.\text{key}$ ). If  $x$  is a root, we are done. Otherwise, let  $y$  be the parent of  $x$ . If  $k \geq y.\text{key}$ , we are done. Otherwise, we cut  $x$ . Since  $x$  is now in the root list, we need to update  $H.\text{min}$ , etc. If  $x$  was marked, it is now unmarked. We must now treat  $y$ : if  $y$  is a root, we are done. Otherwise, if  $y$  was unmarked, we mark  $y$  and we are done. The last case is when  $y$  is marked. This is the most interesting case. But before treating it, we note that the other cases so far all have constant cost, and have constant increase in potential. Thus, they can be directly charged.

To treat the case where  $y$  is marked, we recall that this means  $y$  had previously lost a child, and since it just lost a second child, viz.  $x$ , our rules require  $y$  to be cut. In fact, the cut may cascade into more cuts and the cost may be unbounded. Here is code fragment to “recursively cut”  $y$ :

```

RECURSIVECUT( $y, H$ ):
  If ( $y.mark = \text{false}$  and  $y \neq \text{root}$ )
     $y.mark \leftarrow \text{true}$ ;
  elif ( $y \neq \text{root}$ )  $\triangleleft$  So  $y$  is marked
     $y.mark \leftarrow \text{false}$  and  $\text{Cut}(y, H)$ ;
     $\text{RecursiveCut}(y.parent, H)$ .

```

Note that if  $c \geq 1$  is the number of cuts, then  $t(H)$  is increased by  $c$ , but  $m(H)$  is decreased by  $c - 1$  or  $c$  (the latter iff  $x$  was marked). The worst case is when  $m(H)$  decreases by  $c - 1$ , giving us  $\Delta\Phi \leq c - 2(c - 1) = 2 - c$ . By charging 2 units for this operation,

$$\text{COST} = 1 + c, \quad \text{CHARGE} = 3, \quad \text{CREDIT} = 2 - c, \quad \Delta\Phi \leq 2 - c$$

and so the credit-potential invariant is verified.

In summary, we have achieved our goal of charging  $O(1)$  units to every operation except for **deleteMin** which is charged  $2 + 2D(n)$ . We next turn to bounding  $D(n)$ . We remark on an unusual feature in our marking scheme: each non-root  $y$  can suffer the loss of at most two children before  $y$  itself is cut and made a root. But if  $y$  is a root, it may lose an unlimited number of children.




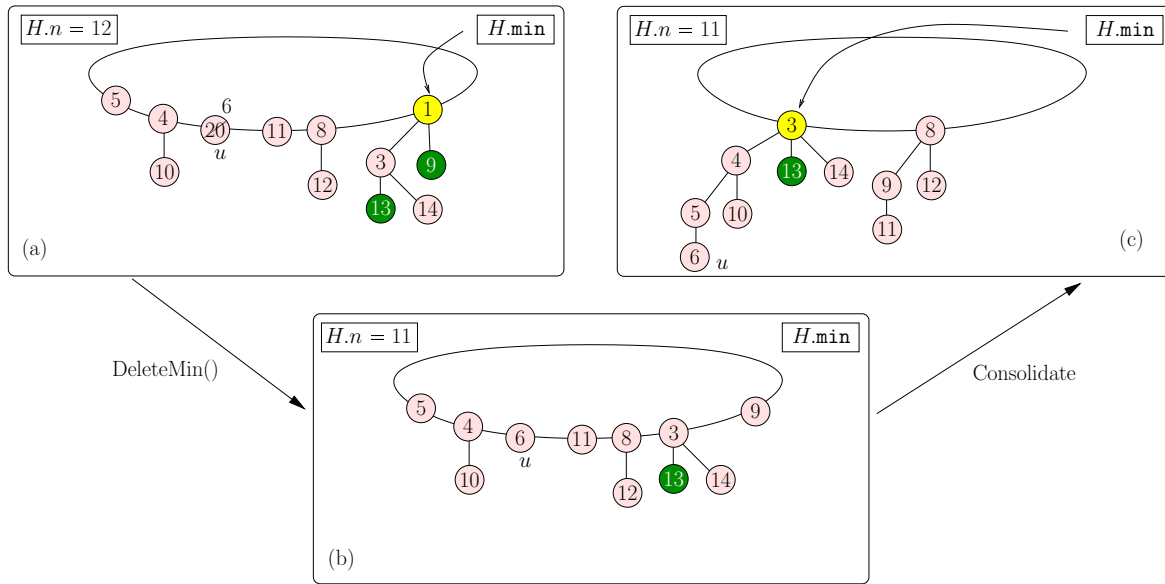
Figure 23: A Fibonacci heap

**¶33. Example.** Consider the Fibonacci heap in Figure 23, with keys written inside each node. Four of the nodes (keys 8, 9, 11, 13) are “marked”. We do not show sibling pointers to avoid clutter (e.g., 8, 3, 9 are siblings).

- Consider the operation **decreaseKey**( $u, 6$ ): the results in the data structure in Figure 24(a): we have to cut the node  $u$ , and this leads to a recursive cut of the marked keys 11 and 8. Thus, we add three new roots to the root-list: node  $u$  (6, formerly 20), node 11 and node 8.
- Consider the further operation of **deleteMin**( $\cdot$ ). The key 1 is deleted, but its children 3 and 9 are now added to the root-list. We now do consolidation of the list of roots

5, 4, 6, 11, 8, 3, 9.

It is not hard to see that we end up with the two roots 3 and 8. Moreover,  $H.min \leftarrow 3$ .

Figure 24: (a) After `decreaseKey(u, 6)` (c) After `deleteMin()`

¶34. Degree Bound Let us now show that  $D(n) = O(\log n)$ .

Recall the  $i$ th Fibonacci number  $i = 0, 1, 2, \dots$  is defined by  $F_i = i$  if  $i = 0, 1$  and  $F_i = F_{i-1} + F_{i-2}$  for  $i \geq 2$ . Thus the sequence of Fibonacci numbers starts out as

$$0, 1, 1, 2, 3, 5, 8, \dots$$

We will use two simple facts:

**Lemma 13**

(a)  $F_i = 1 + \sum_{j=1}^{i-2} F_j$  for  $i \geq 1$ .

(b)  $F_{j+2} \geq \phi^j$  for  $j \geq 0$ , where  $\phi = (1 + \sqrt{5})/2 > 1.618$ .

*Proof.* Part(a) follows easily by induction, or better still, by “unrolling” the recurrence for  $F_i$ . For part(b), we observe that  $\phi$  is a solution to the equation  $x^2 - x - 1 = 0$  so  $\phi^2 = 1 + \phi$ . Clearly  $F_2 = 1 \geq \phi^0$  and  $F_3 = 2 \geq \phi^1$ . Inductively,

$$F_{j+2} = F_{j+1} + F_j \geq \phi^{j-1} + \phi^{j-2} = \phi^{j-2}(\phi + 1) = \phi^j.$$

**Q.E.D.**

Let  $x$  be a node in a Fibonacci heap with  $n$  items, and let

$$y_1, y_2, \dots, y_d \tag{30}$$

be the children of  $x$ , given in the order in which they are linked to  $x$ . So  $x.\text{degree} = d$  and  $y_1$  is the earliest child (among  $y_1, \dots, y_d$ ) to be linked to  $x$ .

**Lemma 14**

$$y_i.\text{degree} \geq \begin{cases} 0 & \text{if } i = 1, \\ i - 2 & \text{if } i \geq 2. \end{cases}$$

*Proof.* This is clearly true for  $i = 1$ . For  $i \geq 2$ , note that when  $y_i$  was linked to  $x$ , the degree of  $x$  is at least  $i - 1$  (since at least  $y_1, \dots, y_{i-1}$  are children of  $x$  at the moment of linking). Hence, the degree of  $y_i$  at that moment is at least  $i - 1$ . This is because we only do linking during consolidation, and we link two roots only when they have the same degree. But we allow  $y_i$  to lose at most one child before cutting  $y_i$ . Since  $y_i$  is not (yet) cut from  $x$ , the degree of  $y_i$  is at least  $i - 2$ . **Q.E.D.**

**Lemma 15** Let  $\text{SIZE}(x)$  denote the number of nodes in the subtree rooted at  $x$  and  $d$  the degree of  $x$ . Then

$$\text{SIZE}(x) \geq F_{2+d}, \quad d \geq 0.$$

*Proof.* This is seen by induction on  $d$ . The result is true when  $d = 0$  or  $d = 1$ . If  $d \geq 2$ , let  $y_1, \dots, y_d$  be the children of  $x$  as in (30). Then

$$\begin{aligned} \text{SIZE}(x) &= 1 + \sum_{i=1}^d \text{SIZE}(y_i) \\ &\geq 1 + \sum_{i=1}^d F_{2+y_i.\text{degree}} && \text{(by induction)} \\ &\geq 2 + \sum_{i=2}^d F_i && \text{(by last lemma, and since } d \geq 2) \\ &= 1 + \sum_{i=1}^d F_i = F_{d+2}. && \text{(by Lemma 13(a))} \end{aligned}$$

**Q.E.D.**

It follows that if  $x$  has degree  $d$ , then

$$n \geq \text{SIZE}(x) \geq F_{d+2} \geq \phi^d.$$

where the last inequality is from Lemma 13(b). Taking logarithms, we immediately obtain:

**Lemma 16**

$$D(n) \leq \log_{\phi}(n).$$

This completes our analysis of Fibonacci heaps.

*So that is where the name “Fibonacci” arises!*

Historical remarks. Prior to Fibonacci heaps, **binomial heaps** from Vuillemin (1978) were the best data structure for the mergeable queue ADT. Indeed, Fibonacci heaps has vestiges of binomial heaps as can be seen when we do consolidation. There is some interest to improve the amortized complexity of Fibonacci heaps to worst case complexity bounds. This was achieved by Brodal in 1996.



**Exercise 3.1:** Suppose you insert the following sequence of items (represented by their keys) into an initially empty Fibonacci heap: 2, 3, 5, 7, 11, 13, 17, 19, 23, 29, 31. Then you do a deleteMin. Next, you insert the following sequence of items into another Fibonacci heap: 37, 41, 43, 47, 53, 57, 59. Then you do a deleteMin. (a) Draw the resulting data structure. (b)

**Exercise 3.2:** On marking items.

- (a) Which of the mergeable queue operations calls for marking?  
 (b) What is the rationale for marking items?

**Exercise 3.3:** “Second chance Fibonacci heap.” Suppose that instead of cutting a node just after it loses a second child, we only cut a node just after it loses a third child. Carry out the analysis as before.

**Exercise 3.4:** Suppose  $x > 1$  and  $x^3 = x^2 + 1$ , and  $\phi$  is the Golden ratio. Prove that  $x^2 > \phi > x$ .

**Exercise 3.5:** Generalize the above Exercises to the  $c$ -chance Fibonacci heap, for any integer  $c > 1$ . The constants  $\phi$  for  $c = 1, 2, 3, 7$  are plotted in Figure 25.




Figure 25: Constants for  $c = 1, 2, 3, 7$

**Exercise 3.6:**

- (a) Determine  $\hat{\phi}$ , the other root of the equation  $x^2 - x - 1 = 0$ . Numerically compute  $\hat{\phi}$  to 5 decimal places.  
 (b) Determine  $F_i$  exactly in terms of  $\phi$  and  $\hat{\phi}$ . HINT:  $F_i = A\phi^i + B\hat{\phi}^i$  for constants  $A, B$ .  
 (b) What is the influence of the  $\hat{\phi}$ -term on the relative magnitude of  $F_i$ ?

END EXERCISES

### §3.2. Pointer Model of Computation

There is an aesthetically displeasing feature in our consolidation algorithm: its use of array indexing does not seem to conform to the style used in the other operations. Intuitively the reason is that, unlike the other operations, indexing does not fit within the “pointer model” of computation. In this section, we will give the pointer-based solution to consolidation. We also provide an elegant formalization of the pointer model, going back to Schönhage. This model can be used for a formal theory of computability and complexity. It has many advantages over the standard Turing machines model.

¶35. **Pointer based consolidation.** We outline a purely pointer-based method for consolidation. We rely on the reader’s understanding of pointers as found in conventional programming languages such as C or C++.

For each Fibonacci heap  $H$ , suppose we maintain a doubly-linked list of nodes

$$(R_0, R_1, \dots, R_k)$$

where  $k \leq D(n)$  is the maximum degree of any node seen so far in  $H$ . We call this list the **degree register** because each node in  $H$  of degree  $i$  is required to have a pointer to node  $R_i$  (i.e., that node is “registered” at  $R_i$ ). Here  $k$  is the largest degree of a node that has been seen so far. Note that when we link  $x$  to  $y$  then the degree of  $y$  increments by one and when we cut  $x$ , then the degree of  $x$ .parent decrements by one, and these are the only possibilities. If item  $x$  has its degree changed from  $i$  to  $i \pm 1$  then we can re-register  $x$  by pointing it to  $R_{i \pm 1}$  in constant time. Occasionally, we have to extend the length of the register by appending a new node  $R_{k+1}$  to the doubly-linked list (when some node attains a degree  $k + 1$  that is larger than any seen so far). It is thus easy to maintain this degree register.

Now suppose we must consolidate a root list  $J$ . By going through the items in  $J$ , we can create (with the help of the degree register) a list of lists

$$(L_0, L_1, \dots, L_k)$$

where list  $L_i$  comprises the roots of degree  $i$  in  $J$ . This takes  $\mathcal{O}(D(n) + t)$  operations if  $J$  has  $t$  elements. It is now easy to consolidate the lists  $L_0, \dots, L_k$  into one list in which no two trees have the same degree, using  $\mathcal{O}(t)$  time. The cost of this procedure is  $\mathcal{O}(D(n) + t)$ , as in the solution that uses array indexing.

But we can take this idea further: we can reinterpret the root list of  $H$  as a sublist of the degree register, and we now require that the degrees of roots in the root-list must be distinct. This ensure that the length of the root list is  $\leq D(n)$ , not  $n$  in the worst case. Whenever we add a new root of degree  $i$  to the root-list, we must check if there is already a root of degree  $i$ . If so, we link these two roots, and recursively register the result as a root of degree  $i + 1$ . The worst case cutting a node is thus  $\Theta(\lg n)$  because of this recursion. Nevertheless, the amortized cost is  $\mathcal{O}(1)$ , using the same ideas as the Counter Example analysis (see §1). We get several benefits in this approach:

- We have described a purely pointer-based implementation of Fibonacci heaps.
- We have eliminated the explicit consolidation process.
- The operation `deleteMin( $H$ )` amounts to a simple search through the degree register; hence its worst case time is no longer  $\mathcal{O}(n)$  but  $\mathcal{O}(\log n)$ .

But there is a down side: when we merge two heaps, the cost is no longer  $\mathcal{O}(1)$  but  $\mathcal{O}(\log n)$  because of the need to merge their degree registers. But many applications of Fibonacci heaps (e.g., in the implementation of Prim's Algorithm for MST) do not need to merge heaps. In this case, as a degree register is no burden.

**¶36. The Pointer Computational Model.** We now give a formal model of the pointer model. A **pointer program**  $\Pi$  consists of a finite sequence of instructions that operate on an implicit potentially infinite digraph  $G$ . All program variables in  $\Pi$  are of type **POINTER**, but we also manipulate integer values via these pointers. Each pointer points to some node in  $G$ . Each node  $N$  in  $G$  has four components:

(integer-value, 0-pointer, 1-pointer, 2-pointer).

These are accessed as  $P.\text{Val}$ ,  $P.0$ ,  $P.1$  and  $P.2$  where  $P$  is any pointer variable that points to  $N$ . There is a special node  $N_0 \in G$  and this is pointed to by the **nil** pointer. By definition,  $\text{nil}.\text{Val} = 0$  and  $\text{nil}.i = \text{nil}$  for  $i = 0, 1, 2$ . Note that with 3 pointers, it is easy to model binary trees.

**¶37. Pointer Expressions.** In general, we can specify a node by a **pointer expression**,  $\langle \text{pointer-expr} \rangle$ , which is either the constant **nil**, the **NEW()** operator, or has the form  $P.w$  where  $P$  is a pointer variable and  $w \in \{0, 1, 2\}^*$ . The string  $w$  is also called a **path**. Examples of pointer expressions:

$\text{nil}, \text{NEW}(), P, P.0, Q.1, P.2, Q.1202, P.2120120$

where  $P, Q$  are pointer variables. The **NEW()** operator (with no arguments) returns a pointer to a “spanking new node”  $N$  where  $N.0 = N.1 = N.2 = \text{nil}$  and  $N.\text{Val} = 1$ . The only way to access a node or its components is via such pointer expressions.

The integer values stored in nodes are unbounded and one can perform the four arithmetic operations; compare two integers; and assign to an integer variable from any integer expression (see below).

We can compare two pointers for equality or inequality, and can assign to a pointer variable from another pointer variable or the constant **nil** or the function **NEW()**. Assignment to a **nil** pointer has no effect. Note that we are not allowed to do pointer arithmetic or to compare them for the “less than” relation.

The assignment of pointers can be explained with an example:

$P.0121 \leftarrow Q.20002$

If  $N$  is the node referenced by  $P.012$  and  $N'$  is the node referenced by  $Q.20002$ , then we are setting  $N.1$  to point to  $N'$ . If  $N$  is the **nil** node, then this assignment has no effect.

Naturally, we use the result of a comparison to decide whether or not to branch to a labeled instruction. Assume some convention for input and output. For instance, we may have two special pointers  $P_{in}$  and  $P_{out}$  that point (respectively) to the input and output of the program.

To summarize: a pointer program is a sequence of instructions (with an optional label) of the following types.

- Value Assignment:  $\langle \text{pointer-expr.Val} \rangle \leftarrow \langle \text{integer-expr} \rangle$ ;
- Pointer Assignment:  $\langle \text{path-expr} \rangle \leftarrow \langle \text{pointer-expr} \rangle$ ;
- Pointer Comparison: If  $\langle \text{pointer-expr} \rangle = \langle \text{pointer-expr} \rangle$  then goto  $\langle \text{label} \rangle$ ;
- Value Comparison: If  $\langle \text{integer-expr} \rangle \geq 0$  then goto  $\langle \text{label} \rangle$ ;
- Halt

Integer expressions denote integer values. For instance

$$(74 * P.000) - (Q.21 + P)$$

where  $P, Q$  are pointer variables. Here,  $P.000, Q.21, P$  denotes the values stored at the corresponding nodes. Thus, an integer expression  $\langle \text{integer-expr} \rangle$  is either

- Base Case: any literal integer constant (e.g., 0, 1, 74, -199), a  $\langle \text{pointer-expr} \rangle$  (e.g.,  $P.012, Q, \text{nil}$ ); or
- Recursively:

$$(\langle \text{integer-expr} \rangle \langle op \rangle \langle \text{integer-expr} \rangle)$$

where  $\langle op \rangle$  is one of the four arithmetic operations. Recall that  $\text{nil.Val} = 0$ . Some details about the semantics of the model may be left unspecified for now. For instance, if we divide by 0, the program may be assumed to halt instantly.

For a simple complexity model, we may assume each of the above operations take unit time regardless of the pointers or the size of the integers involved. Likewise, the space usage can be simplified to just counting the number of nodes used.

One could embellish it with higher level constructs such as While-loops. Or, we could impoverish it by restricting the integer values to Boolean values (to obtain a better accounting of the bit-complexity of such programs). In general, we could have pointer models in which the value of a node  $P.\text{Val}$  comes from any domain. For instance, to model computation over a ring  $R$ , we let  $P.\text{Val}$  be an element of  $R$ . We might wish to have an inverse to  $\text{NEW}()$ , to delete a node.

**¶38. List reversal example.** Consider a pointer program to reverse a singly-linked list of numbers (we only use 0-pointer of each node to point to the next node). Our program uses the pointer variables  $P, Q, R$  and we write  $P \leftarrow Q \leftarrow R$  to mean the sequential assignments “ $P \leftarrow Q$ ;  $Q \leftarrow R$ ”.

```

REVERSELIST:
Input:  $P_{in}$ , pointer to a linked list.
Output:  $P_{out}$ , pointer to the reversal of  $P_{in}$ .
   $P \leftarrow \text{nil}; Q \leftarrow P_{in};$ 
  If  $Q = \text{nil}$  then goto E;
   $R \leftarrow Q.0 \leftarrow P;$ 
L:  If  $R = \text{nil}$  then goto E;
T:   $P \leftarrow Q \leftarrow R \leftarrow Q.0 \leftarrow P;$ 
    goto L;
E:   $P_{out} \leftarrow Q.$ 

```

This program is easy to grasp once the invariant preceding Line T is understood (see Figure 26 and Exercise).




Figure 26: List Reversal Algorithm: the transformation at Line T.

**Remark:** This model may be more convenient than Turing machines to use as a common basis for discussing complexity theory issues. The main reservation comes from our unit cost for unbounded integers operations. In that case we can either require that all integers be bounded, or else charge a suitable cost  $M(n)$  for multiplying  $n$ -bit integers, etc, reflecting the Turing machine cost. Of course, the use of pointers is still non-elementary from the viewpoint of Turing machines, but this is precisely the convenience we gain.

#### EXERCISES

**Exercise 3.7:** State the invariant before line T in the pointer reversal program; then proving the program correct.  $\diamond$

**Exercise 3.8:** Write the pointer program for the consolidation.  $\diamond$

**Exercise 3.9:** Implement in detail all the Fibonacci heap algorithms using our pointer model,  $\diamond$

**Exercise 3.10:** Write a sorting program and a matrix multiplication program in this model. What is the time complexity of your algorithms?  $\diamond$

#### END EXERCISES

## §4. Application to Algorithms of Prim and Dijkstra

An original application of Fibonacci heaps is in computing minimum spanning trees (MST). We saw in §V.4 that there are several known algorithms for MST. Here, we are interested in efficient implementation of Prim's algorithm. Basically, we will reduce Prim's algorithm is a sequence of requests on a mergeable queue ADT. It turns out that Prim's algorithm has the same algorithmic framework as another famous algorithm, namely, Dijkstra's shortest path<sup>8</sup> algorithm. Hence we shall give them a unified treatment in this section.

*Sorting Student:  
hey, I have seen  
this before!*

The input for MST is a connected bigraph  $G = (V, E; C)$  with edge cost function  $C : E \rightarrow \mathbb{R}$ . The fact that  $G$  is connected is not essential, but simplifies our discussion.

The shortest path problem solved by Dijkstra is for an input digraph  $G = (V, E; C, s)$  where  $C : E \rightarrow \mathbb{R}_{\geq 0}$  and  $s \in V$ . This is slightly different from the input for MST: we now treat a digraph instead of a bigraph, and edge costs must be non-negative. Also, we are given a vertex  $s \in V$  called the **source**. The goal is to compute the min-cost paths from  $s$  to each  $v \in V$ , where the cost of a path  $v_0 - v_1 - \dots - v_k$  is just  $\sum_{i=1}^k C[v_{i-1}, v_i]$ . The fact that edge costs for non-negative is critical for Dijkstra's algorithm, but not essential for MST. Basically, if we add a fixed large positive value to every edge cost of the MST, we get an equivalent MST instance with positive costs. Hence, we shall assume all costs are non-negative in this section.

*Edges have  
non-negative costs*

**§39. The Generic Framework.** Both algorithms can be viewed as growing a set  $S \subseteq V$  of vertices, adding one vertex to  $S$  at a time. The algorithms stop when  $S = V$ . We will initialize  $S \leftarrow \{v_0\}$  where  $v_0$  is arbitrary for MST but  $v_0$  is the source for the shortest path instance. Define  $U := V \setminus S$  as the complement of  $S$ . To choose the next vertex  $u$  from  $U$  to be added to  $S$ , we put the vertices of  $U$  into a min-priority queue  $Q$ . To keep track of the membership in  $U$ , let  $\text{inU}[v \in V]$  be a global Boolean array where  $\text{inU}[v] = \text{true}$  iff  $v \in U$ . Our choice for the next vertex  $u$  is just the minimum vertex in  $Q$ . The priority of each  $v \in V$  will be defined by a global **min-cost array**

*$S$  as the source set,  
and  $U$  as the  
unknown set.*

$$\text{minC}[v \in V]$$

that associates a cost  $\text{minC}[v]$  with each  $v \in V$ . As discussed above, costs are non-negative:  $\text{minC}[v] \geq 0$ . The interpretation of  $\text{minC}[v]$  will vary for the two applications. We can initialize  $\text{minC}[v]$  to the cost of the edge  $v_0 - v$  for both MST and Dijkstra.

Suppose we just added vertex  $u$  to  $S$ . We must now update the array  $\text{minC}$  from  $u$ . Let us postulate a subroutine  $\text{UPDATE}(u)$  for this task. Two properties of  $\text{UPDATE}(u)$  are important to note: (1) We only need to update  $\text{minC}[v]$  for those vertices  $v$  that are adjacent to  $u$ . (2) Any updated value  $\text{minC}[v]$  must be a strict decrease in value. Since  $\text{minC}[v]$  is used as the priority of  $v$  in  $Q$ , it means that we need to call  $\text{decreaseKey}(v, \text{minC}[v], Q)$ .

#### GENERIC ALGORITHM:

##### ▷ *INITIALIZATION*

1. Initialize the arrays  $\text{inU}, \text{minC}$

##### ▷ *MAIN LOOP*

2. While  $Q \neq \emptyset$  do  
 3.  $u^* \leftarrow Q.\text{deleteMin}()$   
 4.  $\text{UPDATE}(u^*)$   
 5.  $\text{POSTPROCESS}(\text{minC})$

<sup>8</sup>The general treatment of min-cost paths may be found in Chapter XIV.

After the main loop, Line 5 will construct the necessary output using the data in the min-cost array `minC`. Usually, Prim's algorithm must output a min-cost spanning tree, and Dijkstra's algorithm must output a shortest path from  $s$  to each  $v \in V$ . But for the simplified versions of Prim and Dijkstra, we reduce our tasks to computing only the respective minimum costs: min-cost of a spanning tree, and min-cost of the paths  $s \cdots v$ . It will be easy to convert our min-cost algorithms into algorithms that actually construct some min-cost objects (Exercise). For this simplified version, the post-processing of `minC` is trivial in Dijkstra: we just output the array `minC` because `minC[v]` is just the min-cost of the path from  $s$  to  $v$ . For Prim, postprocessing amounts to computing  $\sum_{v \in V} \text{minC}[v]$ .

**¶40. Updates.** Let us now develop the details specific to Prim's Algorithm. This amounts to filling the details for `UPDATE(u)`. At the end, we will mention changes to achieve the same for Dijkstra's algorithm.

Recall that the notion of a set  $S \subseteq V$  being "Prim-good" (§V.6). If  $S$  is a singleton then  $S$  is Prim-good. Suppose  $S$  is Prim-good and we ask how  $S$  might be extended to a larger Prim-good set. There is an ordering of the vertices in  $S$  such that  $S = \{v_0, v_1, \dots, v_t\}$  ( $t = |S| + 1$ ) where  $v_i$  is added to  $S$  in the  $i$ th iteration. Write  $S_i = \{v_0, \dots, v_i\}$ . Since  $S_i$  is Prim-good, there is a set  $E_i$  of edges that forms the MST for  $G|_{S_i}$ . Let  $e_i$  be the edge such that  $E_i = E_{i-1} \cup \{e_i\}$ . In particular  $e_i$  has the form  $v_j - v_i$  for some  $j < i$ . We do not plan to explicitly maintain any of these information, but this implicit information will be necessary for understanding the correctness. Let us maintain the following invariant relative to the set  $S$ :

- For each  $v_i \in S = \{v_0, \dots, v_t\}$ , let `minC[vi]` be the cost of the edge  $e_i$ .
- For each  $u \in U$ , let

$$\text{minC}[w] := \min\{C[v, w] : (u-v) \in E, v \in S\}.$$

In order to find a node  $u^* \in V - S$  with the minimum `minC`-value, we will maintain  $V - S$  as a **single**<sup>9</sup> mergeable queue  $Q$  in which the least cost `minC[u]` serves as the key of the node  $u \in V - S$ . Hence extending the Prim-good set  $S$  by a node  $u^*$  amounts to a `deleteMin` from the mergeable queue. After the deletion, we must update the information `mst[S]` and `minC[v]` for each  $v \in V - S$ . But we do not really need to consider every  $v \in V - S$ : we only need to update `minC[v]` for those  $v$  that are adjacent to  $u^*$ . The following code fragment captures our intent.

```

UPDATE(u, S):
1.  inU[u] ← false.
2.  for w adjacent to u,
2.1    If inU[w] and (minC[w] > C[u, w]) then
2.2      minC[w] ← C[u, w].
2.3      DecreaseKey(w, minC[w], Q).

```

For Dijkstra's Algorithm, we just change Steps 2.1 and 2.2 to:

<sup>9</sup>So we are not using the full power of the mergeable queue ADT which can maintain several mergeable queues. In particular, we never perform the union operation in this application.

```

2.2      minC[w] ← minC[u] + C[u, w]}.
2.1      If inU[w] and (minC[u] + C[u, w]) then
2.2      minC[w] ← minC[u] + C[u, w] .

```

The correctness of Dijkstra's Algorithm will be proved in §XIV.3.

**¶41. Complexity Analysis.** The complexity parameters are  $n := |V|$  and  $m := |E|$ . We assume the mergeable queue is implemented by a Fibonacci heap. The initialization takes  $O(n)$  time. In the main while-loop, we perform  $n - 1$  iterations to add vertices to  $S$ . Each iteration calls `deleteMin` once, for a total cost of  $O(n \log n)$ . We account for the calls to `UPDATE(u)` separately: when calling `UPDATE(u)` subroutine, we call the `DecreaseKey` operation for each edge in  $u$ 's adjacency list. This gives a total of  $m$  such calls in all the `Update` calls. This has total cost  $O(m)$ . The other costs of `UPDATE(u)` is  $O(1)$  and can be charged to  $u$ , for an overall cost of  $O(n)$ . Thus the total cost to do the updates is  $O(m + n)$ . In the main procedure, we make  $n - 1$  passes through the Whileloop. So we perform  $n - 1$  `deleteMin` operations, and as the amortized cost is  $O(\log n)$  per operation, this has total cost  $O(n \log n)$ . The postprocessing work is  $O(n)$ . We have proven:

**Theorem 17** *The complexity of Prim's and Dijkstra's algorithm on a graph with  $n$  vertices and  $m$  edges is in  $O(m + n \log n)$ .*

## EXERCISES

Students should be able to demonstrate understanding of Prim's algorithm by doing hand simulations. The first exercise illustrates a simple tabular method for hand simulation.

**Exercise 4.1:** Use the Tabular Method to hand simulate Prim's algorithm on the following graph (Figure 27) beginning with  $v_1$ :




Figure 27: Graph of a House

The Tabular Method amounts to filling in the following table, row by row. We have filled in the first three rows already.



$i$	$v_2$	$v_3$	$v_4$	$v_5$	$v_6$	$v_7$	$v_8$	$v_9$	$v_{10}$	$v_{11}$	$v_{12}$	$mst[S]$	New Edge
1	<u>2</u>	3	2	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$	2	$(v_1, v_2)$
2	*	"	"	8	"	"	"	"	"	"	"	4	$(v_1, v_4)$
3	*	"	*	"	9	"	1	9	"	"	"	7	$(v_1, v_3)$

Note that the minimum cost in each row is underscored, indicating the item to be removed from the priority queue.  $\diamond$

**Exercise 4.2:** Repeat the simulation of the previous exercise, but this time, we implement Prim's Algorithm using Fibonacci heaps. and run the algorithm on the Graph in Figure 14 (Lect.VI). Show the state of your Fibonacci heap in each stage of the algorithm.  $\diamond$

**Exercise 4.3:** Let  $G_n$  be the graph with vertices  $\{1, 2, \dots, n\}$  and for  $1 \leq i < j \leq n$ , we have an edge  $(i, j)$  iff  $i$  divides  $j$ . For instance,  $(1, j)$  is an edge for all  $1 < j \leq n$ . The **cost** of the edge  $(i, j)$  is  $j - i$ .

(a) Hand simulate using the Tabular Method (see the first Exercise above). Show the final MST and its cost.




Figure 28:  $G_{10}$ : edges from node 1 are omitted for clarity.

(b) What can you say about the MST of  $G_n$ ? Is it unique? What is the asymptotic cost of the MST?  $\diamond$

**Exercise 4.4:** Modify the above algorithm for computing the minimum cost of a MST into one that constructs the MST.  $\diamond$

**Exercise 4.5:** Modify the above algorithm to compute a minimum spanning forest in case the input graph is not connected.  $\diamond$

**Exercise 4.6:** Let  $G = (V, E; \mu)$  be an edge-costed bigraph and  $T \subseteq E$ ,  $S \subseteq V$ . Let  $V(T) = \{v \in V : \exists u, (u, v) \in T\}$  denote the *vertices of*  $T$ , and  $G|S := (S, E'; \mu)$  where  $E' = E \cap \binom{S}{2}$  denote the *restriction of*  $G$  to  $S$ . We define  $T$  to be *prim-good* if  $T$  is an MST of  $G|V(T)$  and  $T$  can be extended into an MST of  $G$ . We define  $S$  to be *prim-good* if  $S$  is singleton or there exists a prim-good set  $T$  of edges such that  $S = V(T)$ . Show or give a counter-example:

- (a)  $T$  is a tree of  $G|V(T)$  and can be extended into an MST of  $G$  implies  $T$  is prim good.
- (b)  $S$  is prim-good implies every MST of  $G|S$  is prim-good.  $\diamond$

END EXERCISES

## §5. List Update Problem

The splay tree idea originates in the “move-to-front rule” heuristic for following **list update problem**: let  $L$  be a doubly-linked list of **items** where each item has a unique key. For simplicity, we usually write  $L$  as a sequence of keys. This list supports the **access request**. Each access request  $r$  is specified by a key (also denoted  $r$ ), and we satisfy this request by returning a pointer to the item in  $L$  with key  $r$ . (We assume such an item always exist.) We are interested in a special class of algorithms: such an algorithm  $\alpha$ , on an input  $L$  and  $r$ , searches sequentially in  $L$  for the key  $r$  by starting at the head of the list. Upon finding the item with key  $r$ ,  $\alpha$  is allowed to move the item to some position nearer the head of the list (the relative ordering of the other items is unchanged). Here are three alternative rules which specify the new position of an updated item:

- ( $R_0$ ) The **lazy rule** never modifies the list  $L$ .
- ( $R_1$ ) The **move-to-front** rule always make updated item the new head of the list  $L$ .
- ( $R_2$ ) The **transpose rule** just moves the updated item one position closer to the head of the list.

Let  $\alpha_i$  denote the list update algorithm based on Rule  $R_i$  ( $i = 0, 1, 2$ ). For instance,  $\alpha_1$  is the “move-to-front algorithm”. For any algorithm  $\alpha$ , let  $COST_\alpha(r, L)$  denote the cost of an update request  $r$  on a list  $L$  using  $\alpha$ . For  $i = 0, 1, 2$ , we write  $COST_i(r, L)$  instead of  $COST_{\alpha_i}(r, L)$ . We may define  $COST_i(r, L)$  to be  $1 + j$  where  $j$  is the position of the accessed item in  $L$ . If  $\alpha$  is an update algorithm, then  $\alpha(L, r)$  denotes the updated list upon applying  $\alpha$  to  $L, r$ . We extend this notation to a sequence  $U = \langle r_1, r_2, \dots, r_n \rangle$  of requests, by defining

$$\alpha(L, U) := \alpha(\alpha(L, \langle r_1, \dots, r_{n-1} \rangle), r_n).$$

Similarly,  $COST_\alpha(L, U)$  or  $COST_i(L, U)$  denotes the sum of the individual update costs.

**¶42. Example:** Let  $L = \langle a, b, c, d, e \rangle$  be a list and  $c$  an update request. Then  $\alpha_0(L, c) = L$ ,  $\alpha_1(L, c) = \langle c, a, b, d, e \rangle$  and  $\alpha_2(L, c) = \langle a, c, b, d, e \rangle$ . Also  $COST_i(L, c) = 4$  for all  $i = 0, 1, 2$ .

**¶43. Probabilistic Model.** We analyze the cost of a sequence of updates under the lazy rule and the move-to-front rule. We first analyze a probabilistic model where the probability of updating a key  $k_i$  is  $p_i$ , for  $i = 1, \dots, m$ . The lazy rule is easy to analyze: if the list is  $L = \langle k_1, \dots, k_m \rangle$  then the expected cost of a single access request is

$$C(p_1, \dots, p_m) = \sum_{i=1}^m i \cdot p_i.$$

It is easy to see that this cost is minimized if the list  $L$  is rearranged so that  $p_1 \geq p_2 \geq \dots \geq p_m$ ; let  $C^*$  denote this minimized value of  $C(p_1, \dots, p_m)$ .

What about the move-to-front rule? Let  $p(i, j)$  be the probability that  $k_i$  is in front of  $k_j$  in list  $L$ . This is the probability that, if we look at the last time an update involved  $k_i$  or  $k_j$ , the operation involves  $k_i$ . Clearly

$$p(i, j) = \frac{p_i}{p_i + p_j}.$$

The expected cost to update  $k_i$  is

$$1 + \sum_{j=1, j \neq i}^m p(j, i).$$

1574 The expected cost of an arbitrary update is

$$\begin{aligned} \hat{C} &:= \sum_{i=1}^m p_i \left[ 1 + \sum_{j=1, j \neq i}^m p(i, j) \right] \\ &= 1 + \sum_{i=1}^m \sum_{j \neq i}^m p_i \cdot p(i, j) \\ &= 1 + 2 \sum_{1 \leq j < i \leq m} \frac{p_i p_j}{p_i + p_j} \\ &= 1 + 2 \sum_{i=1}^m p_i \sum_{j=1}^{i-1} p(j, i) \\ &\leq 1 + 2 \sum_{i=1}^m p_i \cdot (i-1) \\ &= 2C^* - 1. \end{aligned}$$

1575 This proves

$$\hat{C} < 2C^*. \quad (31)$$

1576 Thus, the move-to-front rule achieves close to optimal bound without having to know in advance  
1577 the frequency distribution of requests, or to sort the lists.

1578 ¶44. **Amortization Model.** Let us now consider the amortized cost of a fixed sequence of  
1579 updates

$$U = (r_1, r_2, \dots, r_n) \quad (32)$$


1580 on an initial list  $L_0$  with  $m$  items. Clearly the worst case cost per update is  $O(m)$ . So, updates  
1581 over the sequence  $U$  costs  $O(mn)$ . This worst case bound cannot be improved if we use the  
1582 lazy rule. The best case for the lazy rule is  $O(1)$  per update, or  $O(n)$  overall.

1583 What about the move-to-front rule? In analogy to equation (31), we show that it is never  
1584 incur more than twice the cost of any update algorithm. In particular, it is never more than  
1585 twice cost of an optimal offline update algorithm  $\alpha_*$ . Note that  $\alpha_*$ , being offline, can determine  
1586 the best position to move each element after it has been accessed *based on the entire sequence*  
1587  *$U$  in (32).* If the cost of  $\alpha_*$  is denoted  $COST_*$ , we prove

$$COST_1(L, U) \leq 2 \cdot COST_*(L, U). \quad (33)$$

1588 We introduce the following potential function on lists. A pair  $(k, k')$  of keys is an **inversion**  
1589 in a pair  $(L, L')$  of lists if  $k$  occurs before  $k'$  in  $L$  but  $k$  occurs after  $k'$  in  $L'$ . We will compare  
1590 the list  $L$  produced by our move-to-front algorithm to the list  $L^*$  obtained from the optimal  
1591 algorithm: the **potential**  $\Phi(L)$  of  $L$  is defined to be the number of inversions in  $(L, L^*)$ . For  
1592 instance,  $\Phi(L) = 5$  in Figure 29 as there is one inversion involving  $a$ , two inversions involving  
1593  $b$  (not counting that with  $a$ ), two inversions involving  $c$  (not counting those with  $a$  or  $b$ ) and 0  
1594 inversion involving  $d$  (not counting those with  $a, b, c$ ).

Consider the  $j$ th request ( $j = 1, \dots, n$ ). Let  $L_j$  (resp.  $L_j^*$ ) be the list produced by the move-to-front (resp. optimal) algorithm after the  $j$ th request. Write  $\Phi_j$  for  $\Phi(L_j)$ . Let  $c_j$  and

Figure 29: How potential changes under update:  $L' = \alpha_1(L, b)$ 

$c_j^*$  denote the cost of serving the  $j$ th request under two algorithms (respectively). Let  $x_j$  be the item accessed in the  $j$ th request and  $k_j$  is the number of items that are in front of  $x_j$  in both lists  $L_j$  and  $L_j^*$ . Let  $\ell_j$  be the number of items that are in front of  $x_j$  in  $L_j$  but behind  $x_j$  in  $L_j^*$ . Hence

$$c_j = k_j + \ell_j + 1, \quad c_j^* \geq k_j + 1.$$

1595 CLAIM: The number of inversions destroyed is  $\ell_j$  and the number of inversions created is at  
 1596 most  $k_j$ .

In illustration, consider the list  $L'$  in Figure 29(iii), produced by the move-to-front algorithm after accessing  $b$  in  $L$ . Thus

$$x_j = b, \quad k_j = 1, \quad \ell_j = 2, \quad c_j = 4, \quad c_j^* = 2.$$

1597 Note that  $\ell_j$  counts the elements  $d$  and  $e$ . They represent the inversions  $\{b, e\}$  and  $\{b, d\}$  in  
 1598  $L$ . Both inversion were **destroyed** when  $b$  moved to the front in  $L'$ . Likewise,  $k_j$  counts the  
 1599 element  $a$ ; it represents the new inversion  $\{a, b\}$  that was **created** in  $L'$  when  $b$  moved to the  
 1600 front. Next, the optimal algorithm  $\alpha^*$  is allowed to update  $L'$  by moving  $b$  closer to the front  
 1601 of its list. Each element that  $b$  moves past when updating  $L'$  will reduce the number of created  
 1602 inversions by the move-to-front algorithm. It should now be clear that the above CLAIM is  
 1603 true.

It follows

$$\Phi_j - \Phi_{j-1} \leq k_j - \ell_j.$$

1604 Combining these two remarks,

$$\begin{aligned} c_j + \Phi_j - \Phi_{j-1} &\leq 2k_j + 1 \\ &\leq 2c_j^* - 1. \end{aligned}$$

1605 Summing up over all  $j = 1, \dots, n$ , we obtain

$$\begin{aligned} \text{COST}_1(L_0, U) &= \left( \sum_{j=1}^n c_j \right) + \Phi_n - \Phi_0 \\ &\leq \sum_{j=1}^n (2c_j^* - 1), \quad (\text{since } \Phi_n \geq 0, \Phi_0 = 0) \\ &= 2\text{COST}_*(L_0, U) - n. \end{aligned}$$

¶45. **Competitive Algorithms.** Let  $\beta(k)$  be a function of  $k$ . We say an algorithm  $\alpha$  is  $\beta(k)$ -**competitive** if there is some constant  $a$ , for all input lists  $L$  of length  $k$  and for all sequences  $U$  of requests

$$COST_\alpha(L, U) \leq \beta(k) \cdot COST_*(L, U).$$

Here  $COST_*$  is the cost incurred by the optimal offline algorithm.

We have just shown that the Move-to-Front algorithm is 2-competitive. This idea of competitiveness from Sleator and Tarjan is an extremely powerful one as it opens up the possibility of measuring the performance of online algorithms (such as the move-to-front algorithm) without any probabilistic assumption on the input requests.

¶46. **Final Remarks.** Amortization is closely related to two other topics. One is “self-organizing data structures”. Originally, this kind of analysis is undertaken by assuming the input has certain probability distribution. McCabe (1965) is the first to discuss the idea of move-to-front rule. See “An account of self-organizing systems”, W.J. Hendricks, *SIAM J.Comp.*, 5:4(1976); also “Heuristics that dynamically organizes data structures”, James R. Bitner, *SIAM J.Comp.*, 8:1(1979)82-100. But starting from the work of Sleator and Tarjan, the competitive analysis approach has become dominant. Albers and Westbrook gives a survey in [2]. Indeed, competitive analysis is connected to another major topic, “online algorithms”. Albers gives a survey [1]. Chung, Hajela and Seymour [3] determine that cost of the move-to-front rule over the cost of an optimal static ordering of the list (relative to some probability of accessing each item) is  $\pi/2$ . See also Lewis and Denenberg [6] and Purdom and Brown [9].

## EXERCISES

**Exercise 5.1:** We extend the list update problem above in several ways:

- (a) One way is to allow other kinds of requests. Suppose we allow insertions and deletions of items. Assume the following algorithm for insertion: we put the new item at the end of the list and perform an access to it. Here is the deletion algorithm: we access the item and then delete it. Show that the above analyses extend to a sequence of access, insert and delete requests.
- (b) Extend the list update analysis to the case where the requested key  $k$  may not appear in the list.
- (c) A different kind of extension is to increase the class of algorithms we analyze: after accessing an item, we allow the algorithm to transpose any number of pairs of adjacent items, where each transposition has unit cost. Again, extend our analyses above.  $\diamond$

**Exercise 5.2:** The above update rules  $R_i$  ( $i = 0, 1, 2$ ) are memoryless. The following two rules require memory.

- ( $R_3$ ) The **frequency rule** maintains the list so that the more frequently accessed items occur before the less frequently accessed items. This algorithm, of course, requires that we keep a counter with each item.
- ( $R_4$ ) The **timestamp rule** (Albers, 1995) says that we move the requested item  $x$  in front of the first item  $y$  in the list that precedes  $x$  and that has been requested at most once since the last request to  $x$ . If there is no such  $y$  or if  $x$  has not been requested so far, do not move  $x$ .

- (a) Show that  $R_3$  is not  $c$ -competitive for any constant  $c$ .
- (b) Show that  $R_4$  is 2-competitive.  $\diamond$

**Exercise 5.3:** (Bentley, Sleator, Tarjan, Wei) Consider the following data compression scheme based on any list updating algorithm. We encode an input sequence  $S$  of characters by each character's position in a list  $L$ . The trick is that  $L$  is dynamic: we update  $L$  by accessing each of the characters to be encoded. We now have a string of integers. To finally obtain a binary string as our output, we encode this string of integers by using a prefix code for each integer. In the following, assume that we use the move-to-front rule for list update. Furthermore, we use the prefix code of Elias in Exercise V.4.26 that requires only

$$f(n) = 1 + \lfloor \lg n \rfloor + 2 \lfloor \lg(1 + \lg n) \rfloor$$

bits to encode an integer  $n$ .

(a) Assume the characters are  $a, b, c, d, e$  and the initial list is  $L = (a, b, c, d, e)$ . Give the integer sequence corresponding to the string  $S = abaabcdabaabecbaadae$ . Also give the final binary string corresponding to this integer sequence.

(b) Show that if character  $x_i$  occurs  $m_i \geq 0$  times in  $S$  then these  $m_i$  occurrences can be encoded using a total of

$$m_i f(m/m_i)$$

bits where  $|S| = m$ . HINT: If the positions of  $x_i$  in  $S$  are  $1 \leq p_1 < p_2 < \dots < p_{m_i} \leq m$  then the  $j$ th occurrence of  $x_i$  needs at most  $f(p_j - p_{j-1})$ . Then use Jensen's inequality for the concave function  $f(n)$ .

(c) If there are  $n$  distinct characters  $x_1, \dots, x_n$  in  $S$ , define

$$A(S) := \sum_{i=1}^n \frac{m_i}{m} f\left(\frac{m}{m_i}\right).$$

Thus  $A(S)$  bounds the average number of bits per character used by our compression scheme. Show that

$$A(S) \leq 1 + H(S) + 2 \lg(1 + H(S))$$

where

$$H(S) := \sum_{i=1}^n \frac{m_i}{m} \lg\left(\frac{m}{m_i}\right).$$

NOTE:  $H(S)$  is the “empirical entropy” of  $S$ . It corresponds to the average number of bits per character achieved by the Huffman code for  $S$ . In other words, this online compression scheme achieves close to the compression of the offline Huffman coding algorithm.  $\diamond$

END EXERCISES

## §6. Amortized Subdivision Trees

Quadtrees are 2-dimensional generalizations of binary trees. For this purpose, it is useful to re-imagine binary trees as having nodes that are real intervals. Let  $T$  be a binary tree whose root is an interval  $I_0$ . In general, each node of  $T$  is an interval  $I = [a, b]$  and the children of  $I$  are the subintervals  $[a, (a+b)/2]$  and  $[(a+b)/2, b]$ . Thus the leaves of  $T$  is therefore collection of intervals that forms a subdivision of  $I_0$ .

In 2-dimensions, a quadtree has nodes that are boxes, where a box  $B$  is a Cartesian product of two intervals,  $B = I \times J$ . See Figure 30(b) for a quadtree. Each box  $B$  can be subdivided into four congruent subboxes called the **children** of  $B$ . A **quadtree** is a tree whose nodes are boxes, and each internal node  $B$  has four children which are congruent subboxes as described.




Figure 30: (a) Subdivision of a box and (b) associated quadtree

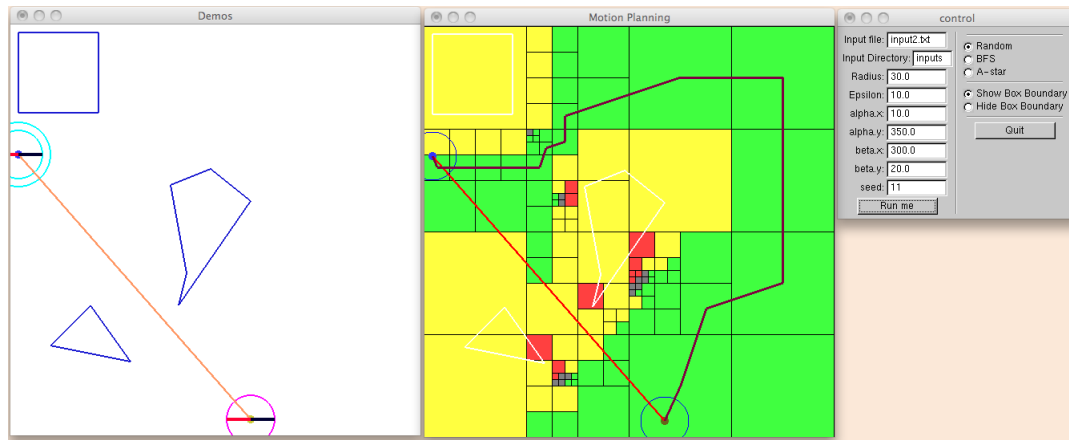


Figure 31: (a) start/goal position of disc robot amidst obstacles, (b) Subdivision Search for Path

The set of leaves of the quadtree forms a **subdivision** of the root box. Figure 30(a) shows the subdivision of the root box of the quadtree in Figure 30(b). One application of quadtrees is in robot motion planning, illustrated in Figure 31.

The boundary of a box can be divided four closed line segments called **sides**. Borrowing a cartographic term, we label the sides after the four (cardinal) compass directions:  $N, S, E, W$ . Two sides meet at a **corner**, which can be label with the four ordinal compass directions:  $NE, NW, SW, SE$ . Two boxes  $B, B'$  are  $k$ -**adjacent** if  $B \cap B'$  is a  $k$ -dimensional set. Note these possibilities:

- 2-adjacent: this means the interiors of  $B$  and  $B'$  intersect. If  $B, B'$  are nodes of a quadtree, this implies a containment relationship,  $B \subseteq B'$  or  $B' \subseteq B$ .
- 1-adjacent: this means  $B \cap B'$  is a line segment of positive length. In a quadtree,  $B \cap B'$  must be the side of either  $B$  or  $B'$  (or both).
- 0-adjacent:  $B \cap B'$  is a common corner of  $B$  and  $B'$ .
- $(-1)$ -adjacent:  $B$  and  $B'$  are disjoint.

For most applications, the 1-adjacency relationship between two boxes are the most important.

We will simply say “adjacent” to mean 1-adjacency. Say  $B'$  is **adjacent in direction**  $D \in \{N, S, E, W\}$  if  $B' \cap B$  is contained in the  $D$ -side of  $B$ . The  $D$ -**neighbors** of  $B$  are those leaf boxes which are adjacent in the direction  $D$ . In general,  $B$  will have at least four neighbors in each of the compass directions, unless  $B$  shares the boundary of the root  $B_0$ . The four directions have some underlying structure: the **opposite** of any direction  $D$  is defined as follows:  $N$  and  $S$  are opposites of each other, as are  $E$  and  $W$ . Two directions that are not opposites of each other are said to be **orthogonal**.

**¶47. Semi-Dynamic Quadrees.** So far, we have viewed a quadtree as a static data structure in which we can make neighbor queries. But most applications require a more dynamic data structure, in particular one that can grow by “splitting” at any leaf  $u$ . To split at  $u$ , we just add its four children to the tree. Now  $u$  is no longer a leaf, but these four children are new leaves. Let  $n$  be the size of the tree,  $n_i$  the number of internal nodes, and  $n_\ell$  the number of leaves. This split operation shows that  $n, n_i, n_\ell$  increases by 4, 1 and 3, respectively. Initially,  $(n, n_i, n_\ell) = (1, 0, 1)$ . Thus after  $k$  splits from the initial root box,  $(n, n_i, n_\ell) = (1+4k, k, 1+3k)$ . E.g., Figure 30 shows a quadtree with  $k = 6$ . Thus  $(n, n_i, n_\ell) = (25, 6, 19)$ . In particular, the ratio of size to number of leaves is  $(1+4k)/(1+3k)$  which approaches  $4/3$  as  $k$  grows. Note that there are some applications where we only need the leaves of the quadtree i.e., the underlying subdivision of the root  $B_0$ . This analysis shows that by maintaining the entire quadtree instead of just the leaves, we incur an overhead of 33.3% in space usage. In the following, we shall assume that it is not an issue.

**Partial Quadrees.** There are applications where may not need even the full set of leaves. For example, in motion planning, boxes represent a set of robot positions. Some robot positions are forbidden, and we classify boxes as “OFF” if all the positions in that box is forbidden, otherwise it is “ON”. If we are interesting in path finding, we may<sup>a</sup> ignore the “OFF” nodes. This is illustrated in Figure 30 where the green boxes are ON, and yellow boxes are OFF.

More generally, let us say that we are given a box predicate  $ON(B)$  that evaluates to true iff the box  $B$  is of interest to the application. We say  $B$  is either “ON” or “OFF” accordingly. Moreover, the predicate is hereditary in the sense that if  $B$  is OFF, then any subbox  $B' \subseteq B$  is also OFF. In this case, we can maintain a **partial quadtree** where we only retain those leaves that are ON; now each internal node between 1 and 4 children. Such a tree might have a path in which all the nodes have only one child. There are techniques for compressing the representation in this case.

<sup>a</sup>This is not entirely true; if we wish to detect non-existence of paths in finite time, we may need to maintain the OFF boxes to detect such conditions.

Let us consider the problem of maintain a binary tree, subject to splitting a leaf, and neighbor queries. In order to answer neighbor queries, we need ma

## EXERCISES

**Exercise 6.1:** Suppose a quadtree has  $L$  leaves and size  $n$ . Give a simple expression for  $n$  as a function of  $L$ .  $\diamond$

**Exercise 6.2:** Suppose for each box  $u$  in our quadtree, we maintain four **adjacency pointers**,



$u.D$  ( $D \in \{N, S, E, W\}$ ) which point to a box adjacent in direction  $D$ . The box  $u.D$  is uniquely defined if we require it to have the largest depth that is at most  $\text{depth}(u)$ . Of course, if  $u$  has no adjacent boxes in direction  $D$ , we let  $u.D = \text{nil}$ . Show that we can find all the neighbors of any box  $u$  in a quadtree in amortized  $O(1)$  time per neighbor.

◇

END EXERCISES

## §7. Continuous Amortization

We consider the highly classical problem of root finding: given a real function  $f : \mathbb{R} \rightarrow \mathbb{R}$ , to compute all the real roots of  $f$ . The **roots** or **zeros** of  $f$  are those  $\alpha \in \mathbb{R}$  such that  $f(\alpha) = 0$ . Equivalently, we want to solve the equation

$$f(x) = 0.$$

There are many variations of this problem: if  $f(x) = ax^2 + bx + c$ , then we know from high school the quadratic formula for the roots:  $\alpha = (-b \pm \sqrt{b^2 - 4ac})/2a$ . But Abel tells us that when  $f$  is a polynomial of degree more than 5, we cannot expect to find analogous formula in terms of radicals. Hence, we seek numerical approximations. Here are some results from the computer algebra system **Maple**: if

$$f := 2x^5 - 11x^4 - 7x^3 + 12x^2 - 4x$$

then calling **fsolve**( $f$ ) will produce the result:

So **Maple** gives us three approximate roots (including the root  $\alpha = 0$  which we could already expect just by looking at  $f$ ). The above output shows 10 digits of significant digits, by default. Here is another example from **Maple**: if

$$g := x^{20} - (20x - 1)^2$$

then **fsolve**( $g$ ) produces four real roots:

$$-1.400398846, 0.05000000000, 0.05000000000, 1.389281264$$

This example show two roots that are numerically identical up to 10 significant digits, both roughly equal to 0.05. Are they identical roots or is their difference less than  $10^{-10}$ ? Of course, if there are not identical, we can keep asking **Maple** for more and more significant digits until the difference is revealed. But if they are identical, this process will not terminate. In algebraic computation, we therefore prefer to approximate  $\alpha$  by an **isolation interval**, namely an interval  $[a, \bar{a}]$  that contains exactly one root of  $f(x)$  and that root is  $\alpha$ . So we pose the **real root isolation problem** as follows: *given a real function  $f$  and an interval of interest  $I_0$ , to compute an isolating interval for each real root of  $f$  in  $I_0$ .*

Several remarks: We may choose  $I_0$  to be the real line in which case we are interested in all the real roots of  $f$ . If we further specialize the function  $f$  to be polynomials, and the coefficients of  $f$  specialized to integers, we will call it the **benchmark problem**. This critical case has occupied theoretical Computer Science since Schönhage's 1982 landmark unpublished paper [10] in which he established that the benchmark problem has bit complexity  $\tilde{O}(d^3 L)$  where  $d$  is the degree of  $f$  and its coefficients have magnitude  $< 2^L$ . When the coefficients of  $f$  are

real numbers for which you can ask for arbitrarily good approximations (this is also known as the bit-stream model) is important for many applications in computer algebra (the coefficients here are actually real algebraic numbers). Of course we can also isolate the complex roots of  $f$ . Another important generalization is to look at multivariate functions. Suppose we have  $m$  polynomials  $f_1, \dots, f_m \in \mathbb{Z}[X_1, \dots, X_n]$  in  $n$  variables. We want to solve the simultaneous equation  $0 = f_1 = f_2 = \dots = f_m$ . The two main cases here are (i)  $m = n$  and there are finitely many zeroes, (ii)  $m = 1$  and the zero set is a  $n - 1$  dimensional surface. We must interpret (ii) correctly: since the zero set is infinite, we can only approximate the surface. The analogue of root isolation is to compute an approximate surface with the correct topology (technically, the approximate surface should be isotopic to the zero set).

**¶48. Generic Subdivision Algorithm.** We shall focus on a very simple class of algorithms for root isolation. Suppose we have a interval predicate  $C_*$  such that for any interval  $I$ , we have  $C_*(I) \in \mathbb{N}$  with the properties

- $C_*(I) = 0$  implies  $I$  contains no zeros of  $f$ .
- $C_*(I) = 1$  implies  $I$  contains one zeros of  $f$ .

Consider the following algorithm that uses two queue data structures,  $Q$  and  $P$ :

```

GENERIC SUBDIVISION ALGORITHM:
Input:  $f$  and  $I_0$ 
Output: A set  $P$  of intervals
   $Q \leftarrow \{I_0\}, P \leftarrow \emptyset$ 
  WHILE  $Q \neq \emptyset$ 
     $I \leftarrow Q.remove()$ 
    If  $(C_*(I) < 2)$ 
       $P.append(I)$ 
    Else
      Let  $I = [a, b]$  and  $m = (a + b)/2$ 
       $Q.append([a, m]), Q.append([m, b])$ 

```

This algorithm creates a **subdivision tree** rooted at  $I_0$  such that each internal node is an interval  $[a, b]$  with two children  $[a, (a + b)/2]$  and  $[(a + b)/2, b]$ . At any moment, the set of the intervals in  $P \cup Q$  forms a **subdivision** of  $I_0$ , i.e.,  $I_0 = \cup I : I \in P \cup Q$  and the interiors of two intervals in  $P \cup Q$  are disjoint.

Upon termination, the intervals in  $P$  forms a partition of  $I_0$ . Every root  $\alpha$  of  $f$  in  $I_0$  will be represented by some isolating interval in  $P$ . But notice that if two intervals in  $P$  shares an endpoint  $m$  that happens to be a root, then we must choose only one of these two intervals for a proper output for the root isolation problem. But by a post processing of  $P$ , we can easily resolve this issue. Another important remark to make is that we assume that it is possible to determine the sign of  $f$  at the endpoint of any subdivision interval. This assumption is quite nontrivial. Assuming that  $I_0$  initially has endpoints that are dyadic numbers, we are assuming that the exact sign of  $f$  at dyadic numbers are possible. The Exercise explores alternatives that avoid this assumption.

The main question is whether this algorithm terminates. Suppose  $C_*$  is convergent in the following sense: for  $p \in \mathbb{R}$ , let  $C_*(p) = 0$  if  $f(p) \neq 0$  and  $C_*(p) = 1$  if  $f(p) = 0$ . for any sequence

of intervals  $I_1, I_2, \dots$  that converges to a point  $p$ , we have  $C_*(I_i) \rightarrow C_*(p)$  (this means that  $C_*(I_i) = C_*(p)$  for all  $i$  large enough. We have the following easy lemma:

**Lemma 18** *If  $C_*$  is convergent, then the Generic Subdivision Algorithm always halts. Upon halting,  $P$  contains (among other things) an isolating interval for each root of  $f$  in  $I_0$ .*

**¶49. Sturm Method.** Suppose  $f$  is a polynomial. The theory of Sturm sequences can provide a predicate  $C_*(I)$  that counts exactly the number of real roots of  $f$  inside  $I_0$ . The predicate is clearly convergent.

Give an Example.

The subdivision tree produced in this case is the minimal subdivision tree possible. It can be shown that for the benchmark problem, the tree size is  $O(d(L + \log d))$ . Unfortunately, in practice, this methods based on Sturm sequences is not competitive.

**¶50. Descartes Method.** Suppose  $f = \sum_{i=0}^d a_i X^i$  is a polynomial. The **sign variation** of  $f$  is the number of sign changes in the sequence  $(a_0, a_1, \dots, a_d)$ , where  $(a_i, a_j)$  counts as a sign change if  $i < j$ ,  $a_i a_j < 0$  and  $0 = a_{i+1} = \dots = a_{j-1}$ . Denote this number by  $Var(f)$ .

**Lemma 19** ([ Descartes' Rule of Sign] *Suppose  $f$  has  $k$  positive real roots. Then  $Var(f) - k$  is a non-negative even number.*

Thus,  $Var(f)$  is an upper bound on the number of real positive roots of  $f$ . What if we are interested in roots of  $f$  in some interval  $I = [a, b]$ , not just the positive real roots of  $f$ ? Well, we can apply a Mobius transformation: let

$$f_I(X) := \frac{1}{(X+1)^d} f\left(\frac{aX+b}{X+1}\right).$$

Then the set of positive real root of  $f_I$  is mapped bijectively to the real roots of  $f$  in  $I$ . We can therefore apply Descartes rule of sign to  $f_I$  to estimate the number of roots of  $f$  in  $I$ .

Example.

We now instantiate the predicate  $C_*(I)$  in the Generic Subdivision Algorithm by the **Descartes predicate**:

$$C_*(I) := Var(f_I).$$

By Lemma 19, if  $C_*(I) = 0$ , there are no roots of  $f$  in  $I$ , and if  $C_*(I) = 1$ , there is exactly one real root of  $f$  in  $I$ . We can further prove:

**Lemma 20** *The Descartes predicate,  $C_*(I) = Var(f_I)$  is convergent.*

Currently, the Descartes method for root isolation is one of the most efficient exact algorithms. Moreover, one can show [4] that the tree size for the Descartes method in the benchmark problem is also  $O(d(L + \log d))$ .

**¶51. Bolzano Method.** The following result is most intuitive: *if  $f$  is a continuous real function, and  $a \neq b$  with  $f(a)f(b) < 0$ , then there exists some  $c$  between  $a$  and  $b$  such that  $f(c) = 0$ .*

This statement goes back to Bolzano (1817) which is a form of the Intermediate Value Theorem. We propose to develop predicate  $C_*(I)$  based on this simple principle. The advantage is that  $f$  can be quite general functions. The basis for such methods is interval arithmetic: if  $f : \mathbb{R} \rightarrow \mathbb{R}$  is any function, then its set theoretic extension (still denoted by  $f$ ) is defined by  $f(S) := \{f(x) : x \in S\}$  for any  $S \subseteq \mathbb{R}$ . Now, a function of the form  $\Box f : \Box\mathbb{R} \rightarrow \Box\mathbb{R}$  is called an **interval version** of  $f$  if for all  $I \in \Box\mathbb{R}$ ,  $f(I) \subseteq \Box f(I)$ . If for all interval sequences  $I_1, I_2, \dots$  that converges to a point  $p \in \mathbb{R}$ , we have  $\Box f(I_i) \rightarrow f(p)$  then we say  $\Box f$  is **convergent**. A **box function** for  $f$  is any interval version of  $f$  that is convergent. We now define

$$C_*(I) = \begin{cases} 0 & \text{if } 0 \notin \Box f(I), \\ 1 & \text{if } 0 \notin \Box f'(I), \text{ and } f(a)f(b) < 0 \text{ where } I = [a, b], \\ 2 & \text{else.} \end{cases} \quad (34)$$

Here  $\Box f$  and  $\Box f'$  are box functions for  $f$ . Using interval arithmetic, it is easy to construct such box functions for polynomials.

By the EVAL algorithm, we mean the Generic algorithm instantiated by (34). Our goal is to sketch out a proof that EVAL also produces an optimal tree size of  $O(d(L + \log d))$  for the benchmark problem. But we need to be careful: the result cannot be true for arbitrary box functions. It is true for the “idea

## References

- [1] S. Albers. Competitive online algorithms. BRICS Lecture Series LS-96-2, BRICS, Department of Computer Science, University of Aarhus, September 1996.
- [2] S. Albers and J. Westbrook. A survey of self-organizing data structures. Research Report MPI-I-96-1-026, Max-Planck-Institut für Informatik, Im Stadtwald, D-66123 Saarbrücken, Germany, October 1996.
- [3] F. R. K. Chung, D. J. Hajela, and P. D. Seymour. Self-organizing sequential search and hilbert’s inequalities. *ACM Symp. Theory of Comput.*, 7, 1985. Providence, Rhode Island.
- [4] A. Eigenwillig, V. Sharma, and C. Yap. Almost tight complexity bounds for the Descartes method. In *31st Int’l Symp. Symbolic and Alge. Comp. (ISSAC’06)*, pages 71–78, 2006. Genova, Italy. Jul 9-12, 2006.
- [5] D. G. Kirkpatrick and R. Seidel. The ultimate planar convex hull algorithm? *SIAM J. Comput.*, 15:287–299, 1986.
- [6] H. R. Lewis and L. Denenberg. *Data Structures and their Algorithms*. Harper Collins Publishers, New York, 1991.
- [7] A. Moffat, G. Eddy, and O. Petersson. Splaysort: Fast, versatile, practical. *Software - Practice and Experience*, 126(7):781–797, 1996.
- [8] M. H. Overmars and J. van Leeuwen. Dynamic multi-dimensional data structures based on quad- and  $k$ -d trees. *Acta Inform.*, 17:267–285, 1982.
- [9] J. Paul Walton Purdom and C. A. Brown. *The Analysis of Algorithms*. Holt, Rinehart and Winston, New York, 1985.

- 
- [10] A. Schönhage. The fundamental theorem of algebra in terms of computational complexity, 1982. Manuscript, Department of Mathematics, University of Tübingen. Updated in 2004 with typo corrections and an appendix of subsequent papers.
- [11] M. Smerk. Self-adjusting  $k$ -ary search trees. In *LNCS*, volume 382, pages 373–380, 1989. Proc. Workshop on Algorithms and Data Structures, Aug. 17-19, 1989, Carleton University, Ottawa, Canada.
- [12] D. D. Sleator and R. E. Tarjan. Self-adjusting binary search trees. *J. ACM*, 32:652–686, 1985.
- [13] R. E. Tarjan. Amortized computational complexity. *SIAM J. on Algebraic and Discrete Methods*, 6:306–318, 1985.